

Live Coding

```
import pyspark as ps
import random

# initialize SparkContext
sc = ps.SparkContext()

flips = 1000000

# lazy eval
coins = xrange(flips)

# lazy eval, nothing executed
heads = sc.parallelize(coins) \
    .map(lambda i: random.random()) \
    .filter(lambda r: r < 0.51) \
    .count()

print "{0} heads flipped!".format(heads)

# shutdown SparkContext
sc.stop()
```



What is an RDD?

- **Resilient:** if the data in memory (or on a node) is lost, it can be recreated
- **Distributed:** data is chunked into partitions and stored in memory across the cluster
- **Dataset:** initial data can come from a file or be created programmatically

Note: RDDs are read-only and immutable, we will come back to this later...

Functions Deconstructed

```
import random  
flips = 1000000
```

```
# lazy eval  
coins = range(flips)
```

Python Generator

```
# lazy eval, nothing executed  
heads = sc.parallelize(coins) \
```

Transformations →

```
.map(lambda i: random.random()) \  
.filter(lambda r: r < 0.51) \  
.count()
```

Create RDD

Action (materialize result)



Functions Deconstructed

```
import random
flips = 1000000

# lazy eval
coins = range(flips)

# lazy eval, nothing executed
heads = sc.parallelize(coins) \
    .map(lambda i: random.random()) \
    .filter(lambda r: r < 0.51) \
    .count()

# create a closure with the lambda function
# apply function to data
```

Closures

Functions Deconstructed

```
import random
flips = 1000000

# local sequence
coins = range(flips)

# distributed sequence
coin_rdd = sc.parallelize(coins)
flips_rdd = coin_rdd.map(lambda i: random.random())
heads_rdd = flips_rdd.filter(lambda r: r < 0.51)

# local value
head_count = heads_rdd.count()
```

Functional Programming Primer

- Functions are applied to **data (RDDs)**
- RDDs are Immutable: $f(\text{RDD}) \rightarrow \text{RDD2}$
- Function application necessitates creation of new **data**

Spark Functions

Transformations

Lazy Evaluation
(does not immediately evaluate)

Returns new RDD

Actions

Materialize Data
(evaluates RDD lineage)

Returns final value
(on driver)

Transformations

```
# Every Spark application requires a Spark Context
# Spark shell provides a preconfigured Spark Context called `sc`
nums = sc.parallelize([1,2,3])

# Pass each element through a function
squared = nums.map(lambda x: x*x) # => {1, 4, 9}

# Keep elements passing a predicate
even = squared.filter(lambda x: x % 2 == 0) # => [4]

# Map each element to zero or more others
nums.flatMap(lambda x: range(x)) # => {0, 0, 1, 0, 1, 2}
```

Actions

```
nums = sc.parallelize([1, 2, 3])
```

```
# Retrieves RDD contents as a local collection  
nums.collect() # => [1, 2, 3]
```

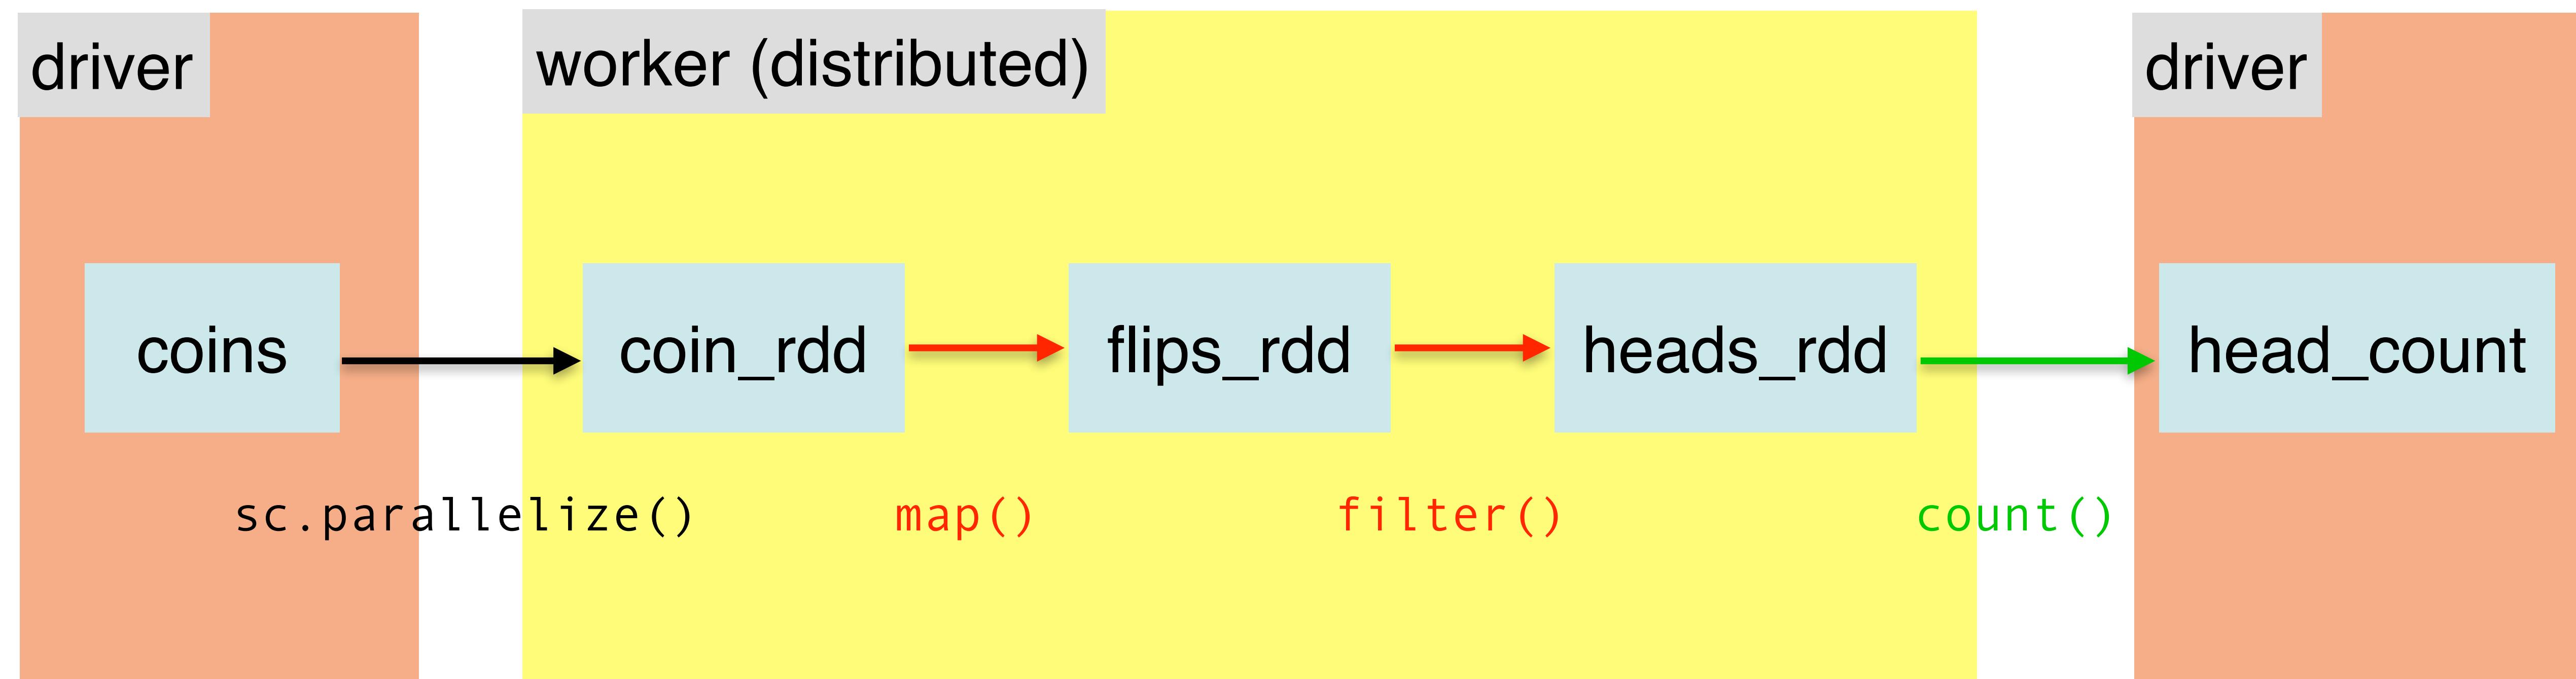
```
# Returns first K elements  
nums.take(2) # => [1, 2]
```

```
# Count number of elements  
nums.count() # => 3
```

```
# Merge elements with an associative function  
nums.reduce(lambda: x, y: x + y) # => 6
```

```
# Write elements to a text file  
nums.saveAsTextFile("hdfs://file.txt")
```

RDD Lineage



Functions Deconstructed

```
import random  
flips = 1000000
```

```
# lazy eval  
coins = xrange(flips)
```

nothing runs here

```
# lazy eval, nothing executed  
heads_rdd = sc.parallelize(coins) \  
    .map(lambda i: random.random()) \  
    .filter(lambda r: r < 0.51)
```

```
head_count = heads_rdd.count()
```



What is an RDD?

An *Abstraction!*

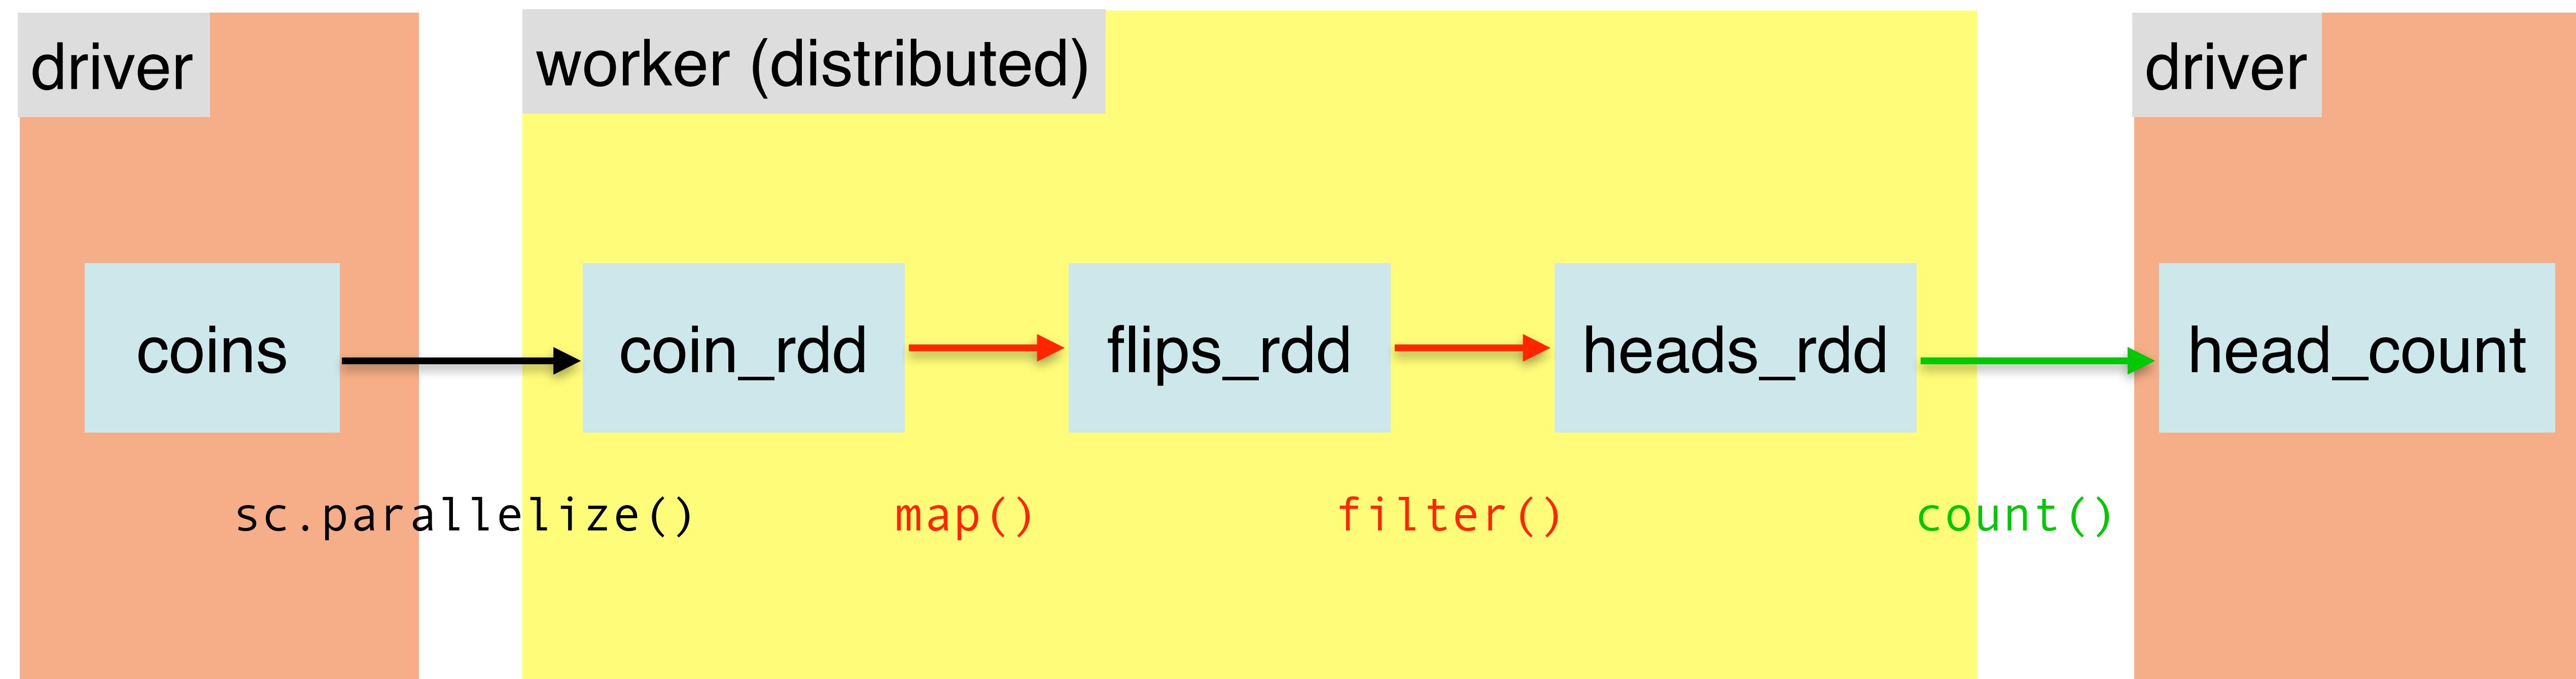
What is an RDD?

Lineage (required)

1. Set of partitions for current RDD (data)
 2. List of dependencies
 3. Function to compute partitions (functional paradigm)
-
4. Partitioner to optimize execution
 5. Potential preferred location for partitions

Optimization (optional)

RDD Lineage



What is an RDD?

Lineage (required)

1. Set of partitions for current RDD (data)
 2. List of dependencies
 3. Function to compute partitions (functional paradigm)
-
4. Partitioner to optimize execution
 5. Potential preferred location for partitions

Optimization (optional)

RDD as an interface

<i>Operation</i>	<i>Meaning</i>
<code>partitions()</code>	Return a list of Partition objects
<code>dependencies()</code>	Return a list of dependencies
<code>compute(p, parent)</code>	Compute the elements of Partition p given its parent Partitions
<code>partitioner()</code>	Return metadata specifying whether this RDD is hash/range partitioned
<code>preferredLocations(p)</code>	List nodes where Partition p can be accessed quicker due to data locality

dummy.txt

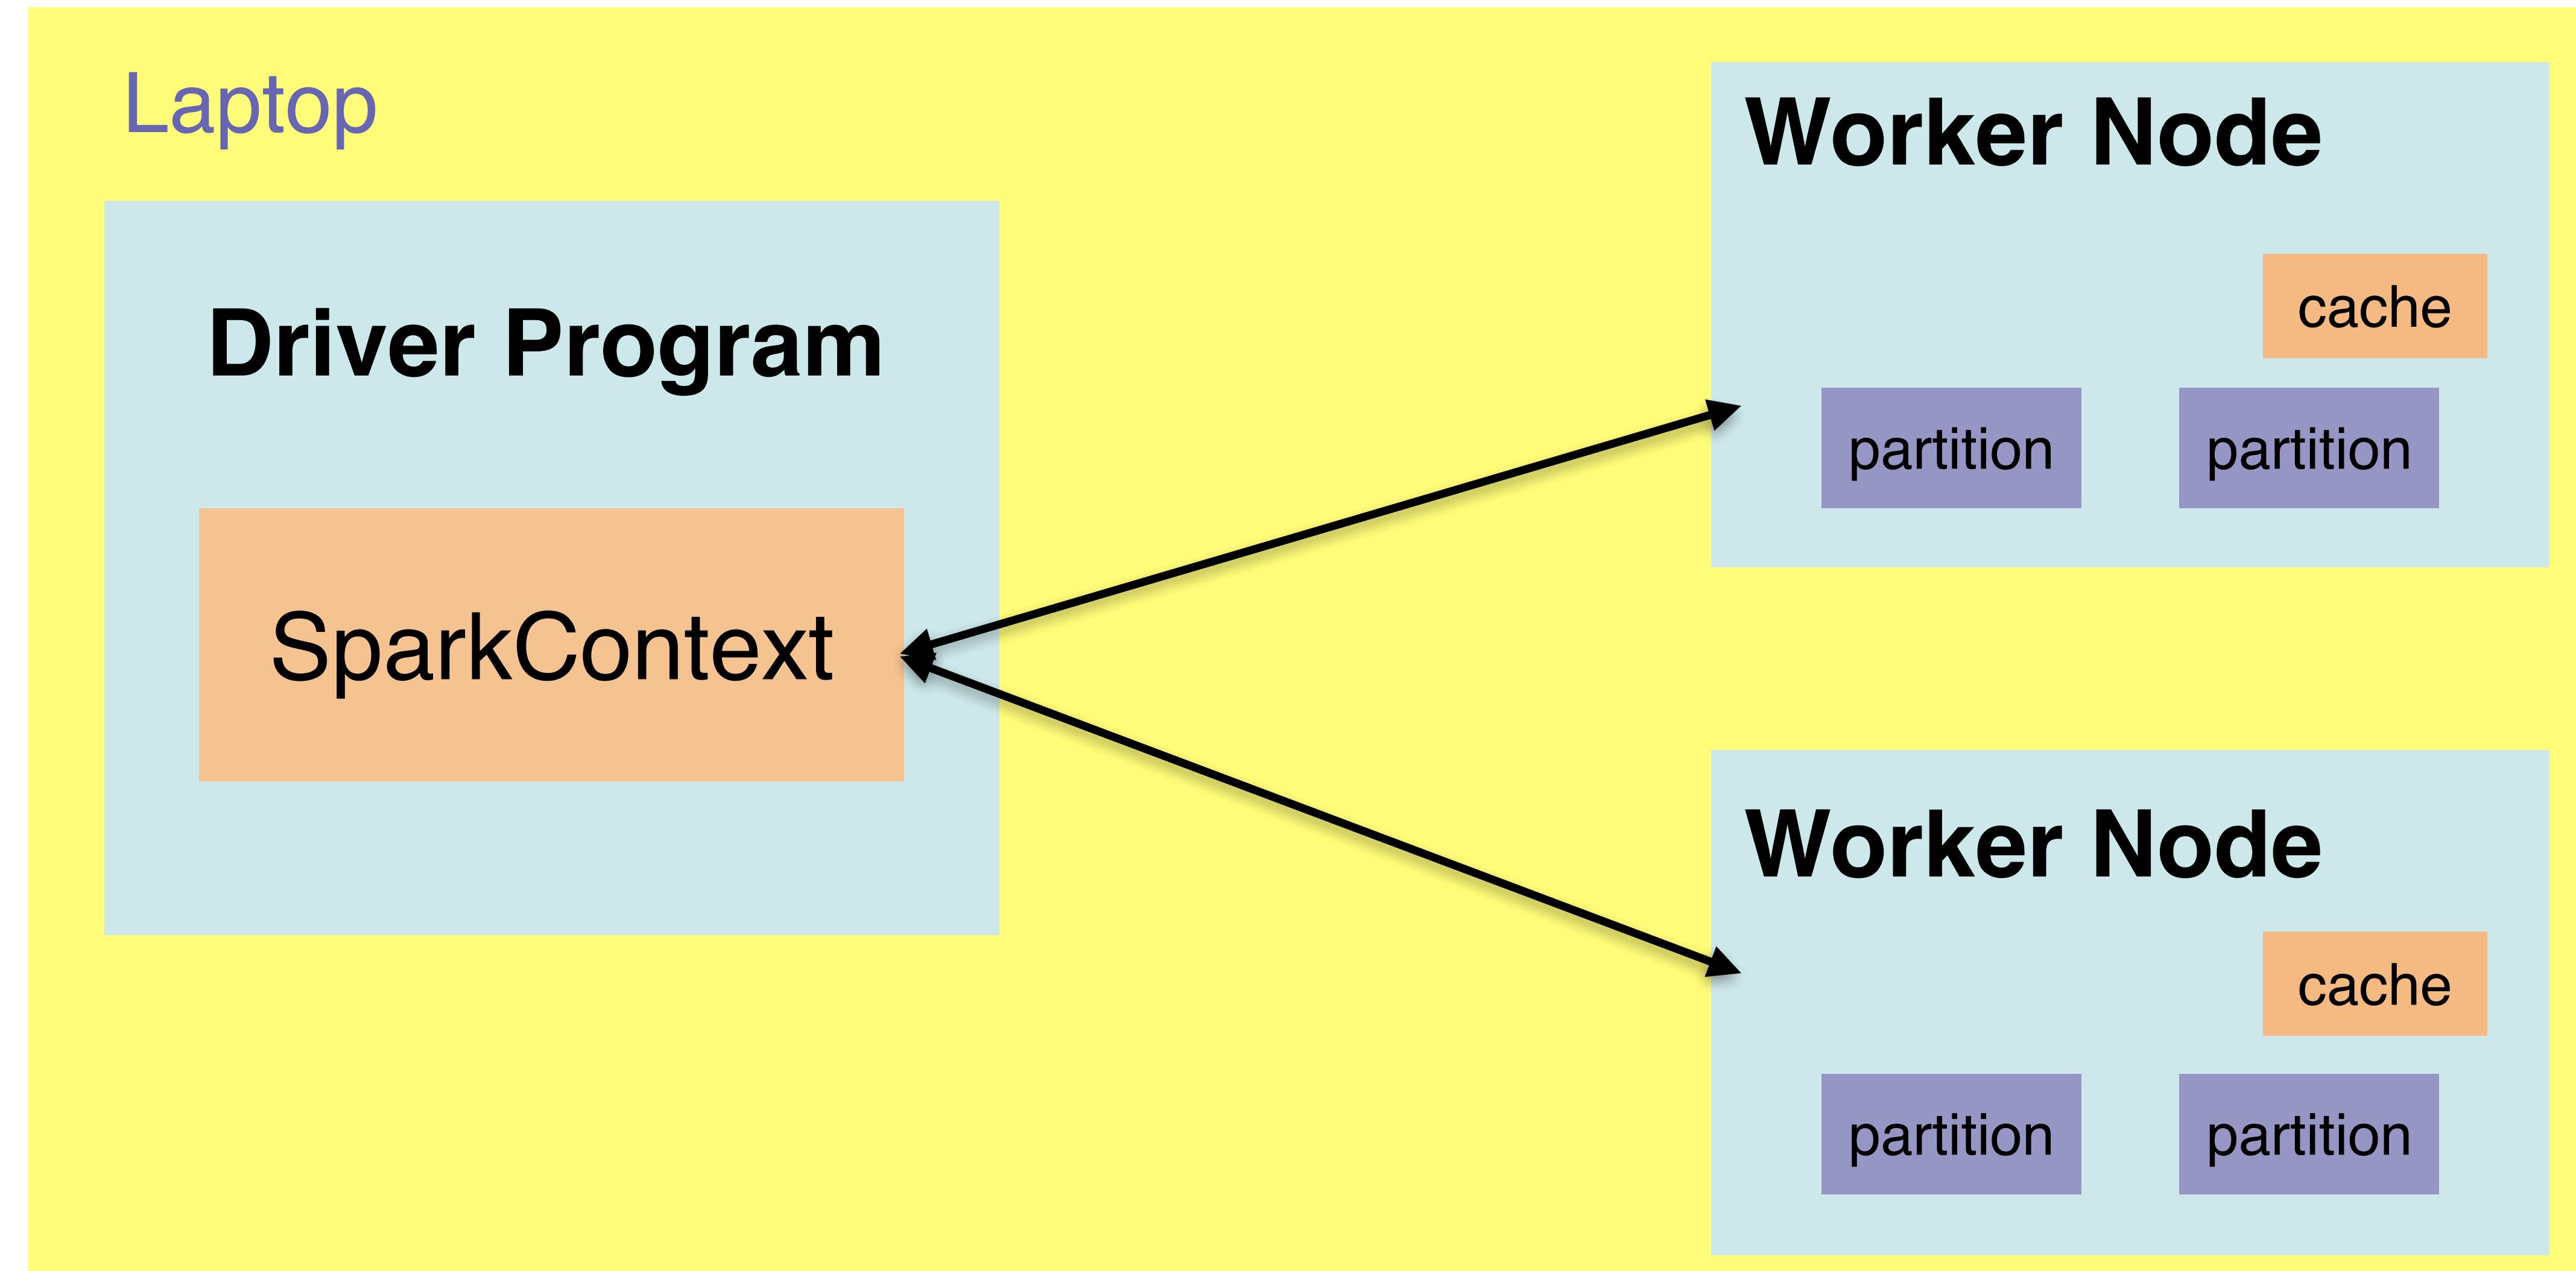
```
file_rdd.map(line => line.split(""))
.map(split => (split(0), split(1).toInt))
.groupByKey()
.mapValues(iter => iter.reduce(_ + _)).collect()
```

jon	2
mary	3
anna	1
jon	1
jesse	3
mary	5

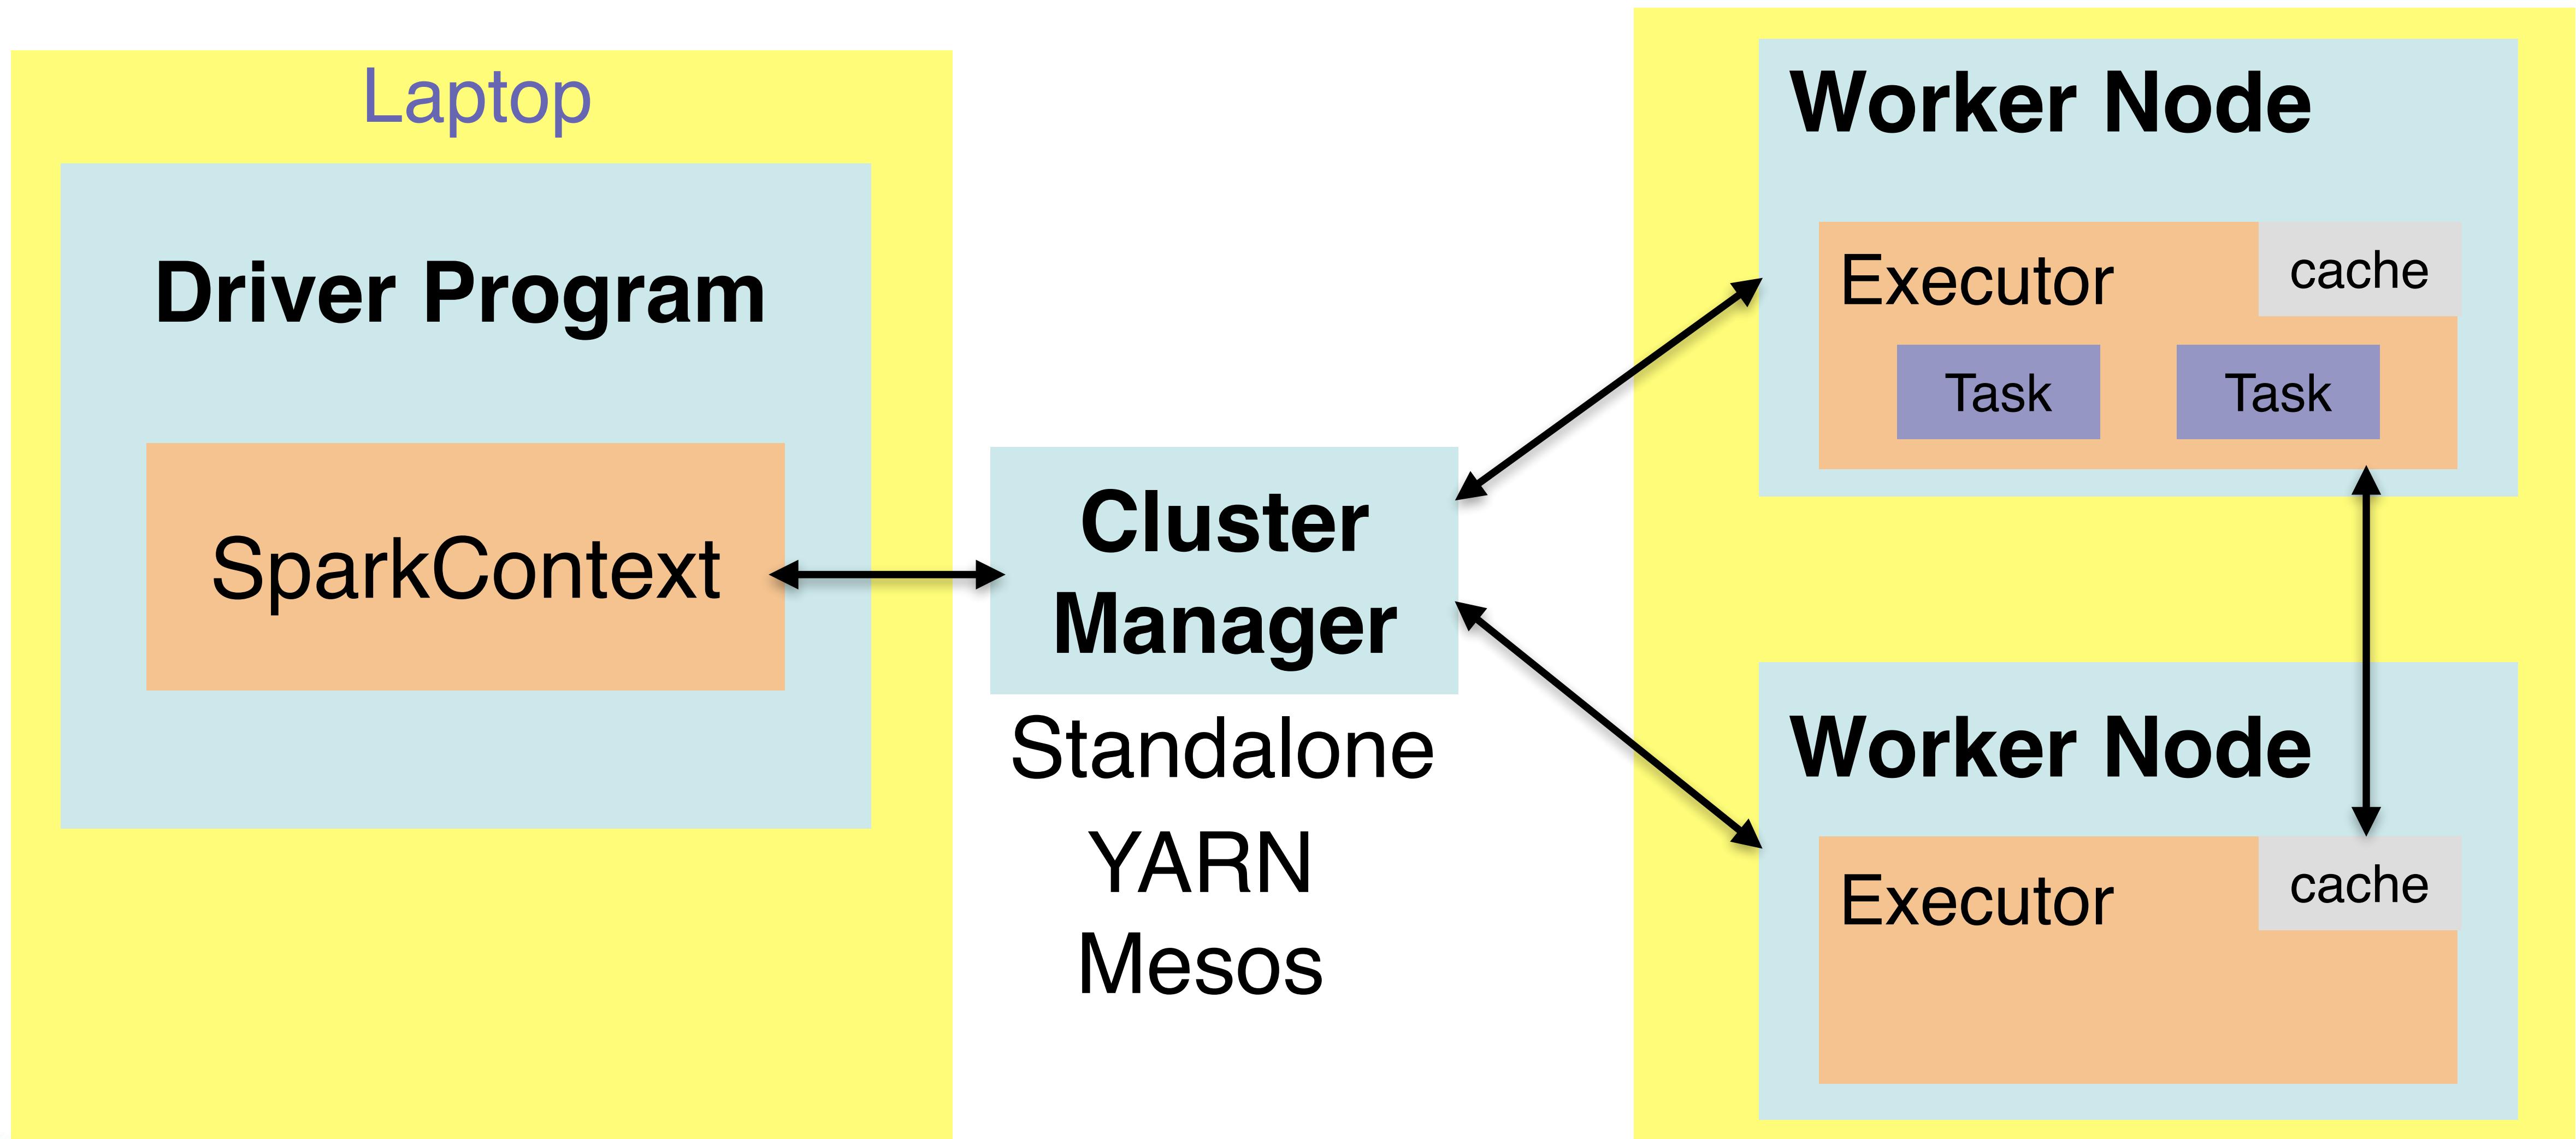
Key-Value Operations

```
pets = sc.parallelize([( "cat", 1), ( "dog", 1), ( "cat", 2)])  
  
pets.reduceByKey(lambda x, y: x + y) # => { (cat, 3), (dog, 1) }  
  
pets.groupByKey() # => { (cat, [1, 2]), (dog, [1]) }  
  
pets.sortByKey() # => { (cat, 1), (cat, 2), (dog, 1) }
```

Local Spark Execution Context



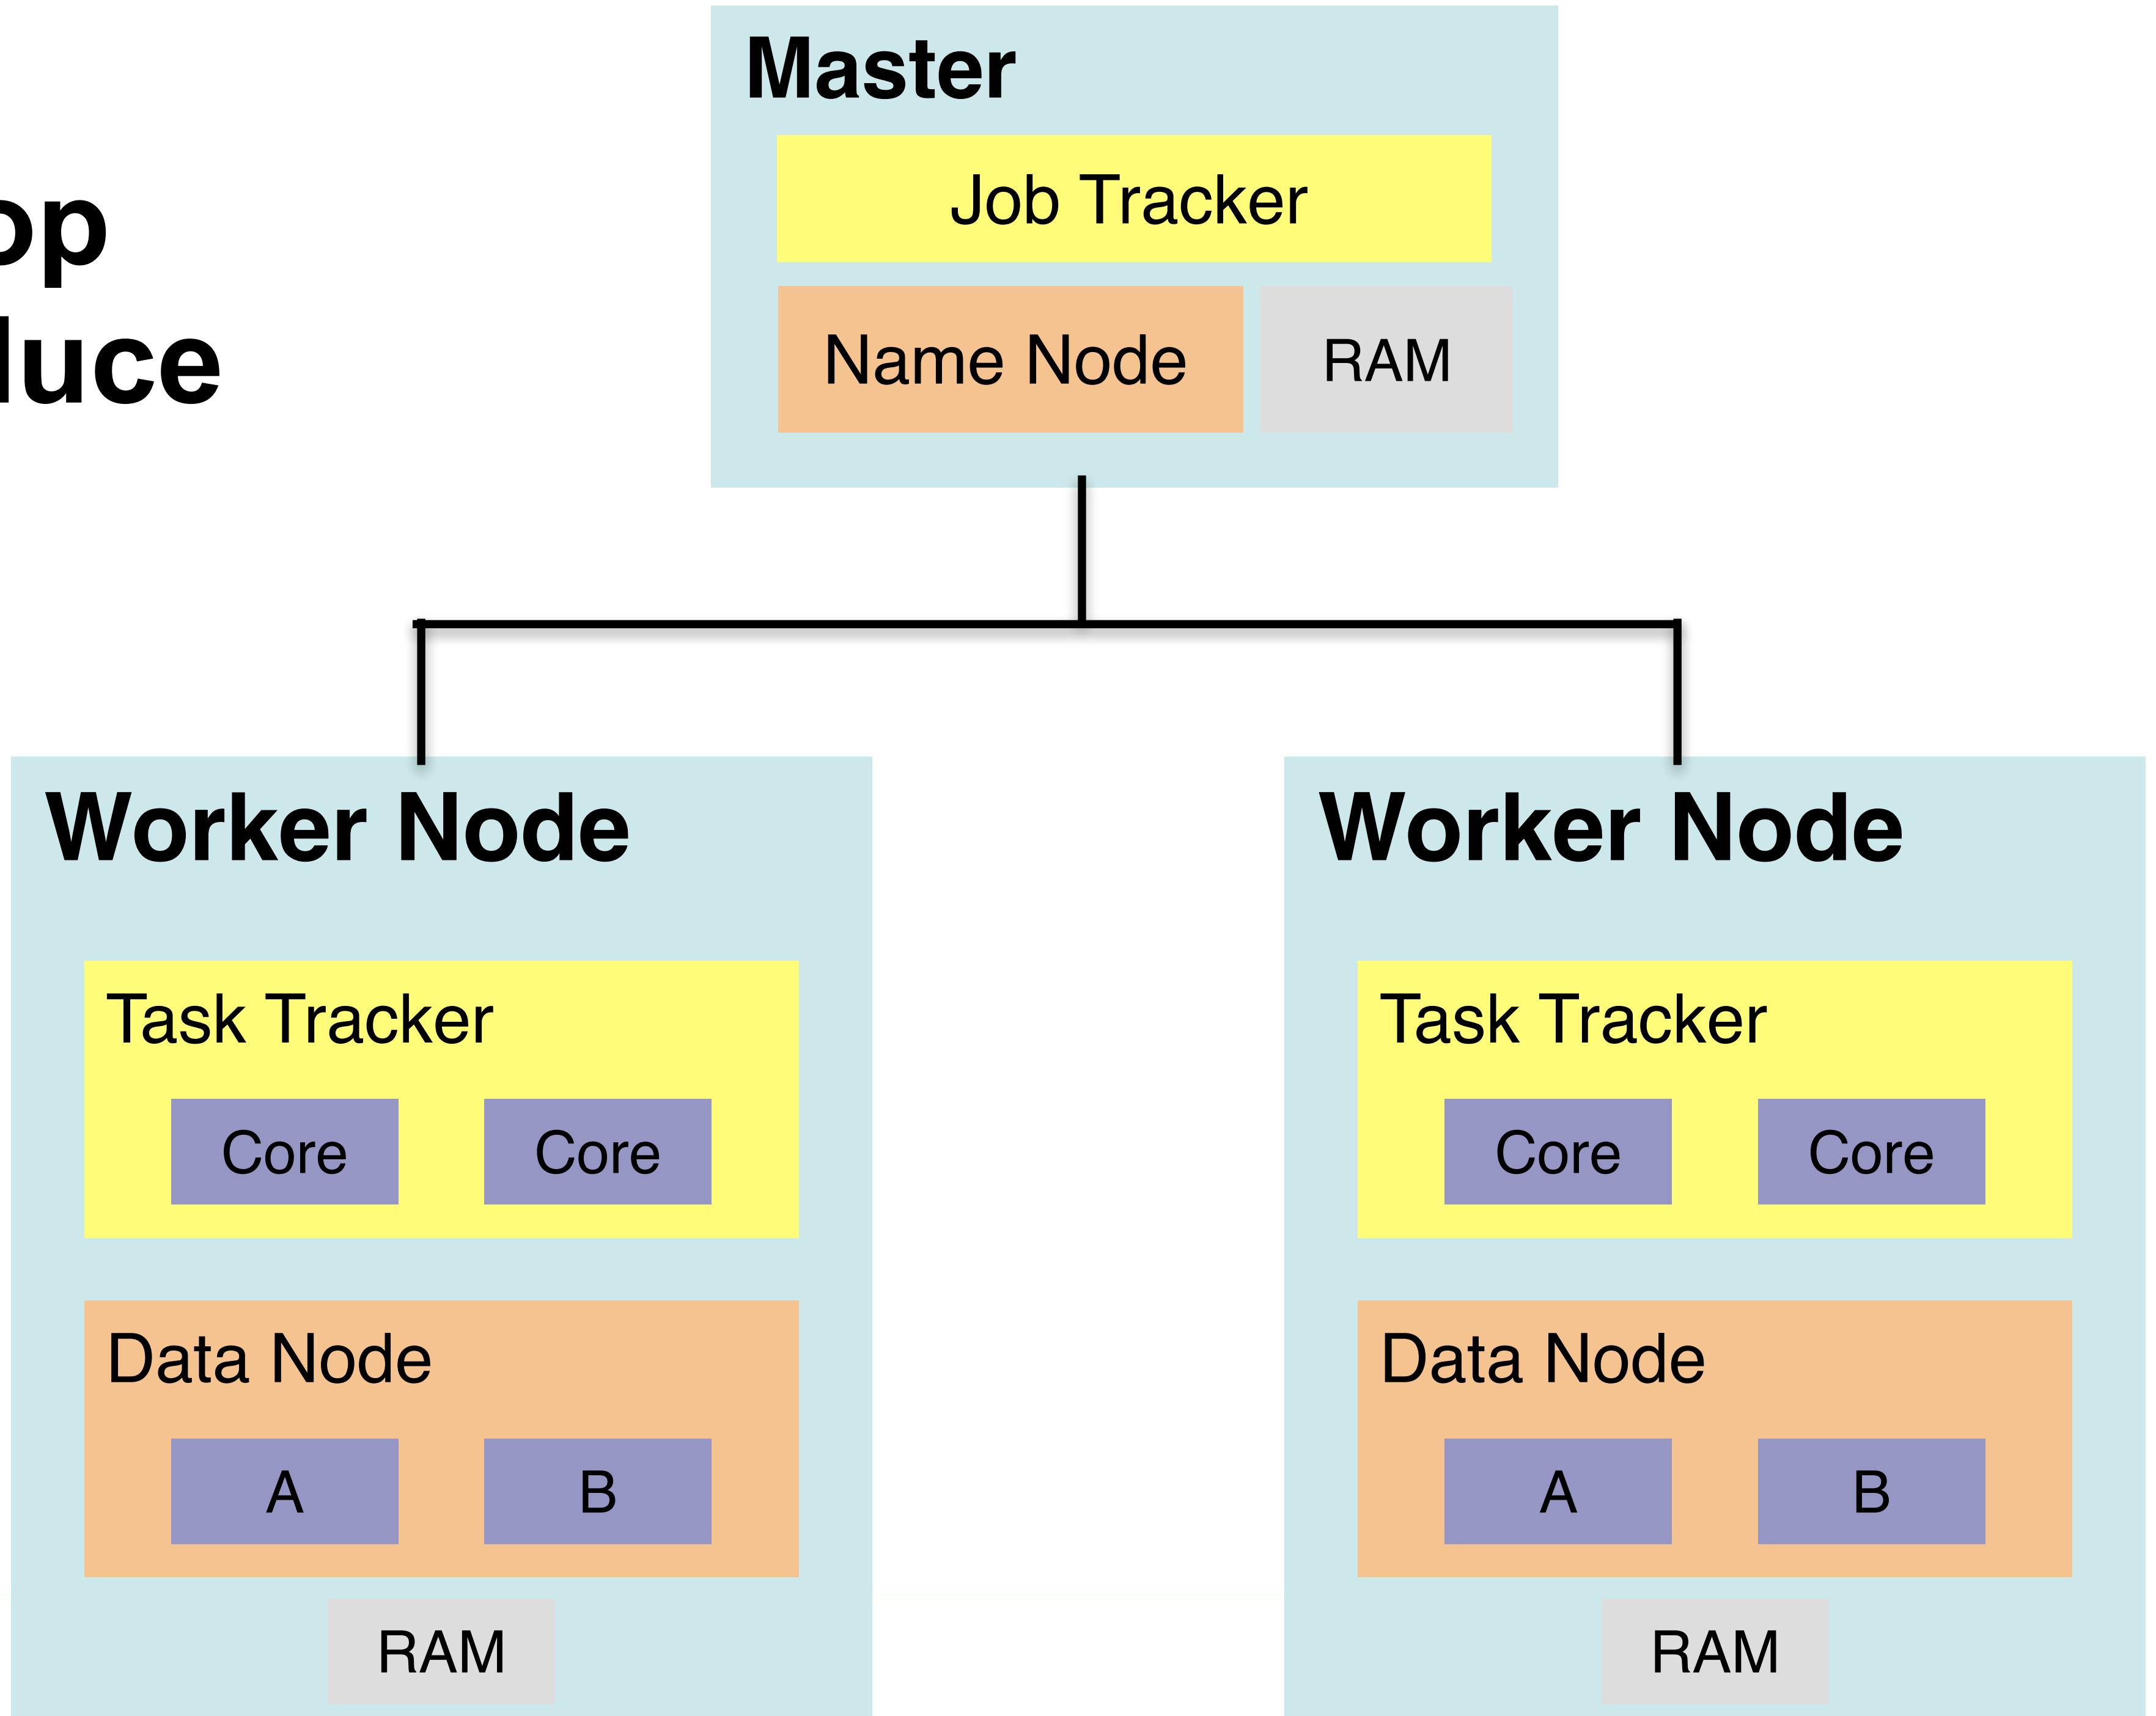
Spark on a Cluster



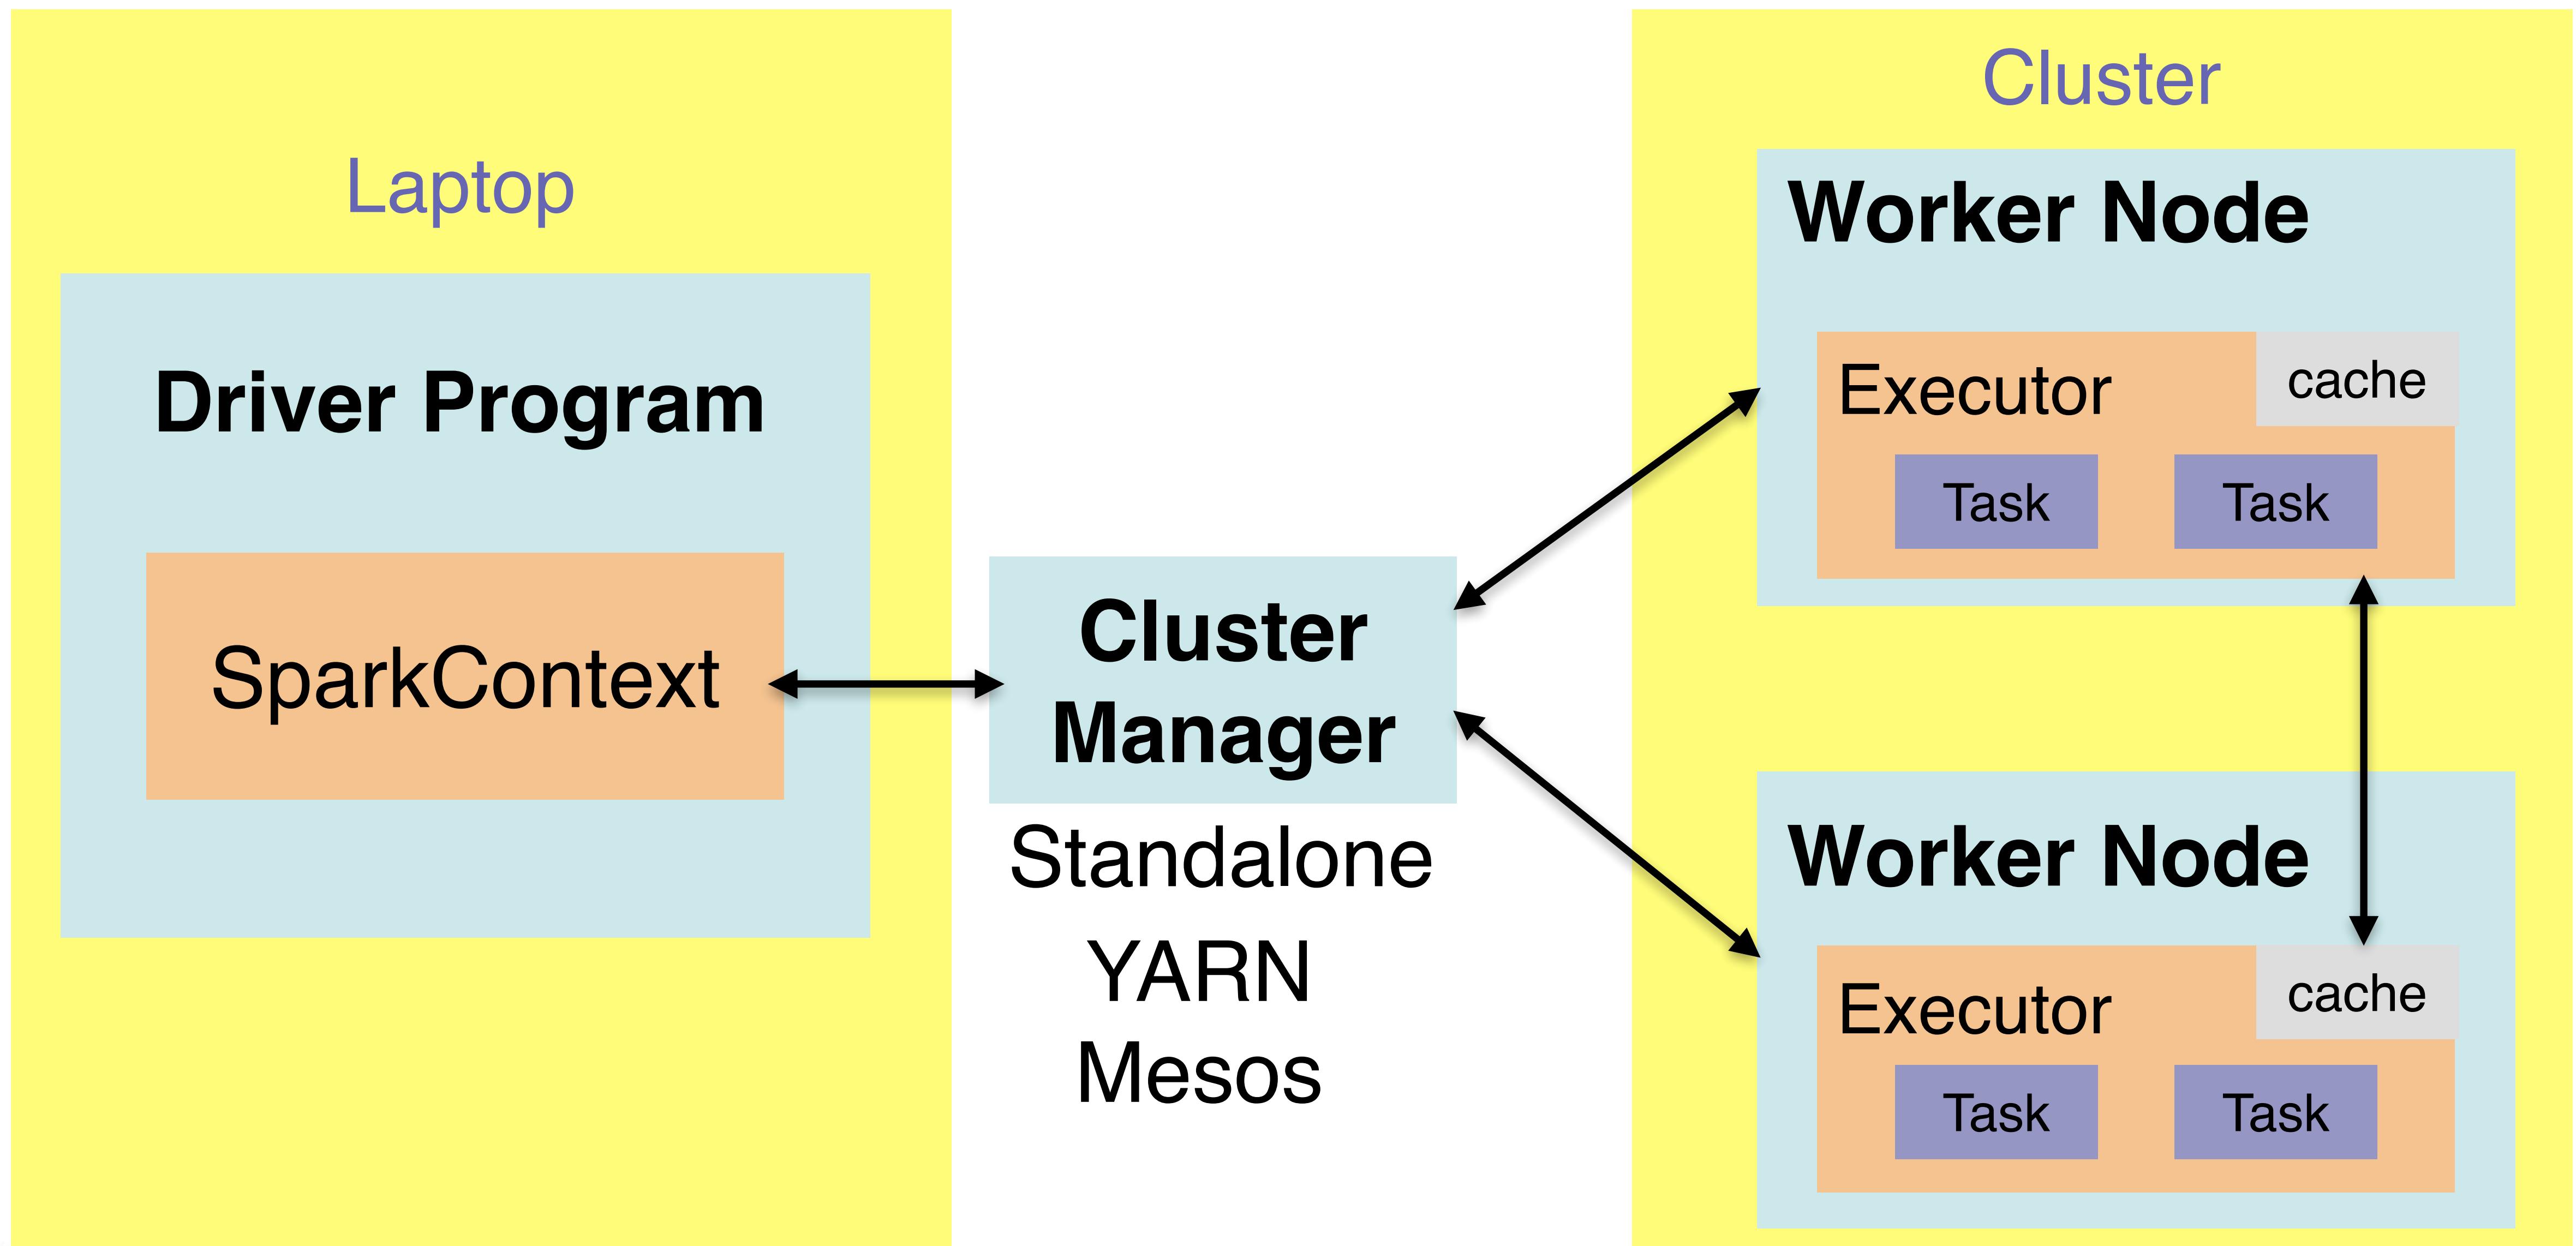
Terminology

Term	Meaning
Driver	<i>Process that contains the SparkContext</i>
Executor	<i>Process that executes one or more Spark tasks</i>
Master	<i>Process that manages applications across the cluster</i>
Worker	<i>Process that manages executors on a particular node</i>

Hadoop MapReduce



Spark Execution Context



Spark vs. Hadoop

- Spark only replaces MapReduce (**computation**)
- Still need a **data store**: HDFS, HBase, Hive, etc.
- Spark has a more **flexible/general** programming model
- Spark often faster for **iterative** computation

A Day in the Life of a Spark Application

1. Determine RDD lineage: construct DAG
2. Create execution plan for DAG: determine stages
3. Schedule and execute tasks



dummy.txt

```
file_rdd.map(line => line.split(""))
.map(split => (split(0), split(1).toInt))
.groupByKey()
.mapValues(iter => iter.reduce(_ + _)).collect()
```

jon	2
mary	3
anna	1
jon	1
jesse	3
mary	5

What's in a Task?

```
sc.textFile('file:///dummy.txt')
```



```
map(line => line.split(" "))
```



```
map(split => (split(0), split(1).toInt))
```



```
groupByKey()
```



```
mapValues(iter => iter.reduce(_ + _))
```



```
collect()
```

1. Create RDDs

```
sc.textFile('file:///dummy.txt')
```



```
map(line => line.split(" "))
```



```
map(split => (split(0), split(1).toInt))
```



```
groupByKey()
```



```
mapValues(iter => iter.reduce(_ + _))
```



```
collect()
```

```
HadoopRDD[ ]
```

```
MapPartitionsRDD[ ]
```



```
MapPartitionsRDD[ ]
```



```
MapPartitionsRDD[ ]
```



```
ShuffledRDD[ ]
```



```
MapPartitionsRDD[ ]
```



```
collect()
```

```
scala> file_rdd.partitions
res51: Array[org.apache.spark.Partition] = Array(org.apache.spark.rdd.HadoopPartition@a17, org.apache.spark.rdd.HadoopPartition@a18)

scala> file_rdd.partitions.size
res52: Int = 2

scala> file_rdd.map(line => line.split(" ")).toDebugString
res53: String =
(2) MapPartitionsRDD[65] at map at <console>:24 []
| MapPartitionsRDD[23] at textFile at <console>:21 []
| file:///Users/jonathandinu/spark-ds-applications/dummy.txt HadoopRDD[22] at textFile at <console>:21 []

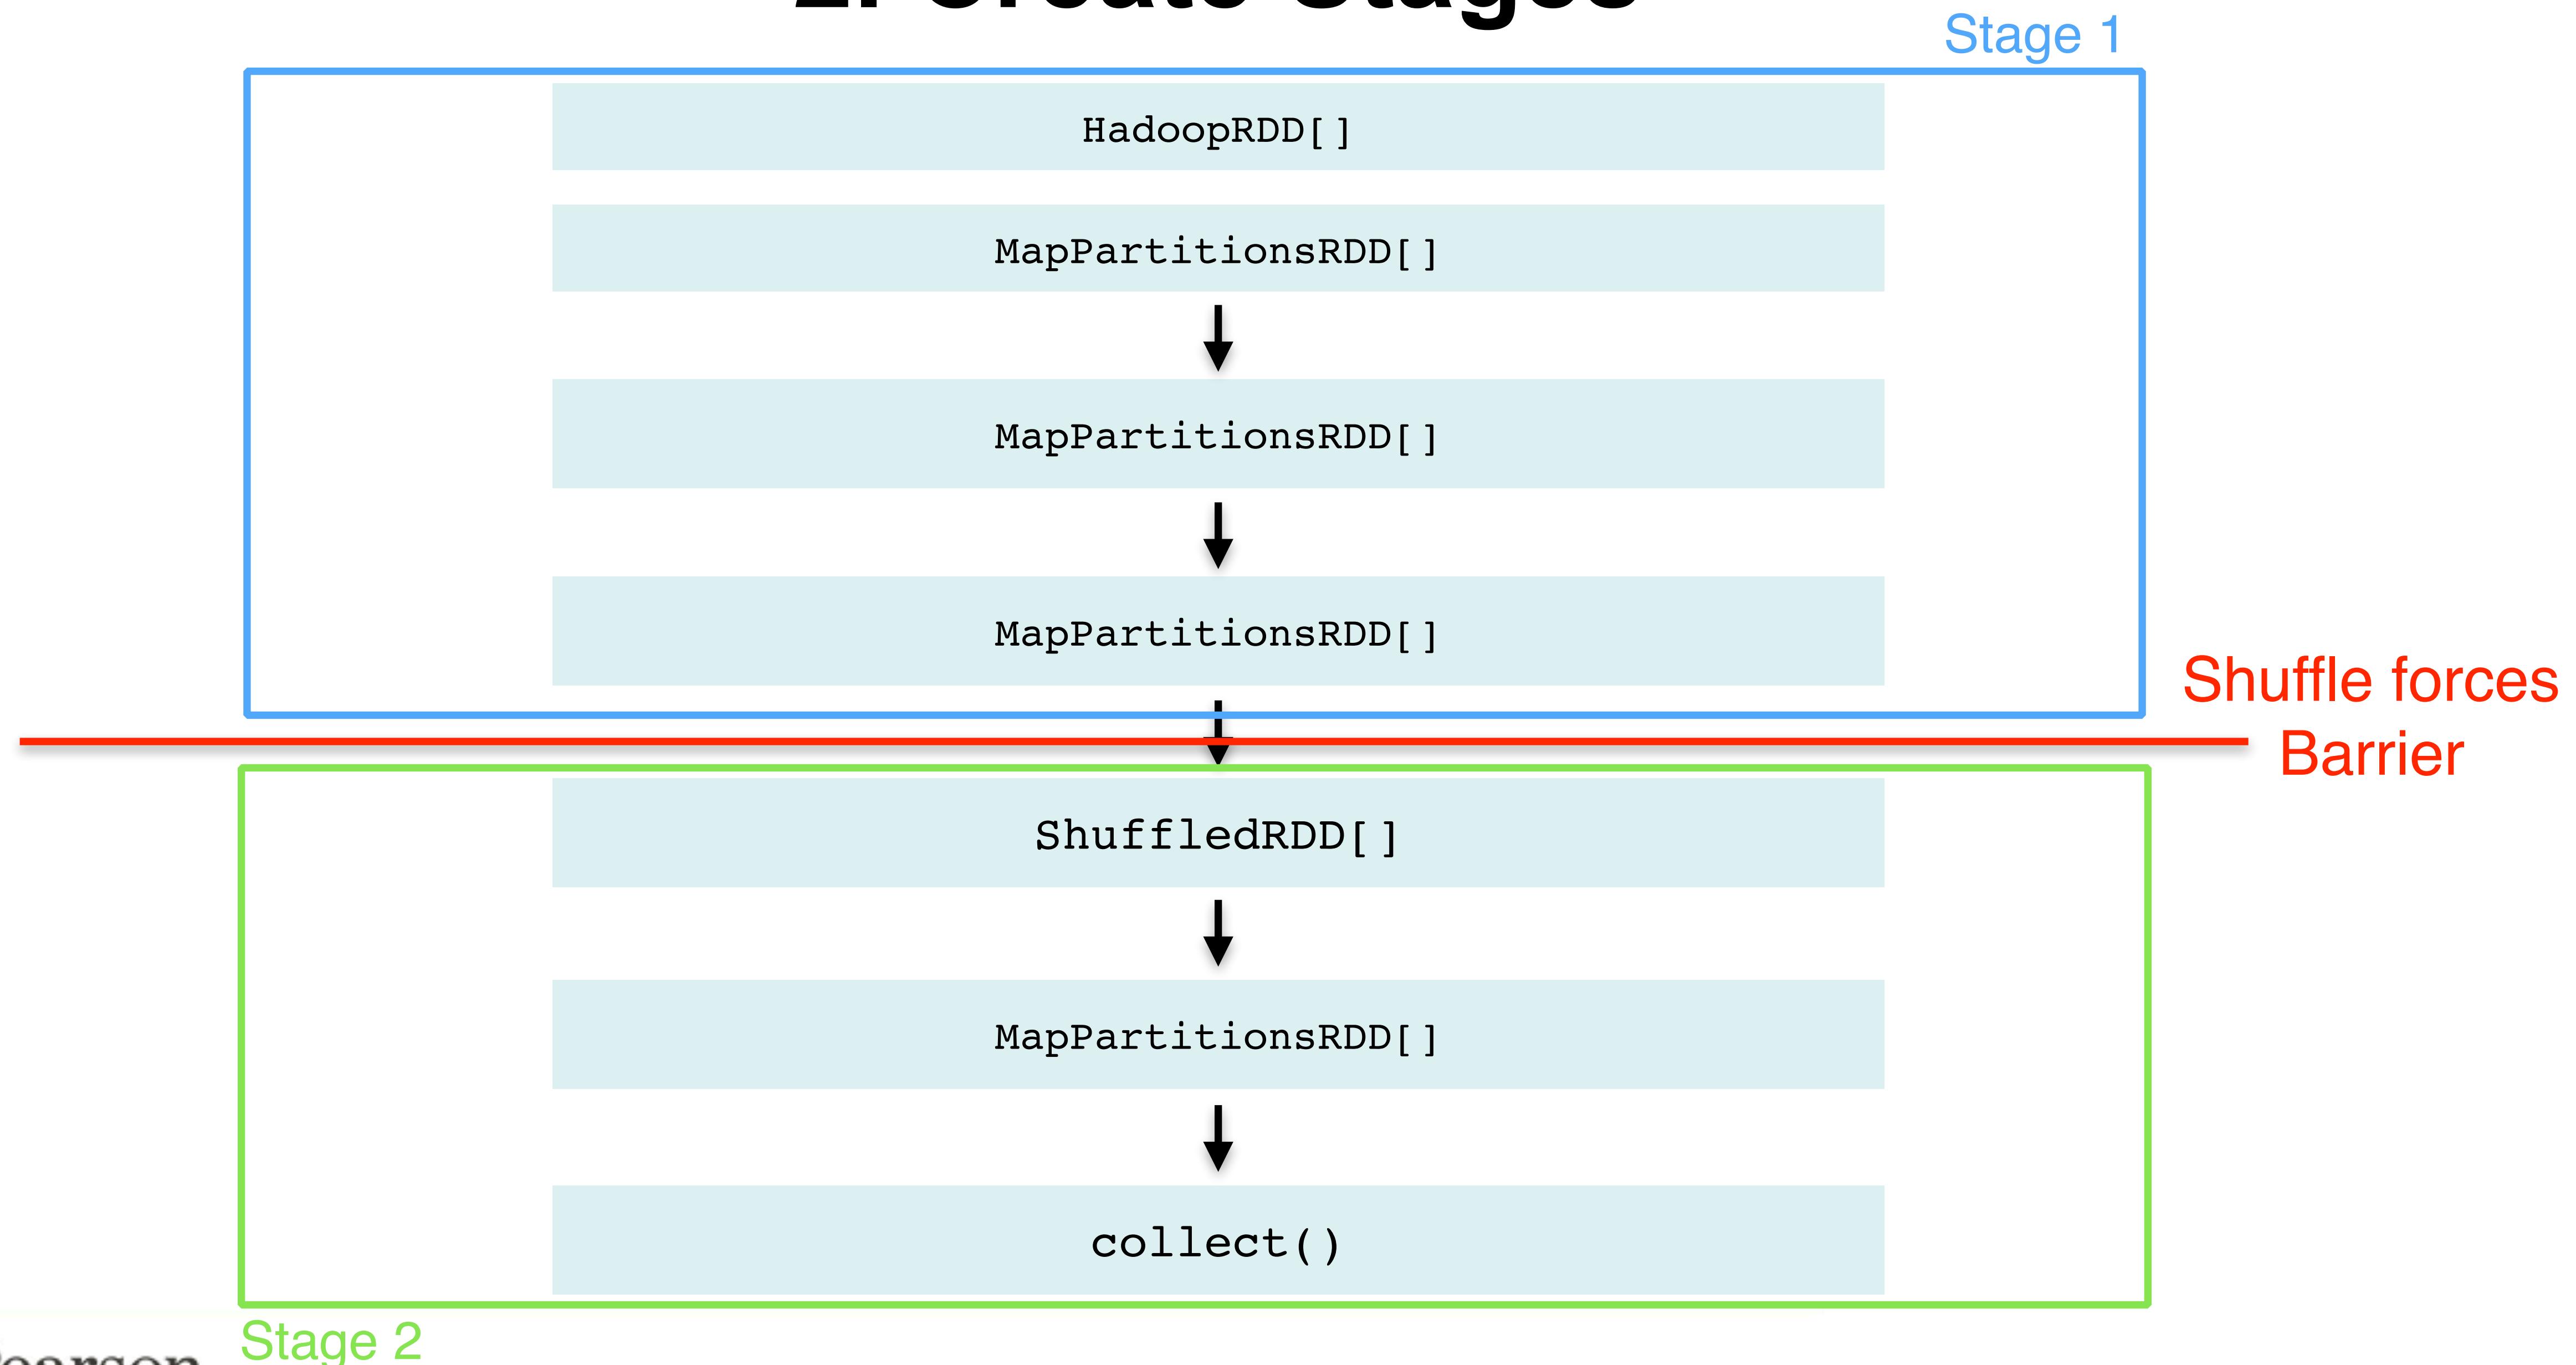
scala> file_rdd.map(line => line.split(" ")).map(split => (split(0), split(1).toInt)).toDebugString
res54: String =
(2) MapPartitionsRDD[67] at map at <console>:24 []
| MapPartitionsRDD[66] at map at <console>:24 []
| MapPartitionsRDD[23] at textFile at <console>:21 []
| file:///Users/jonathandinu/spark-ds-applications/dummy.txt HadoopRDD[22] at textFile at <console>:21 []

scala> file_rdd.map(line => line.split(" ")).map(split => (split(0), split(1).toInt)).groupByKey().toDebugString
res55: String =
(2) ShuffledRDD[70] at groupByKey at <console>:24 []
+- (2) MapPartitionsRDD[69] at map at <console>:24 []
| MapPartitionsRDD[68] at map at <console>:24 []
| MapPartitionsRDD[23] at textFile at <console>:21 []
| file:///Users/jonathandinu/spark-ds-applications/dummy.txt HadoopRDD[22] at textFile at <console>:21 []

scala> file_rdd.map(line => line.split(" ")).map(split => (split(0), split(1).toInt)).groupByKey().mapValues(iter => iter.reduce(_ + _)).toDebugString
res56: String =
(2) MapPartitionsRDD[74] at mapValues at <console>:26 []
| ShuffledRDD[73] at groupByKey at <console>:26 []
+- (2) MapPartitionsRDD[72] at map at <console>:26 []
| MapPartitionsRDD[71] at map at <console>:26 []
| MapPartitionsRDD[23] at textFile at <console>:21 []
| file:///Users/jonathandinu/spark-ds-applications/dummy.txt HadoopRDD[22] at textFile at <console>:21 []

scala> file_rdd.map(line => line.split(" ")).map(split => (split(0), split(1).toInt)).groupByKey().mapValues(iter => iter.reduce(_ + _)).collect()
res57: Array[(String, Int)] = Array((anna,1), (jesse,3), (jon,3), (mary,8))
```

2. Create Stages



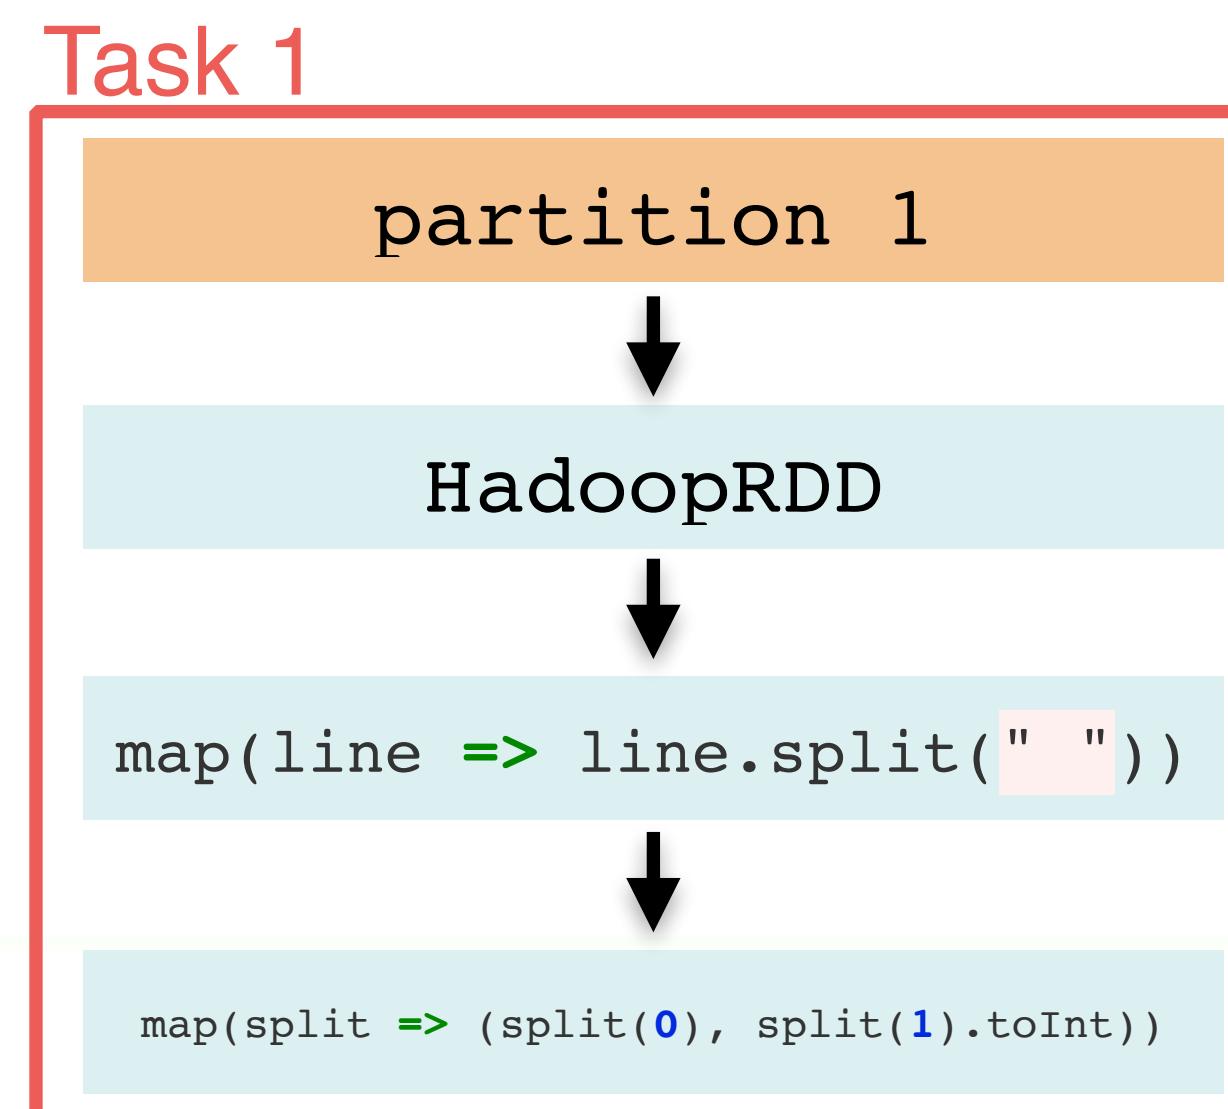
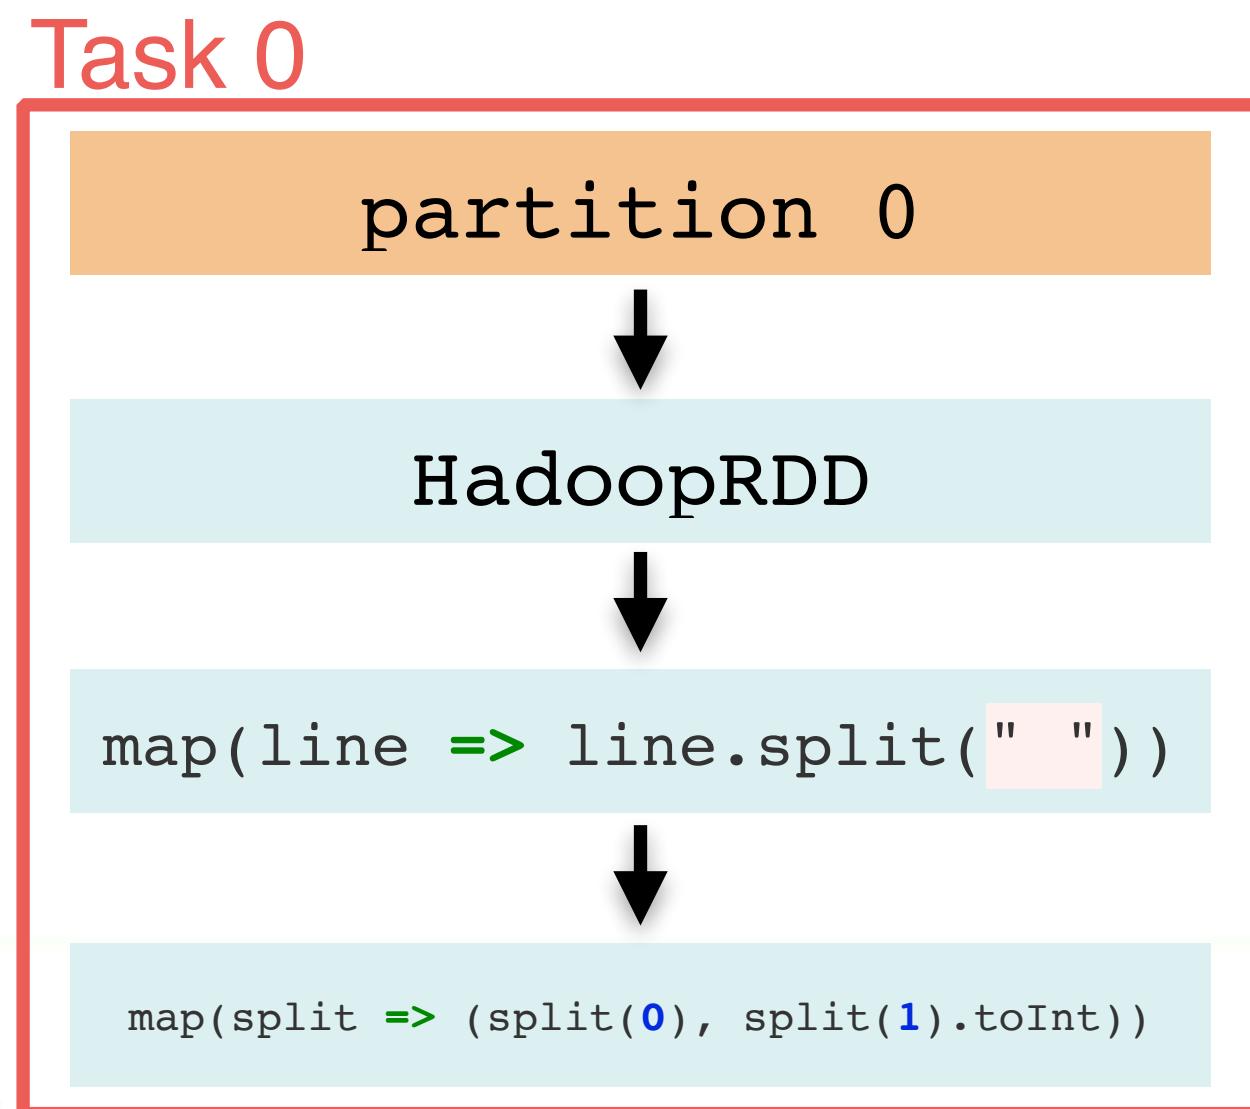
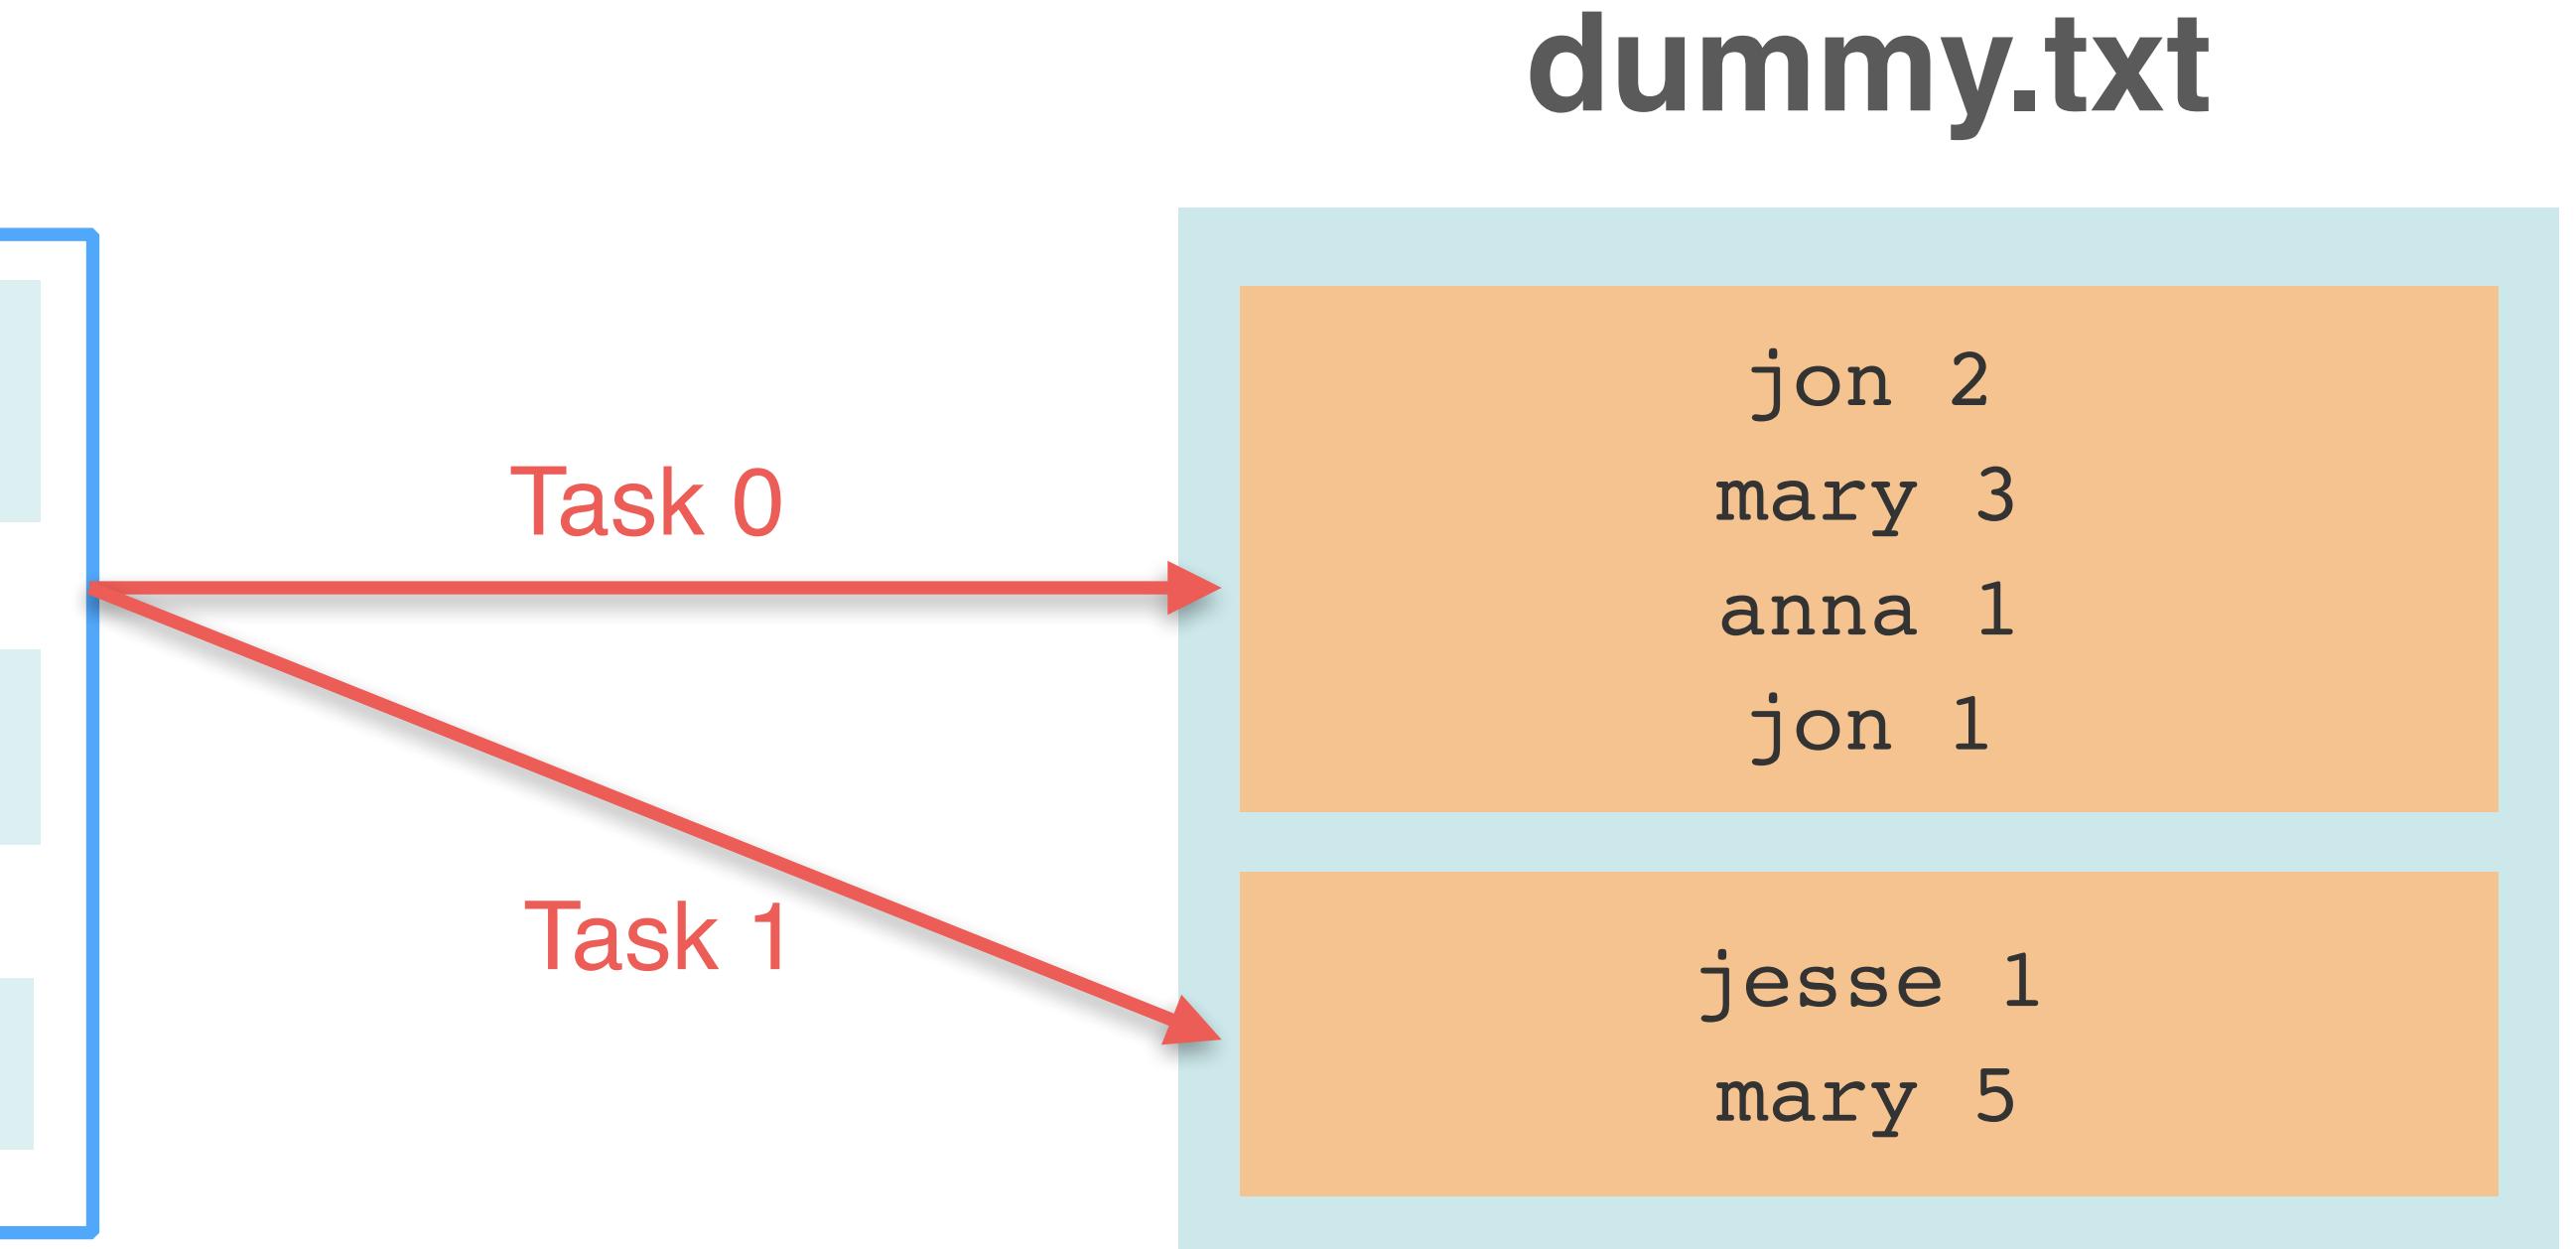
3. Schedule Tasks

1. Split each **stage** into **tasks** to execute
2. Task is simply a **partition** (data) and **computation** (lambda)
3. Execute all **tasks** in a stage before **continuing**

3. Schedule Tasks

Stage 1

```
sc.textFile('file:///dummy.txt')  
      ↓  
map(line => line.split(" " ))  
      ↓  
map(split => (split(0), split(1).toInt))
```



Scheduler

- RDD and partition DAG as **input**
- Tasks (within stages) to execute as **output**

Roles

- Builds stages to execute
- Submit stages to **Cluster Manager**
- **Resubmit** failed stages when output is lost

A Day in the Life of a Spark Application

