

# What is Ray

**Ray provides a simple, universal API for building distributed applications**

# What is Ray

**Ray provides a simple, universal API for building distributed **applications****

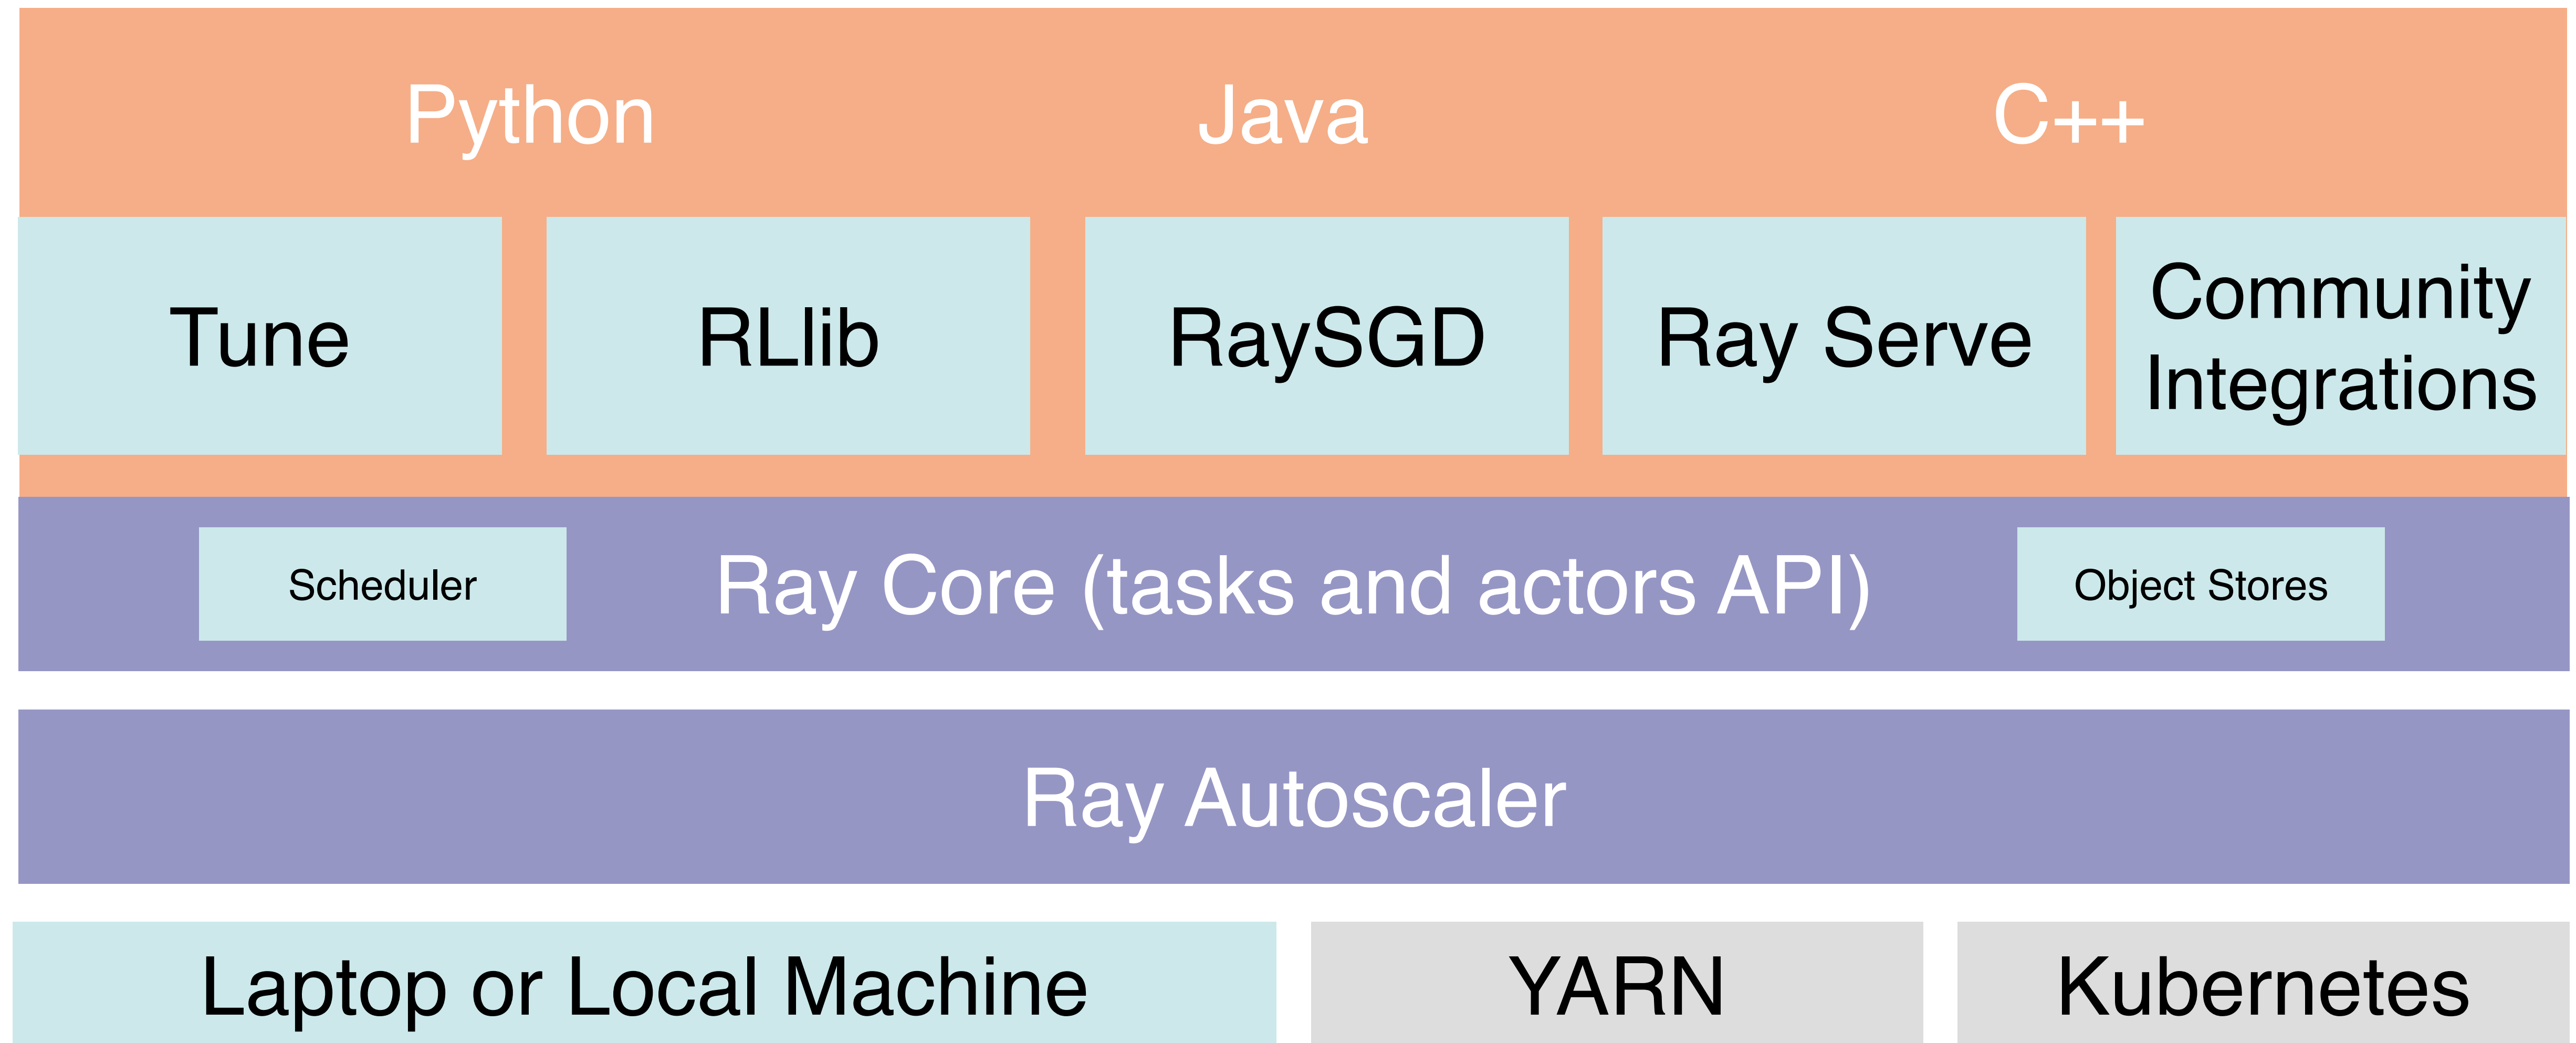
# What is Ray

- Primitives for running fault-tolerant distributed applications
- Parallelize single machine code with little to zero code changes
- Large ecosystem of applications/libraries/tools on top of Ray Core

# What is Ray

multiprocess **meets** **RPC**...

# Ray Layer Cake



# The Ray Programming Model

## Tasks

Remote Functions

Similar to Spark closures

## Actors

Remote Classes/Objects

Maintain internal state

# The Ray Programming Model

## Tasks

```
import ray
ray.init()

@ray.remote
def f(x):
    return x * x

futures = [f.remote(i) for i in range(4)]
print(ray.get(futures)) # [0, 1, 4, 9]
```

## Actors

```
import ray
ray.init() # Only call this once.
```

```
@ray.remote
class Counter(object):
    def __init__(self):
        self.n = 0
```

```
    def increment(self):
        self.n += 1
```

```
    def read(self):
        return self.n
```

```
counters = [Counter.remote() for i in range(4)]
[c.increment.remote() for c in counters]
futures = [c.read.remote() for c in counters]
print(ray.get(futures)) # [1, 1, 1, 1]
```

# The Ray Programming Model

## Tasks

- Fine-grained load balancing
- Support for object locality
- High overhead for small updates
- Efficient failure handling

## Actors

- Coarse-grained load balancing
- Poor locality support
- Low overhead for small updates
- Overhead from checkpointing



## Spark

- Coarse grained
- SQL like abstractions (tables)
- Synchronous
- Functional

## Ray

- Fine grained
- Actor based abstractions (agents)
- Asynchronous
- Object Oriented

## Spark

- Data processing and transformation
- Data heavy workloads
- Iterative MapReduce
- Ex: Pagerank

## Ray

- Distributed asynchronous algorithms
- Compute heavy workloads
- Ex: Reinforcement learning

# Spark

```
text_file = sc.textFile("hdfs://...")

counts = text_file.flatMap(lambda line: line.split(" ")) \
                  .map(lambda word: (word, 1)) \
                  .reduceByKey(lambda a, b: a + b)

counts.saveAsTextFile("hdfs://...")
```

# Ray

```
@ray.remote
def f(x):
    return x * x

futures = [f.remote(i) for i in range(4)]
print(ray.get(futures)) # [0, 1, 4, 9]
```

# Spark

```
def inside(p):  
    x, y = random.random(), random.random()  
    return x*x + y*y < 1  
  
count = sc.parallelize(range(0, NUM_SAMPLES)) \  
    .filter(inside).count()  
  
print("Pi is roughly %f" % (4.0 * count / NUM_SAMPLES))
```

# Ray

```
@ray.remote  
def f(x):  
    return x * x  
  
futures = [f.remote(i) for i in range(4)]  
print(ray.get(futures)) # [0, 1, 4, 9]
```

# Spark

```
# Counts people by age
countsByAge = df.groupBy("age").count()

# Saves countsByAge to S3 in the JSON format.
countsByAge.write.format("json").save("s3a://...")
```

# Ray

```
@ray.remote
def f(x):
    return x * x

futures = [f.remote(i) for i in range(4)]
print(ray.get(futures)) # [0, 1, 4, 9]
```

# Spark

```
# Set parameters for the algorithm.  
# Here, we limit the number of iterations to 10.  
lr = LogisticRegression(maxIter=10)  
  
# Fit the model to the data.  
model = lr.fit(df)  
  
# Given a dataset, predict each point's label.  
model.transform(df).show()
```

# Ray

```
@ray.remote  
class DataWorker(object):  
    def __init__(self):  
        self.model = ConvNet()  
        self.data_iterator = iter(get_data_loader()[0])  
  
    def compute_gradients(self, weights):  
        self.model.set_weights(weights)  
        data, target = next(self.data_iterator)  
  
        self.model.zero_grad()  
        output = self.model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        return self.model.get_gradients()
```


# Differences from Spark

- No built in primitives for partitioning data across cluster.
- Stateful computation with Actors.
- Asynchronous execution.

# Agent Based Modeling

Dead cell

Live cell




1. Any live cell with fewer than two live neighbors dies, as if by underpopulation.
2. Any live cell with two or three live neighbors lives on to the next generation.
3. Any live cell with more than three live neighbors dies, as if by overpopulation.
4. Any dead cell with exactly three live neighbors becomes a live cell, as if by reproduction.



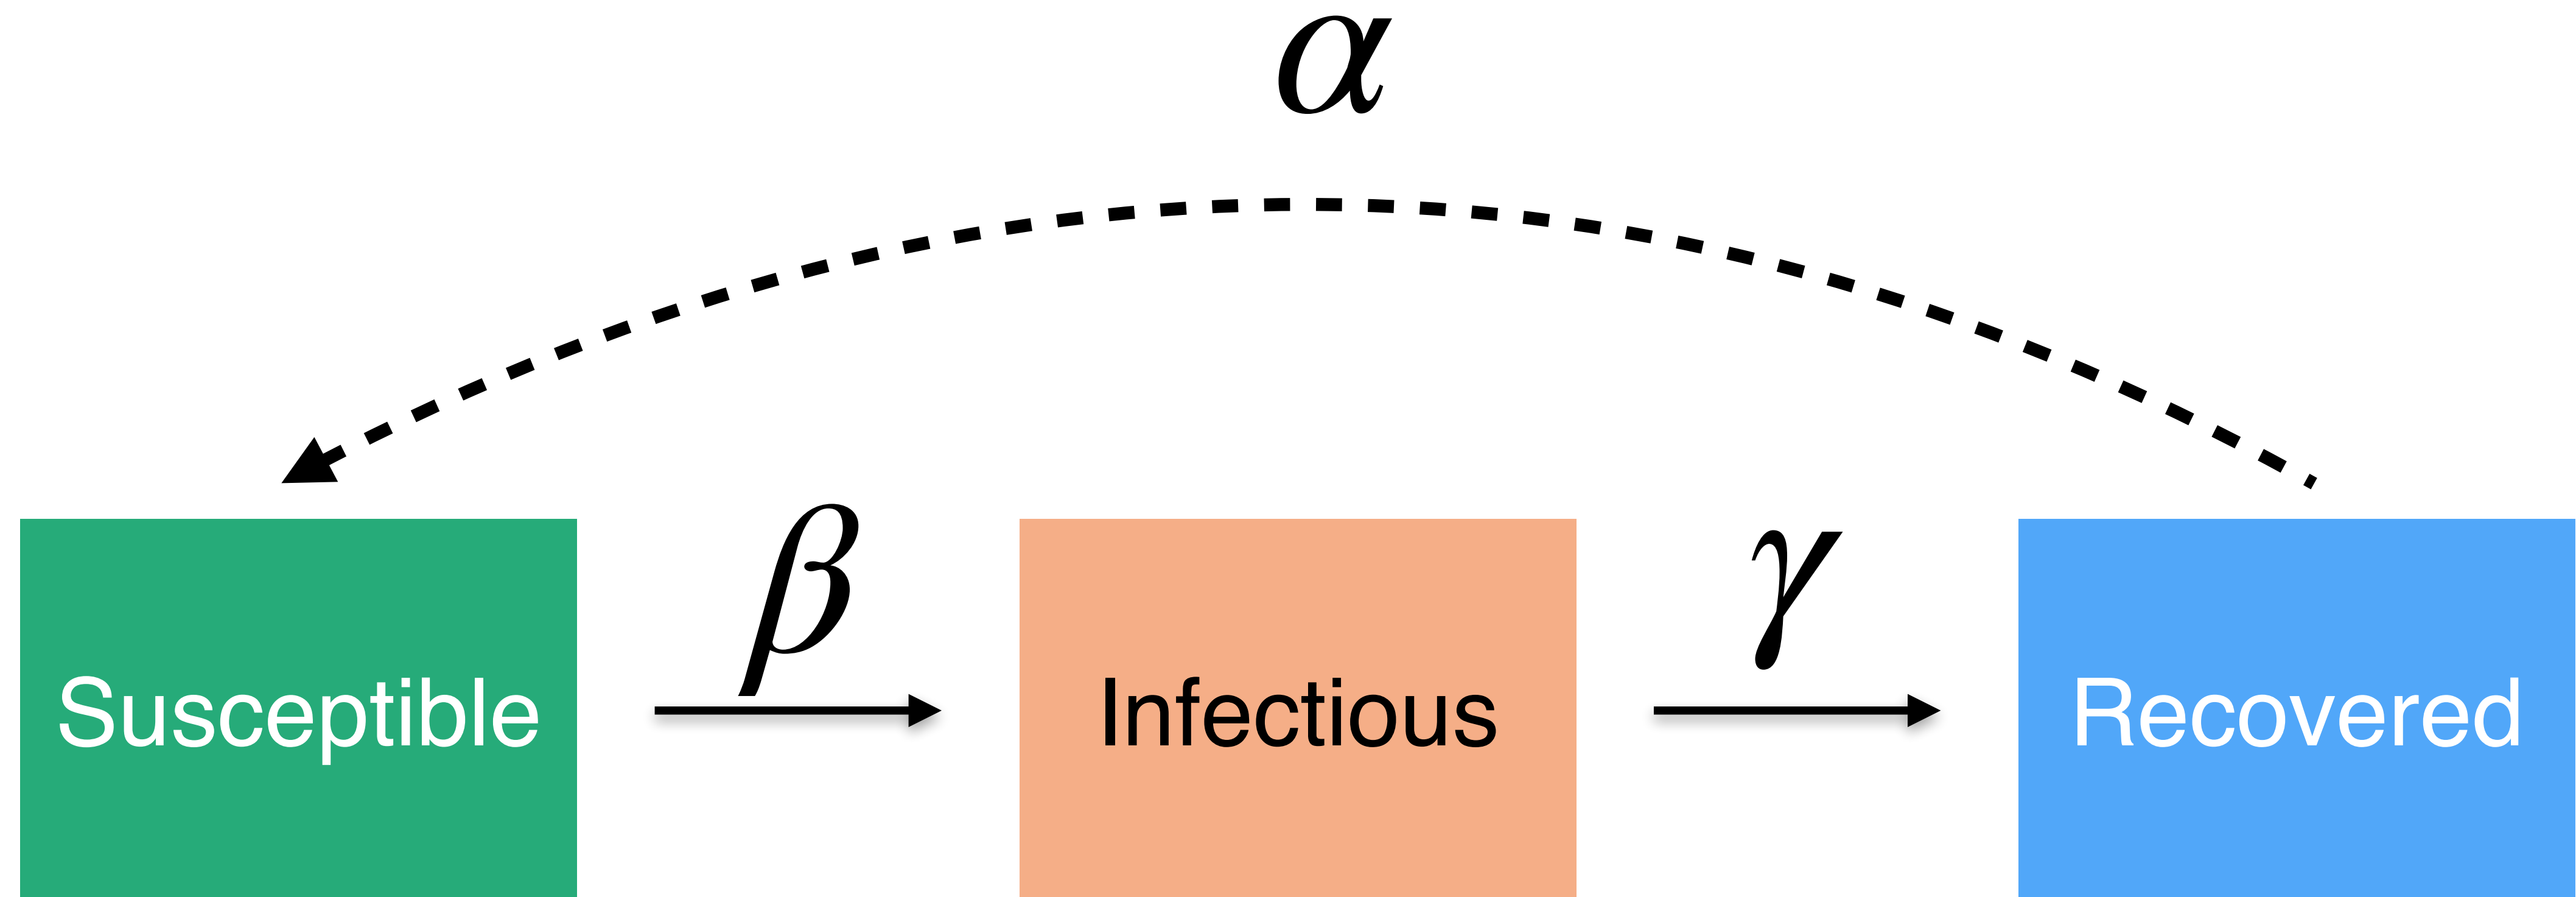
# Schelling's Model of Segregation

- Agents belongs to one of two groups (0 or 1).
- The agents live in a two-dimensional grid world.
- Agents with at least three neighbors in the same group are happy.
- Agent keeps moving one step to new locations until happy.

# Schelling's Model of Segregation



# SIR Epidemic model



$S(t)$ : susceptible individuals who have not yet been infected at time  $t$

$I(t)$ : number of infectious individuals at time  $t$

$R(t)$ : number of individuals who have recovered (and are immune) at time  $t$



# SIR Epidemic model

- Agents belongs to one of three groups (S, I, or R).
- Agents only interact with other agents in their neighborhood.
- An infected agent passes disease to a susceptible one with probability  $\beta$
- An infected agent recovers with probability  $\gamma$