



UPPSALA  
UNIVERSITET

# Lecture-4: Distributed Computing Infrastructures (DCI)

## Part-I

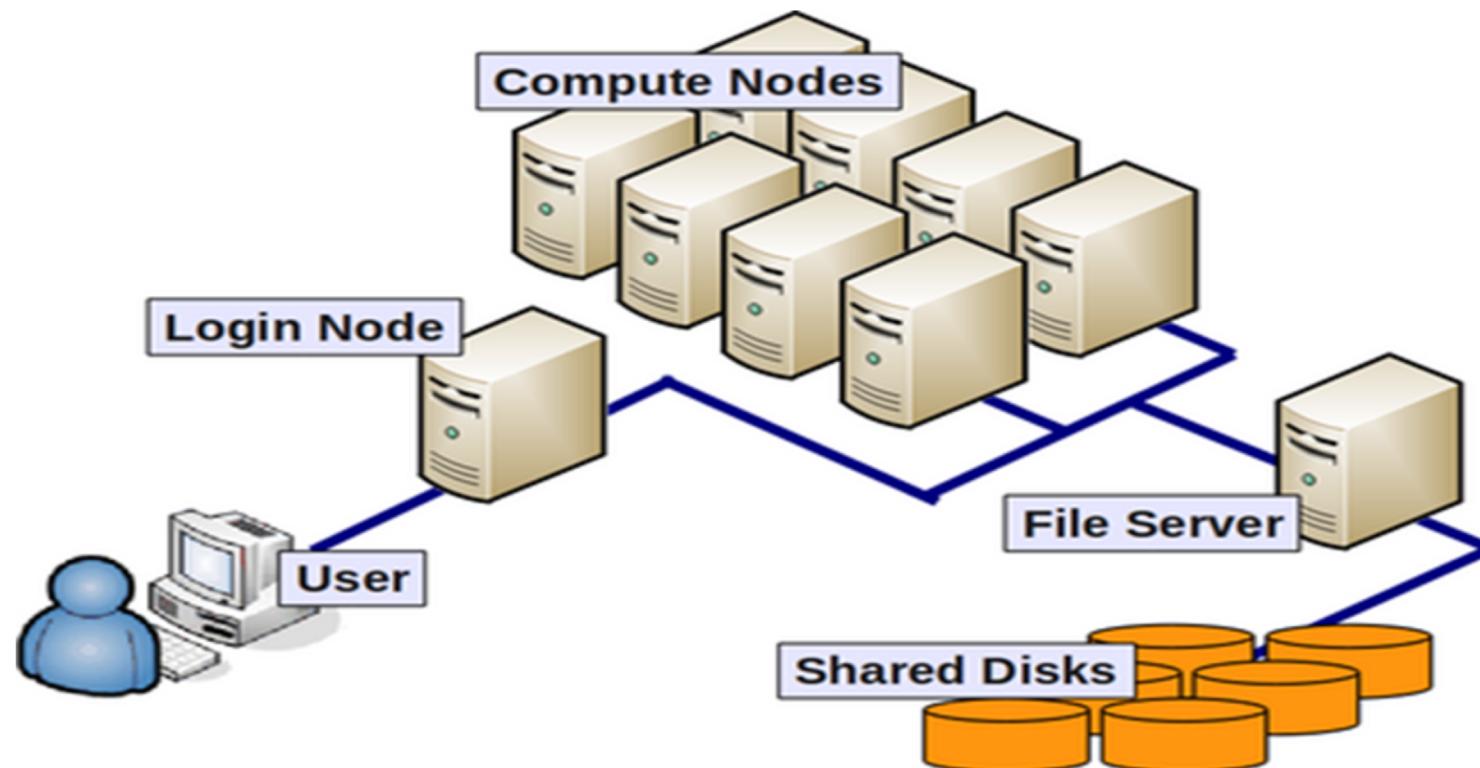
Presenter: Salman Toor  
[Salman.Toor@it.uu.se](mailto:Salman.Toor@it.uu.se)

# Distributed computing e-infrastructures (DCI)

- Broad term used for dedicated and public resources
  - Dedicated resources (clusters and specialized resources)
  - Public resources (public and community clouds)
- Offer compute, storage and network services

# Cluster computing

Dedicated setup designed according to the local needs





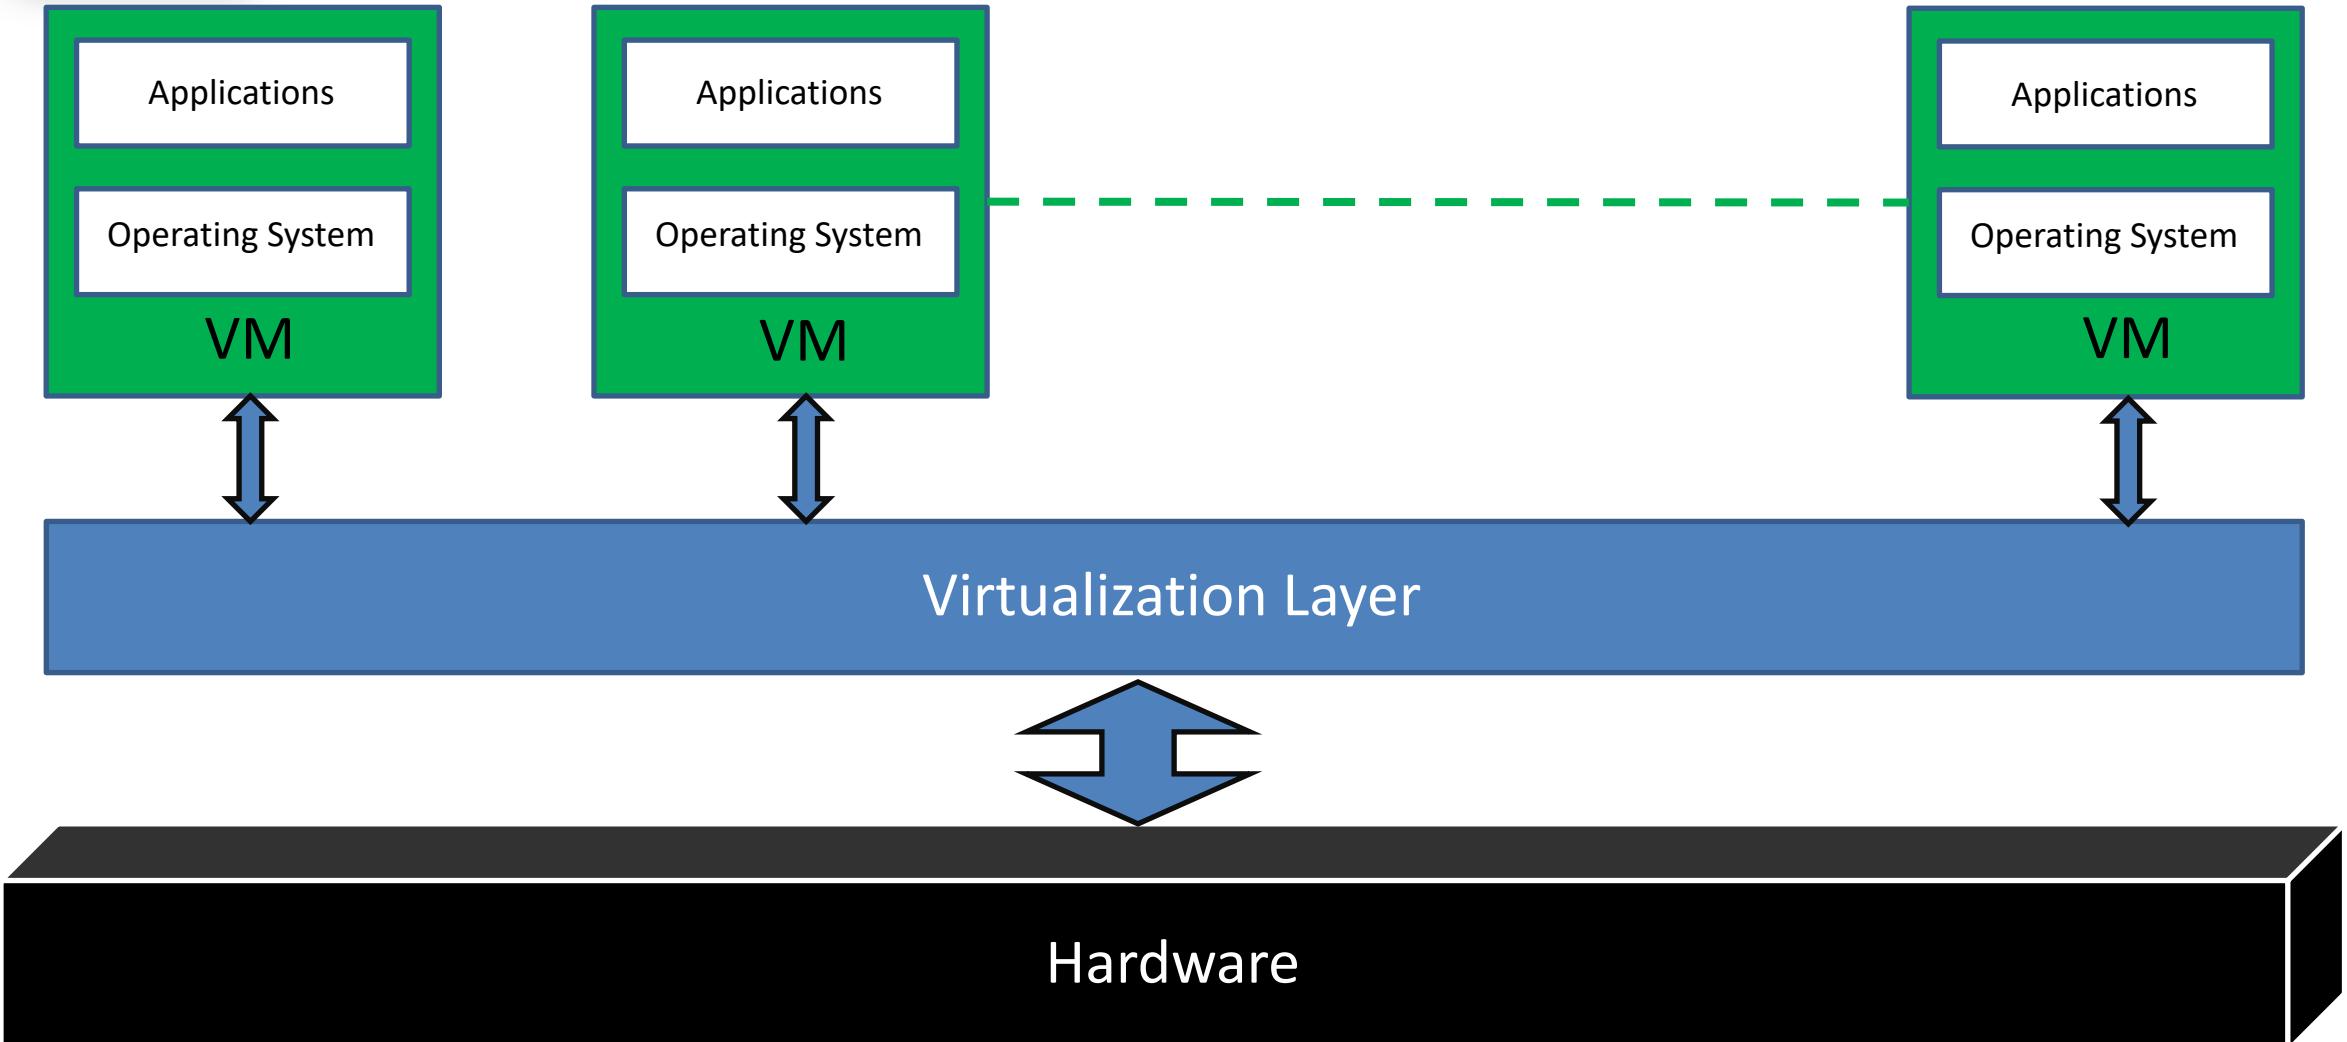
UPPSALA  
UNIVERSITET

# Cloud computing

Large-scale computing both for regular and specialized applications

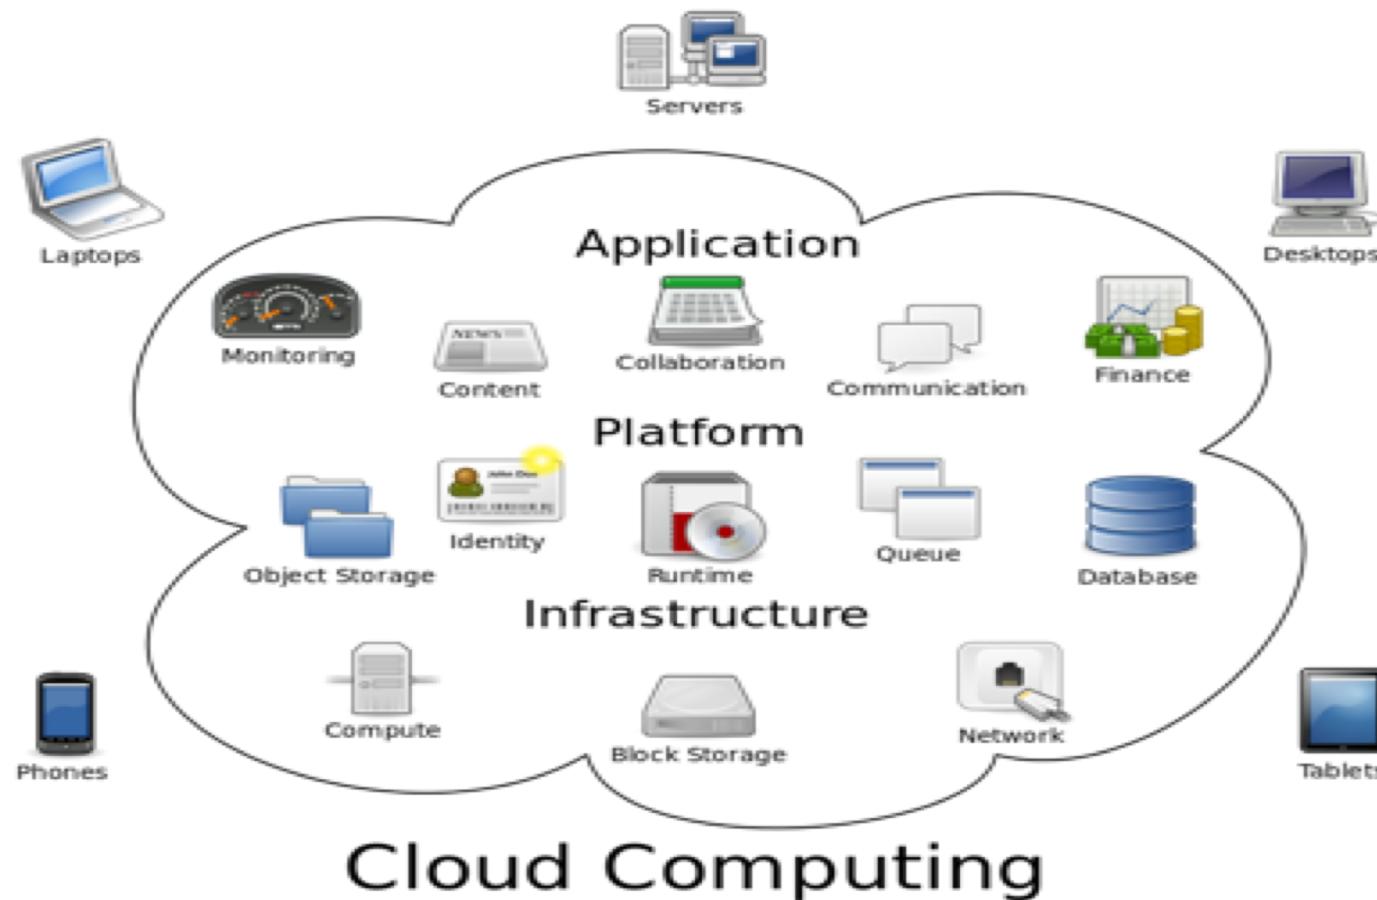


# Basic design of e-infrastructures





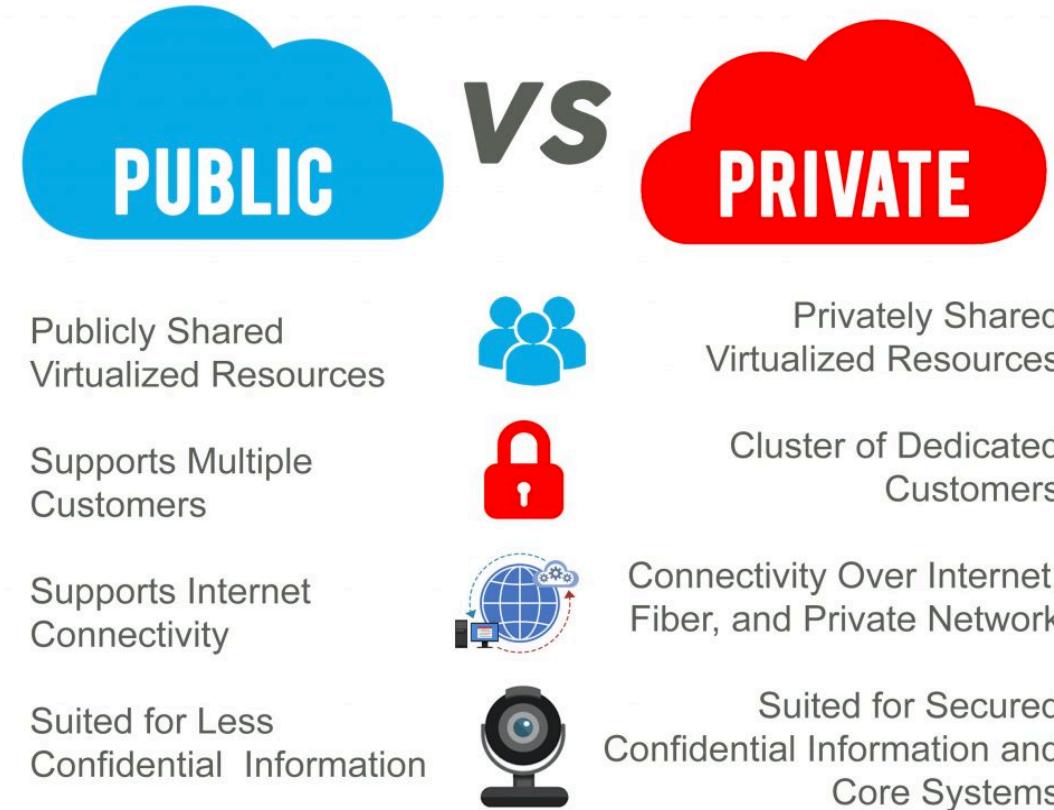
# What is Cloud Computing?



IT-infrastructure and services  
“obscured by a a cloud”

# Deployment Models

- Public
- Private
- Community
- Hybrid



# Features of current e-infrastructures

- On-demand availability of resources
- Pay-as-you-go model
- A platform with a range of services and tools
- Service level agreements (SLAs)

# Cloud computing

- A well-defined economic model
- Complete isolation, direct access and full control of allocated resources
- On demand resource allocation, No job queues!
- Loose coupling with the underlying resources
- “Standard” interface to interact with the cloud resources
- Orchestration of scalable services
- Minimal interaction with service providers

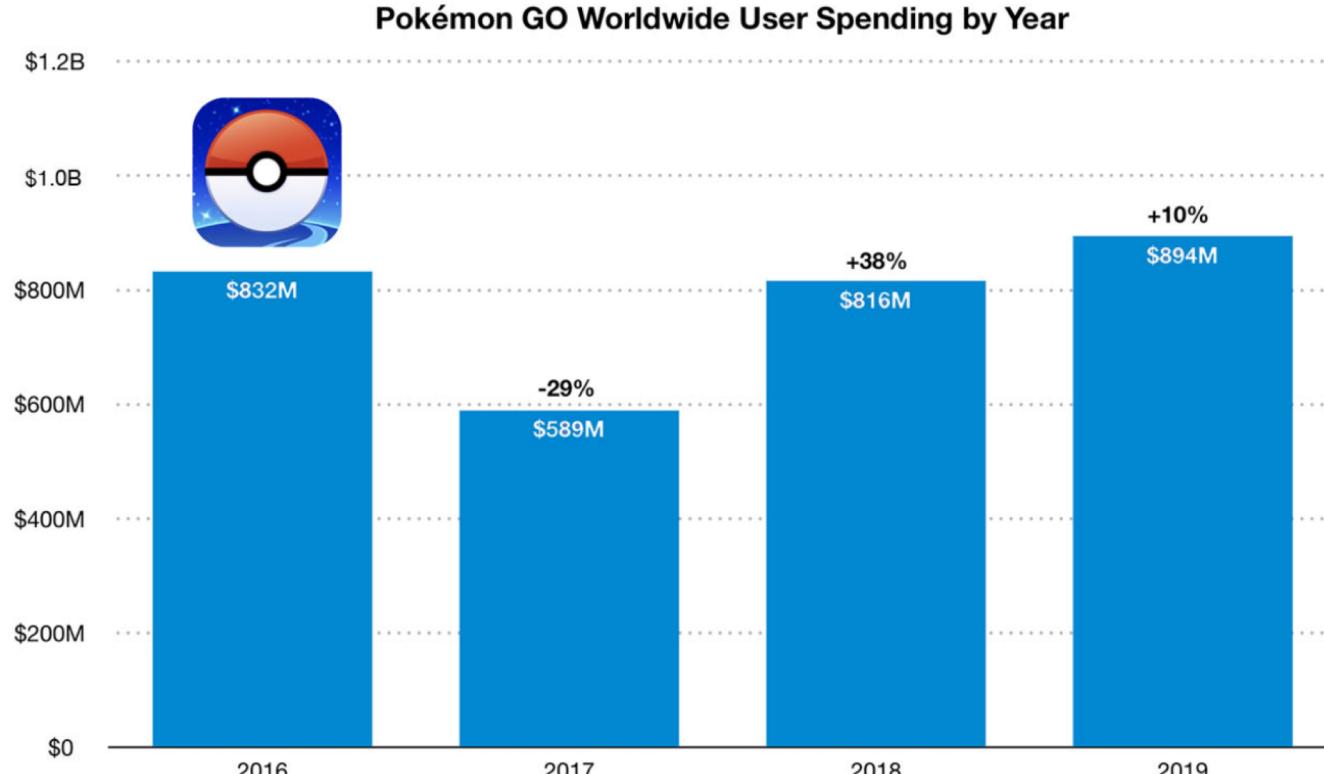
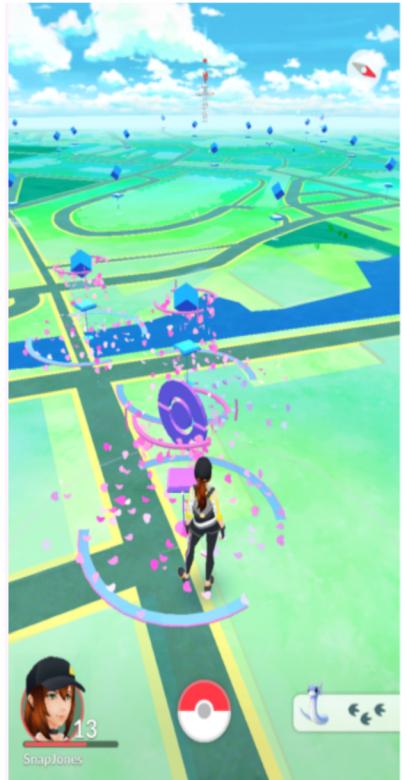
Are legacy applications portable to clouds?

**ANSWER: Yes**



UPPSALA  
UNIVERSITET

# Small companies



Pokémon Go

# Medium and large organizations

- Dropbox
- Google Apps (Gmail, Calendar, Drive ...)
- Netflix
- Uber Cab Service
- Spotify
- etc.

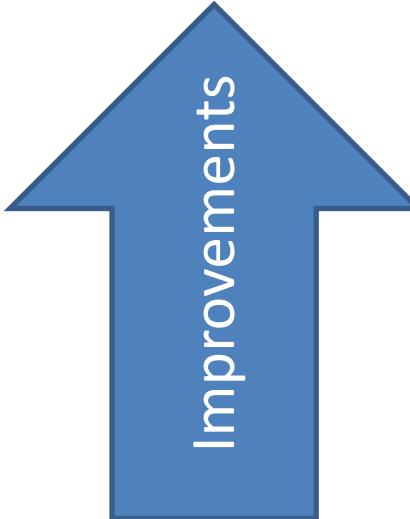


# Progress in last two decades

Cloud systems (Service-oriented, Agile software stack)

Level of abstraction

- Responsibilities
- Roles
- Business model
- ...



Improvements

Technological advancements

- Virtualization
- Communication tools
- Software designs
- ...

Dedicated clusters (Monolithic design)



UPPSALA  
UNIVERSITET

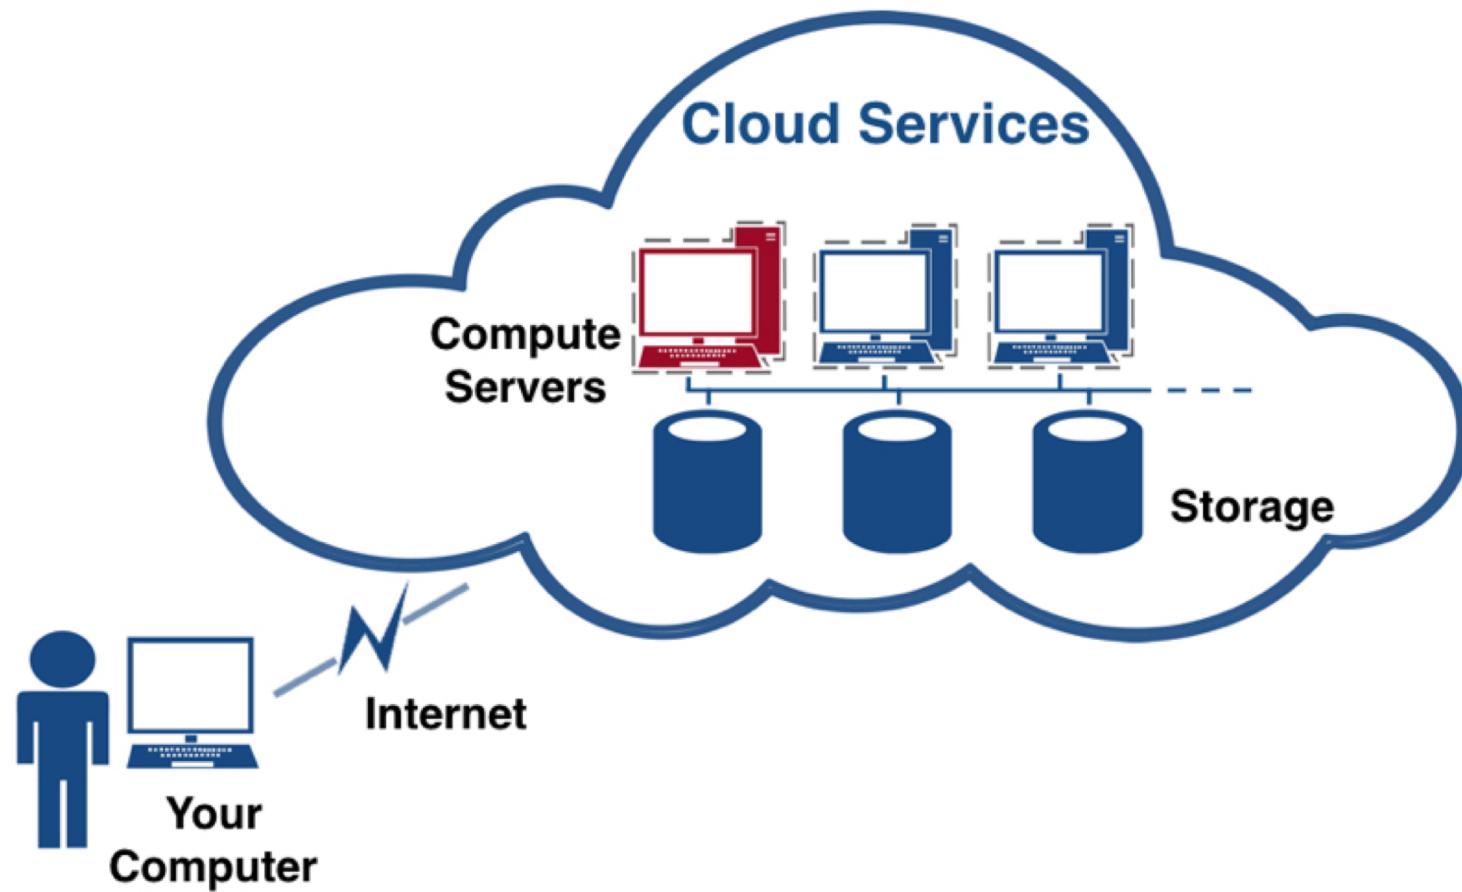
# Cloud 1.0

## Raw virtual resources



UPPSALA  
UNIVERSITET

# Cloud 1.0





UPPSALA  
UNIVERSITET

# Cloud 2.0

# Platforms and applications



UPPSALA  
UNIVERSITET

# Cloud 2.0





UPPSALA  
UNIVERSITET

# Cloud 3.0 Serverless architecture and AI

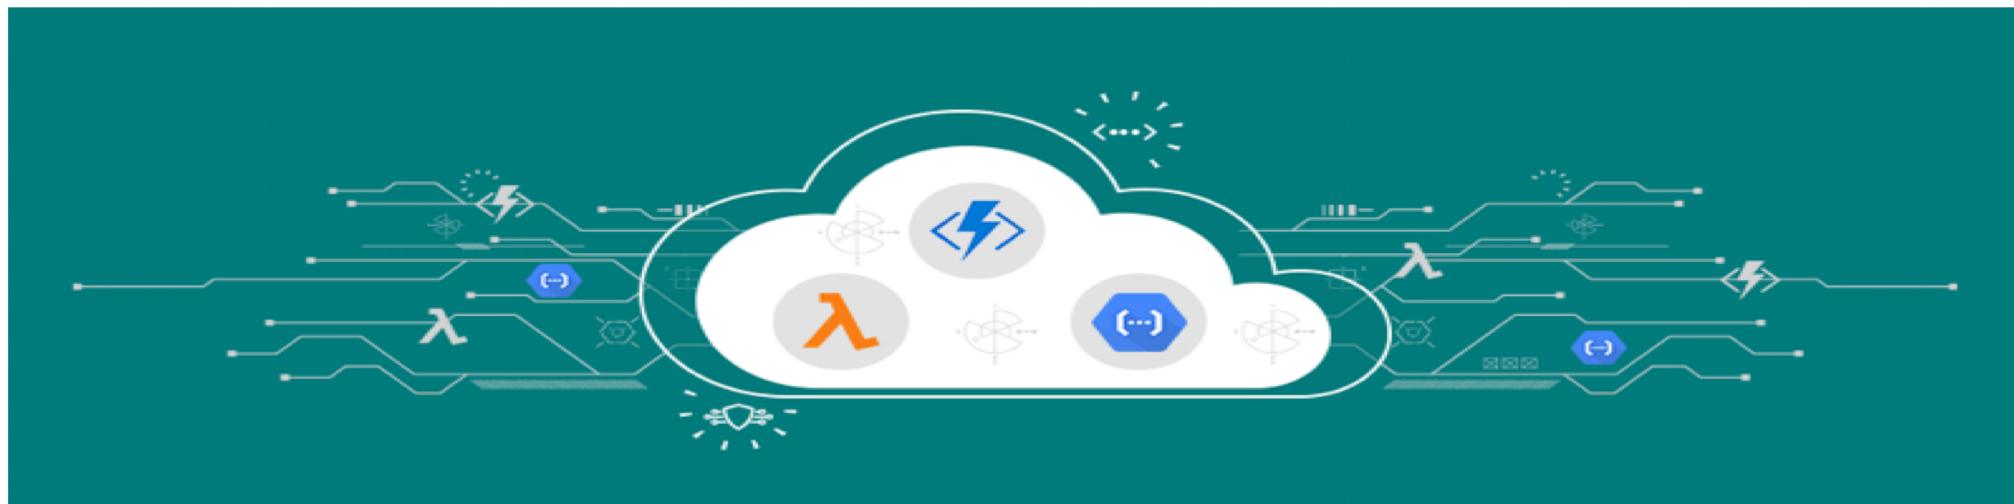


UPPSALA  
UNIVERSITET

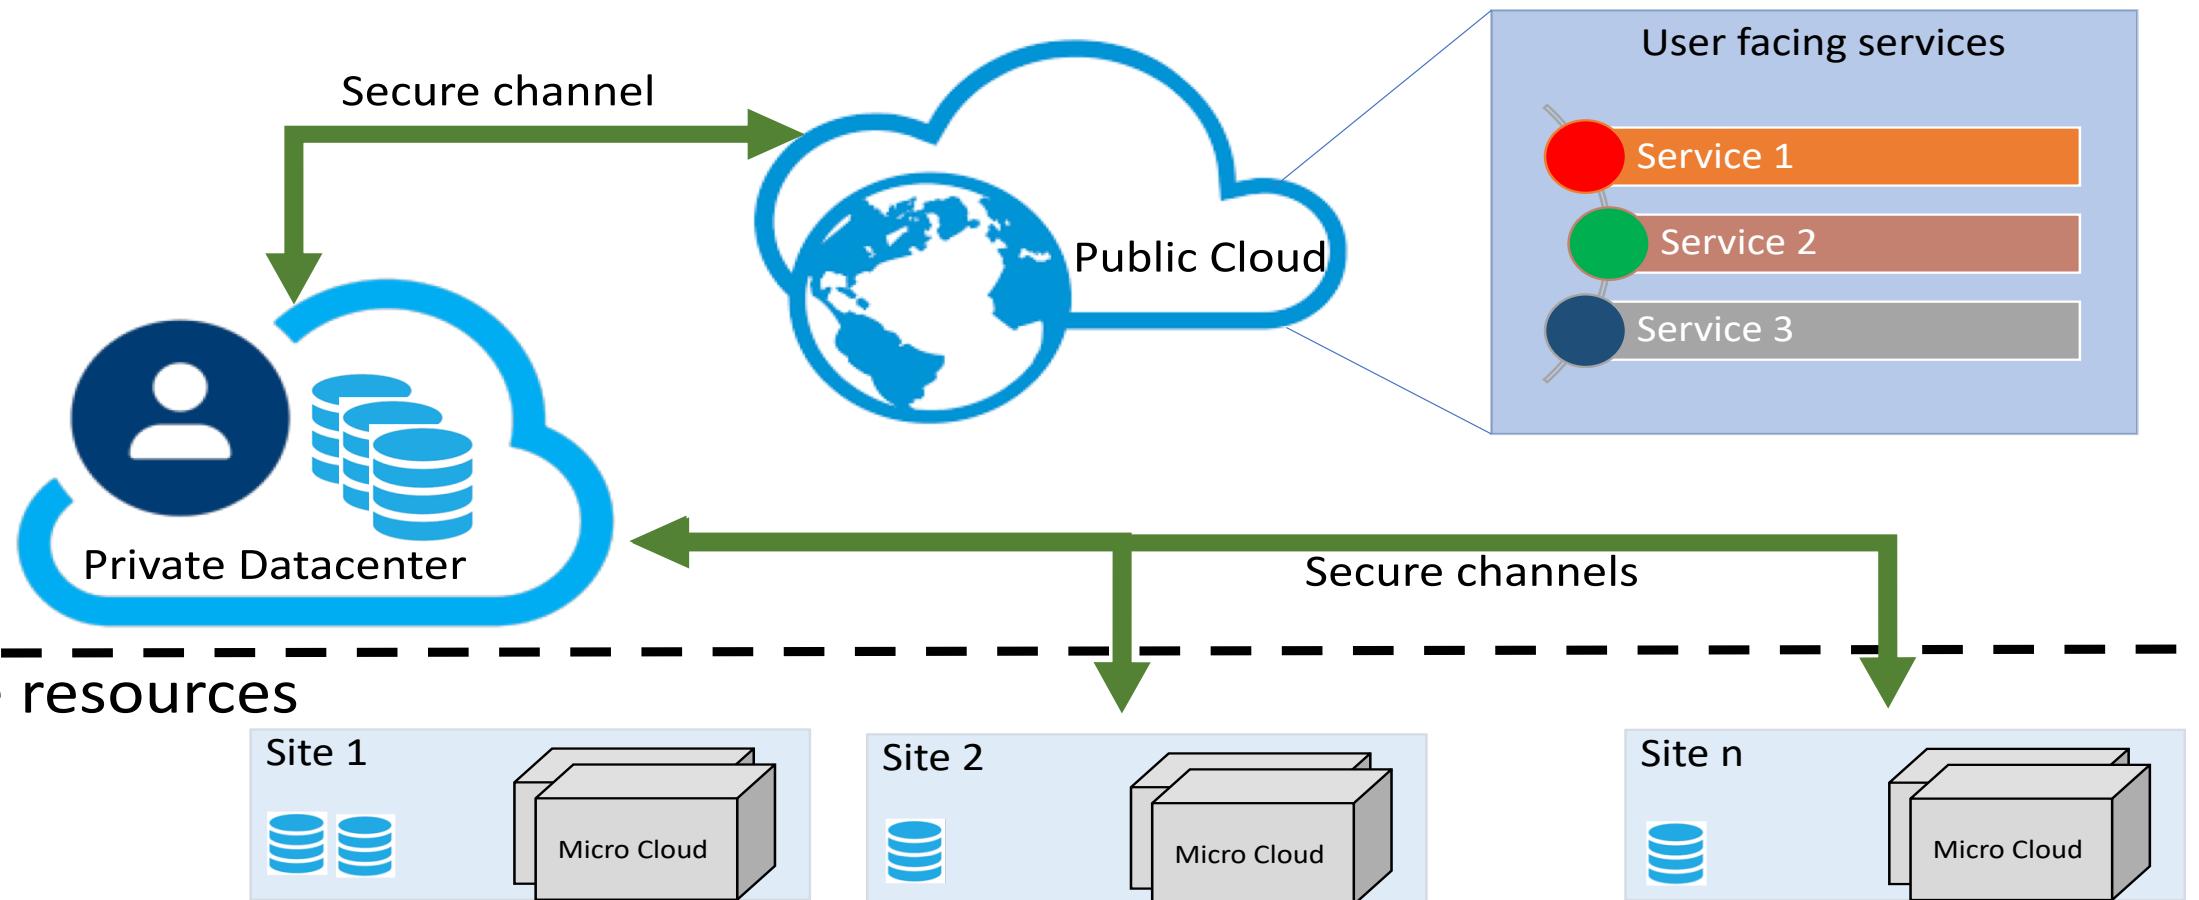
# Cloud 3.0

## Serverless architecture and smart services

- Amazon Lambda
- Google Cloud Functions
- Azure Functions
- OpenFaaS



# Hybrid infrastructure environment





UPPSALA  
UNIVERSITET

# Challenges for Application Experts and Service Providers



# Challenges

- Application design
- Vendor agnostic design



Application Experts

- Data management
- Security and privacy
- Service level agreements
- High availability
- Efficient utilization of resources



Application Experts and  
Service Providers



# Virtualization

- Nutshell: The abstraction of available resources
- Definition:

Virtualization technologies encompass a variety of mechanisms and techniques used to decouple the architecture and user-perceived behavior of hardware and software resources from their physical implementation.

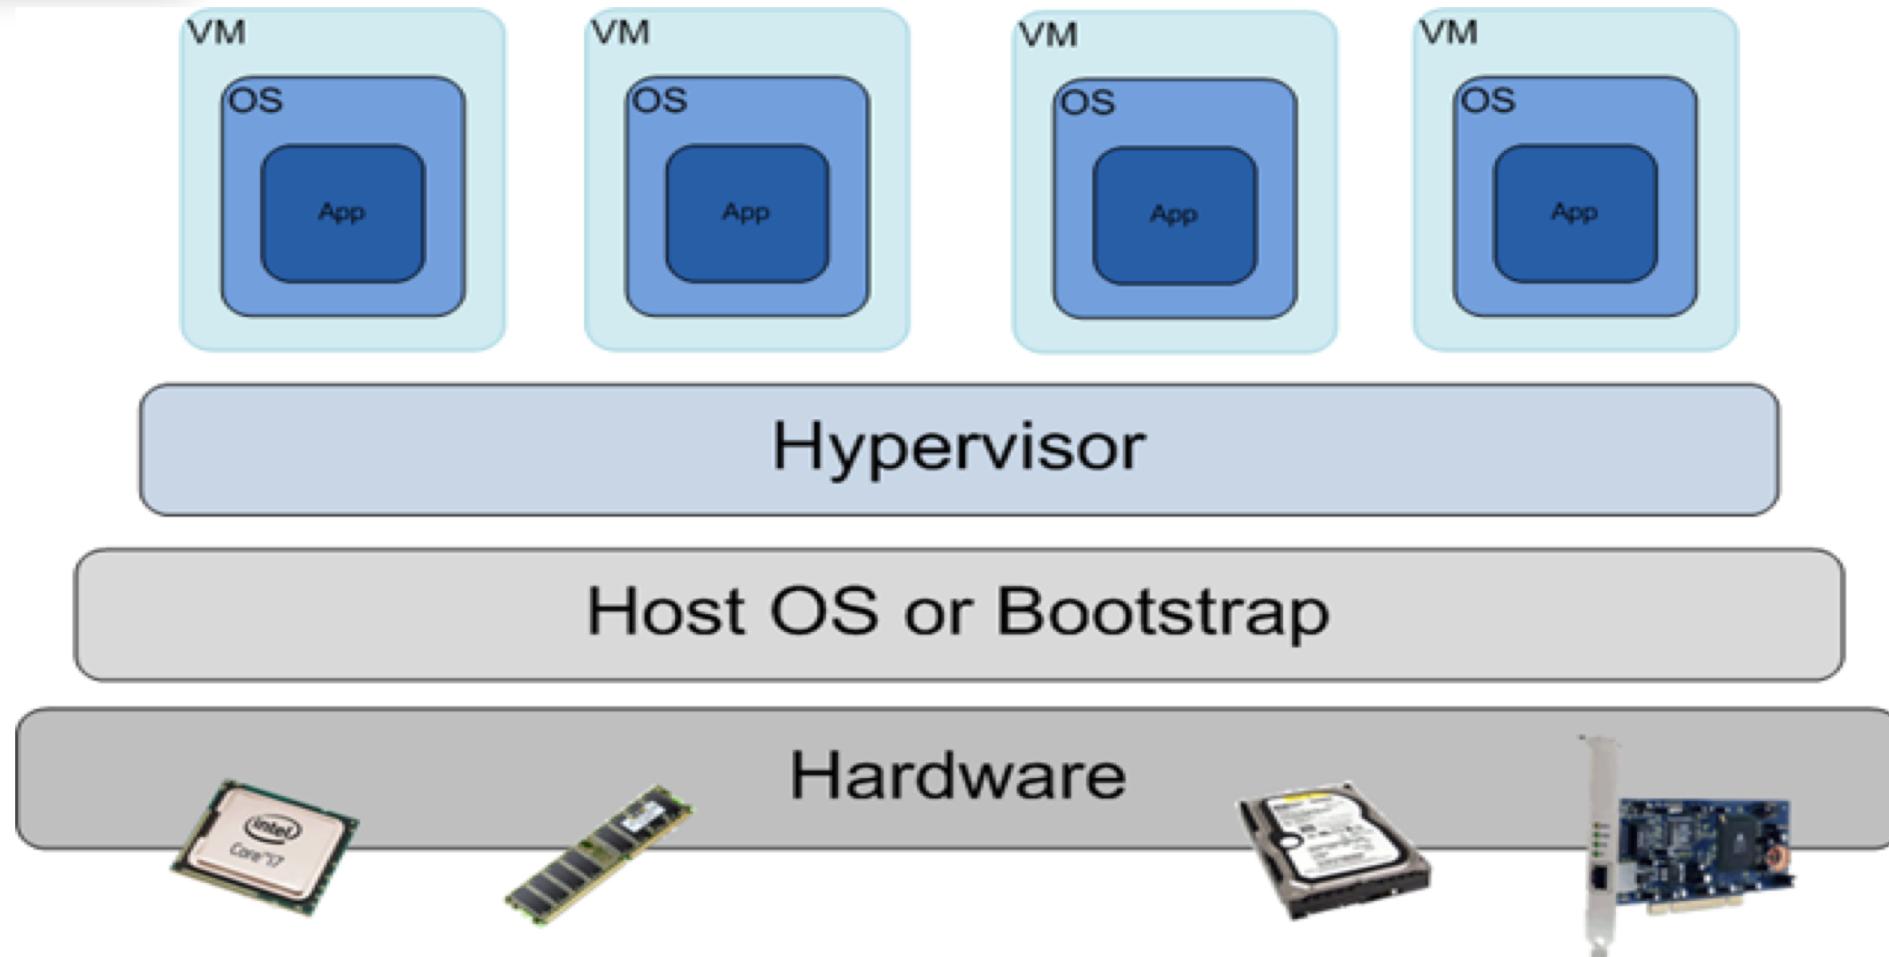
- Whereas, resources can be either compute, storage, network..etc



UPPSALA  
UNIVERSITET

# Virtualization

## Basic illustration





UPPSALA  
UNIVERSITET

# Beyond virtual machines



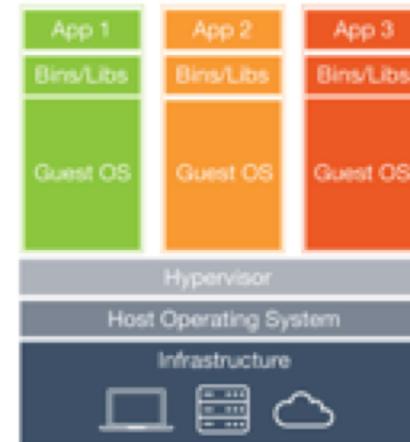
UPPSALA  
UNIVERSITET

# Containers

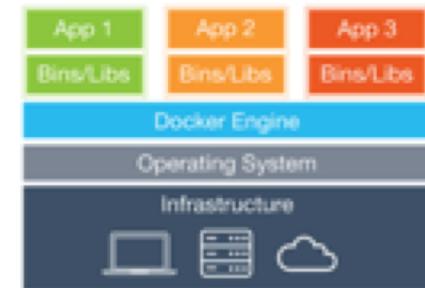
- OS level virtualization environment
- Kernel space is shared
- User space is separate for each Linux system (container)
- A lightweight alternative to Virtual Machines (VM)
- Shared same resources as host OS
- A simple model for packaging applications in Linux



- Docker packages an application together with all its dependencies in the container
- Guarantees that it will always run the same regardless of the environment
- Container based orchestration tool
- Docker Hub, container registry
- Open source



Virtual Machines



Containers

# Images and formats

- Cloud images are customized disks images of OSes for private or public clouds
- Different formats are available:
  - raw: An unstructured disk image format (big in size)
  - vhd: VMware, Xen, Microsoft, VirtualBox, and others
  - vdi: Supported by VirtualBox, QEMU emulator.
  - iso: Archive format for the data contents of an optical disc
  - qcow2: Supported by the QEMU emulator that can expand dynamically and supports Copy on Write.
  - ...

# Contextualization

In cloud computing contextualization means providing customized computing environment

Or

Allows a virtual machine instance to learn about its cloud environment and user requirement (the ‘context’) and configure itself to run correctly

# Contextualization

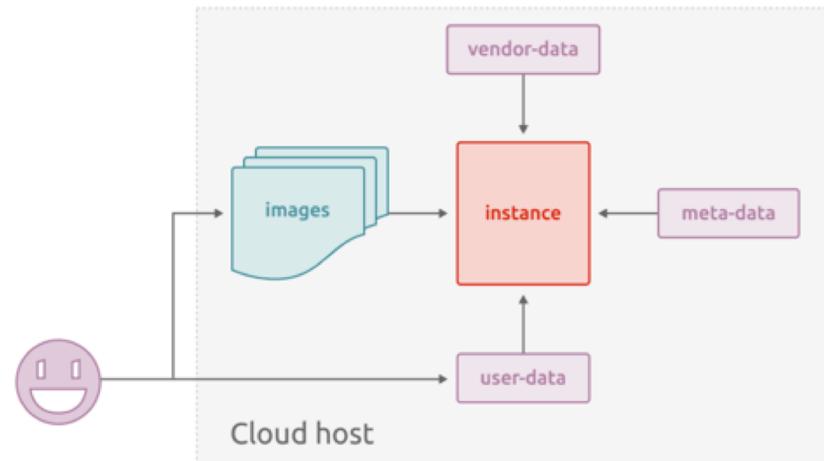
- Provide scalable solution
- No need to manage fat images
- Dynamic configuration
- Typically work in two layers
  - Meta-data : System information handled at cloud level
  - User-defined-data: User specific requirements/settings



UPPSALA  
UNIVERSITET

# CloudInit

Cloud-init is the industry standard multi-distribution method for cross-platform cloud instance initialization. It is supported across all major public cloud providers, provisioning systems for private cloud infrastructure, and bare-metal installations.



<https://cloudinit.readthedocs.io/en/latest/>

<https://www.lucd.info/2019/12/06/cloud-init-part-1-the-basics/>

```
#cloud-config
apt_update: true
apt_upgrade: true
packages:
  - cowsay
  - python-pip
  - python-dev
  - build-essential
  - cowsay
byobu_default: system

runcmd:
  - echo "export PATH=$PATH:/usr/games" >> /home/ubuntu/.bashrc
  - source /home/ubuntu/.bashrc
  - pip install Flask
  - python /home/ubuntu/cowsay-app.py &
```

# CloudInit

# Orchestration

- Orchestration is a process of resource contextualization based on the automation available in the cloud systems.
- A process required for
  - rapid application deployment
  - scalability
  - management
  - high availability
  - agility
- Essential for large complex applications
- A process at the level of Platform as a Service (PaaS)



UPPSALA  
UNIVERSITET

# Orchestration

- Available tools
- Docker Compose <https://www.docker.com/>
- Kubernetes <http://kubernetes.io/>
- Ansible <https://www.ansible.com/>
- Heat <https://wiki.openstack.org/wiki/Heat>



# Containers

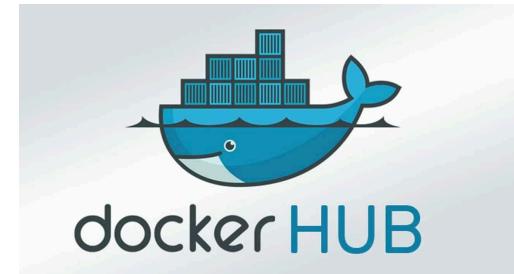
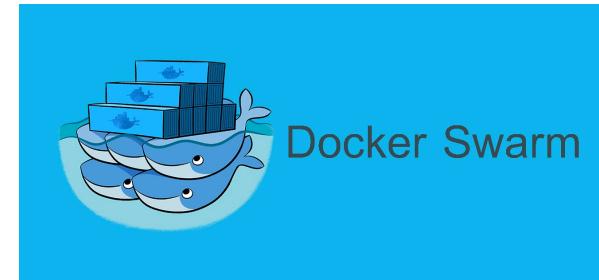
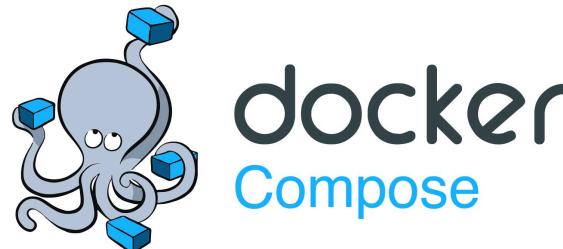
- OS level virtualization environment
  - Kernelspace is shared
  - Userspace is separate for each Linux system (container)
- A lightweight alternative to Virtual Machines (VM)
- Shared same resources as host OS
- A simple model for packaging applications in Linux.



UPPSALA  
UNIVERSITET

# Docker

- Docker packages an application together with all its dependencies in the container
- Guarantees that it will always run the same regardless of the environment
- Container based orchestration tool
- Docker Hub, container registry
- Open source



# Infrastructure as Code (IaC)

Infrastructure as code, also referred to as IaC, is an IT practice that codifies and **manages underlying IT infrastructure as software**. The purpose of infrastructure as code is to enable developers or operations teams to automatically **manage, monitor and provision resources**, rather than manually configure discrete hardware devices and operating systems.

## Infrastructure resources

Compute  
Storage  
Network  
High level services

## Operating systems and Host configurations

Linux  
Windows  
Access services  
Host level security

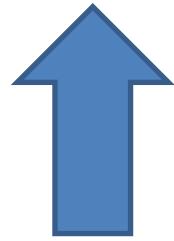
## Application configuration

Application dependencies  
Application configurations  
Data access control  
Application level security  
Application access

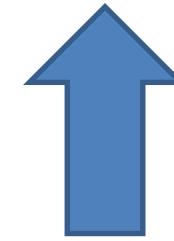


UPPSALA  
UNIVERSITET

# Infrastructure as Code (IaC)



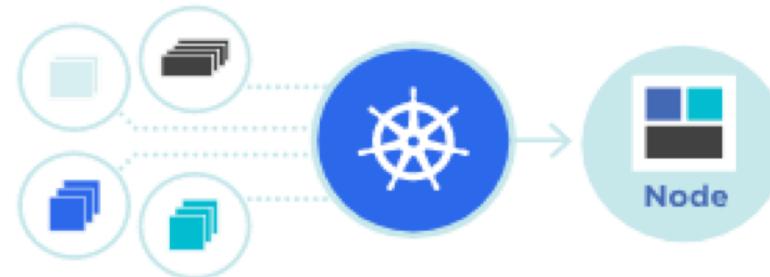
## Contextualization



## Orchestration

# Kubernetes

Kubernetes, also known as K8s, is an open-source system for automating deployment, scaling, and management of containerized applications.



# Kubernetes

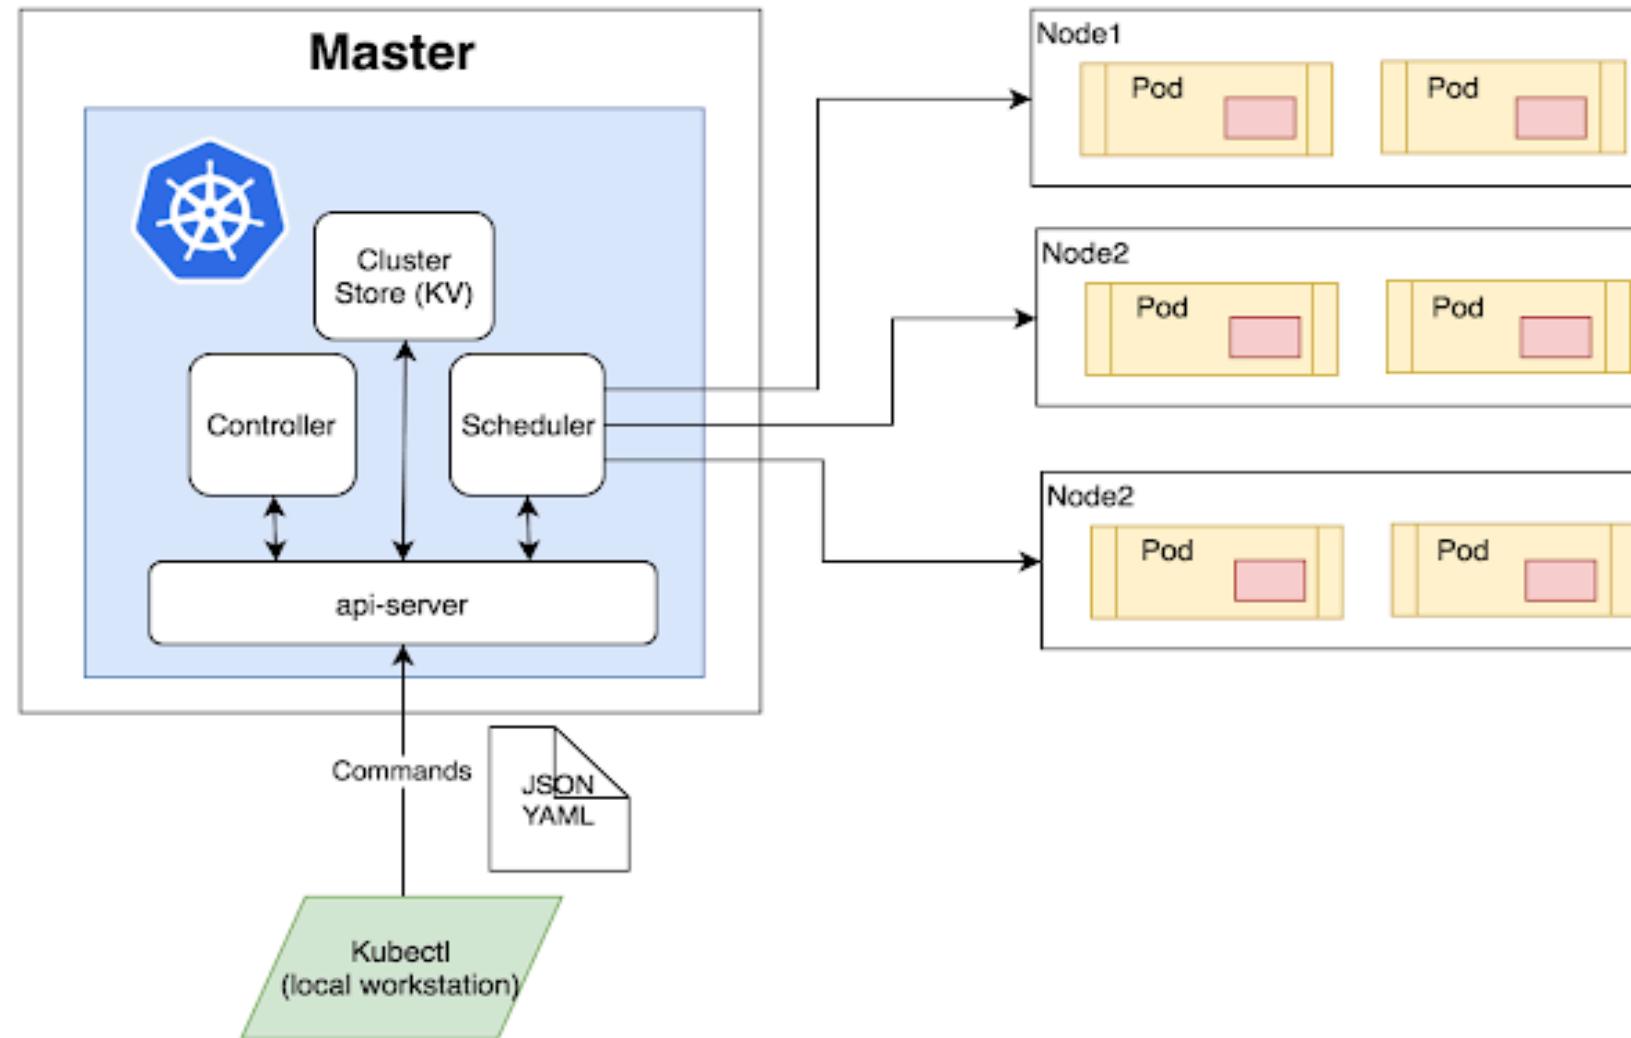
- Features
  - Automation
  - Scheduling
  - Self healing capabilities
  - Horizontal scaling
  - Load balancing
  - Environment consistency
  - Application centric management
  - Built-in services
  - Effective management of credentials
  - Application centric access control

- Basic concepts
  - **Cluster:** A collection of hosts aggregating their resources (CPU, RAM, Storage and Network)
  - **Master:** The main node of the cluster, forms control panel of Kubernetes
  - **Node:** A single host in the cluster, running essential services
  - **Namespaces:** A logical organization of the resources



UPPSALA  
UNIVERSITET

# Kubernetes Architecture



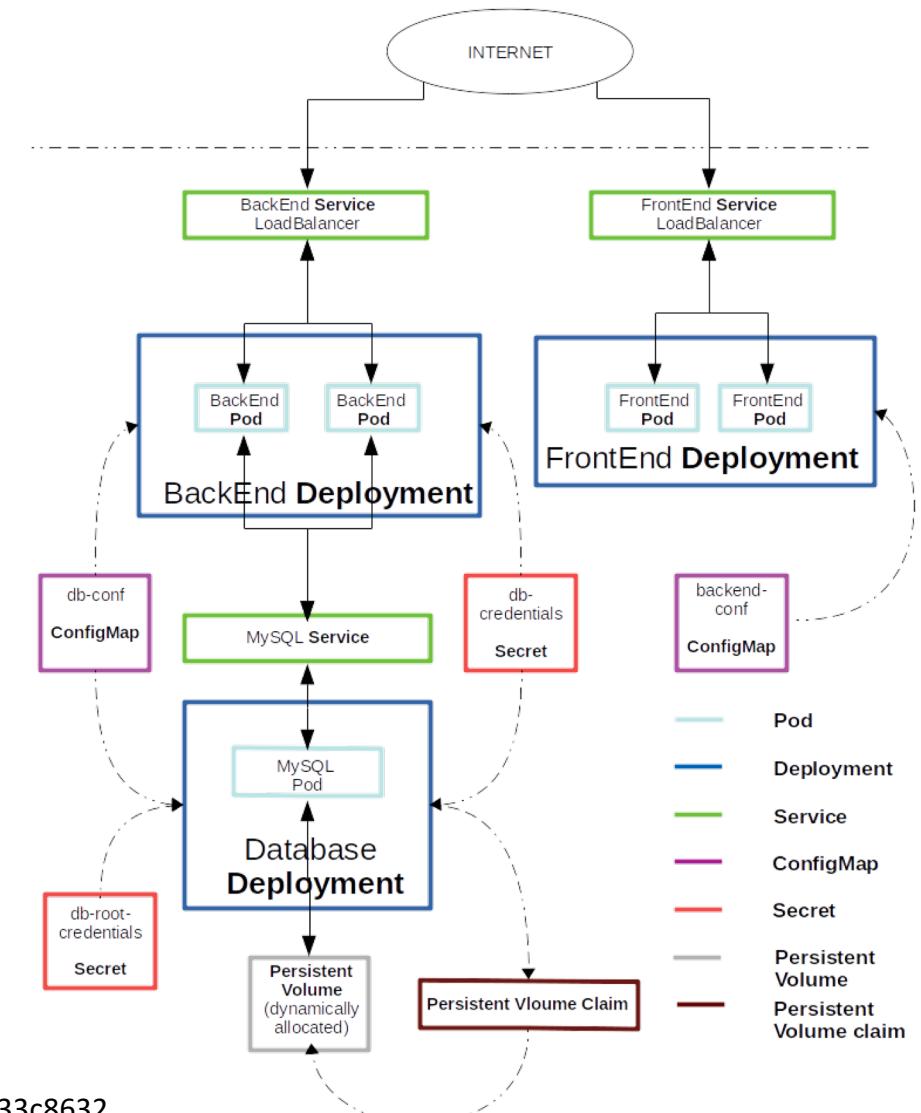
# Kubernetes

- Selected list of K8s objects
  - **Pods:** A group of one or more containers.
  - **Service:** An abstraction which defines a logical set of pods and a policy to access them. Every Kubernetes service has an IP address, but unlike the IP address of a pod, it is stable.
  - **Persistent-Volume and Persistent-Volume Claim:** A persistent-volume is a piece of storage in the cluster that has been provisioned to be used. Whereas a persistent-volume claim is a request by an application to consume the abstract storage resources declared through persistent volume.
  - **ConfigMap:** Another abstraction to decouple environment-dependent application-configuration-data from containerized applications.
  - **Secrets:** An object that contains a small amount of sensitive data such as a password, a token, or a key.
  - **Deployment:** An abstraction to represent the desired state of an actual deployment on Kubernetes



# Kubernetes Example

- Web application:
  - Database (MySQL server)
  - Backend (Java Spring Boot application)
  - Frontend (Angular app)
- Configuration settings:
  - Database, one replica
  - Backend, two replicas
  - Frontend, two replicas



# Dockerfile

## Frontend-dockerfile

```
FROM node:7.7-alpine

# install dependencies
ADD package.json /tmp/package.json
RUN cd /tmp && npm install

# Copy dependencies
RUN mkdir -p /opt/to-do-app && cp -a /tmp/node_modules /opt/to-do-app

# Setup workdir
WORKDIR /opt/to-do-app

RUN mkdir /opt/to-do-app/dist
COPY dist /opt/to-do-app/dist
COPY server.js /opt/to-do-app
COPY initialize.js /opt/to-do-app

# run
EXPOSE 8080
CMD ["sh", "-c", "node initialize && node server"]
```

## Backend-dockerfile

```
FROM openjdk:10-jre-slim
RUN mkdir -p /opt/to-do-app/
COPY ./target/to-do-listEntity-app-0.0.1-SNAPSHOT.jar /opt/to-do-app/
WORKDIR /opt/to-do-app/
EXPOSE 8080
CMD ["java", "-jar", "to-do-listEntity-app-0.0.1-SNAPSHOT.jar"]
```

Upload the containers to dockerHub

```
# docker push image-name
```



UPPSALA  
UNIVERSITET

# Kubernetes

- Installation
  - Kubernetes: <https://kubernetes.io>
  - MiniKube: <https://minikube.sigs.k8s.io/docs/start/>
  - MicroK8s: <https://microk8s.io/>



# Kubernetes

## backend-spring-application.yaml

```
spring:  
  datasource:  
    type: com.zaxxer.hikari.HikariDataSource  
    hikari:  
      idle-timeout: 10000  
    platform: mysql  
    username: ${DB_USERNAME}  
    password: ${DB_PASSWORD}  
    url: jdbc:mysql://${DB_HOST}/${DB_NAME}  
  
  jpa:  
    hibernate:  
      naming:  
        physical-strategy: org.hibernate.boot.model.naming.PhysicalNamingStrategyStandardImpl
```

## db-root-credentials-secret.yaml

```
# Define 'Secret' to store db-credentials  
apiVersion: v1  
kind: Secret  
metadata:  
  name: db-credentials # Name of the 'Secret'  
data:  
  username: dXNlcg== # base64 encoded 'Secret' username  
  password: ZGV2ZWxvcA== # base64 encoded 'Secret' password
```

## mysql-configmap.yaml

```
# Define 'Configmap' to store non-sensitive database configuration  
apiVersion: v1  
kind: ConfigMap  
metadata:  
  name: db-conf # name of ConfigMap, referenced in other files  
data:  
  host: mysql # host address of mysql server, we are using DNS of Service  
  name: to-do-app-db # name of the database for application
```

## db-root-credentials-secret.yaml

```
# Define 'Secret' to store 'root' user Credentials  
apiVersion: v1  
kind: Secret  
metadata:  
  name: db-root-credentials # Name of the Secret  
data:  
  password: bWFnaWM= # base64 encoded 'root' password
```

```
# kubectl apply -f <FILE_NAME>
```



# Kubernetes

## backend-spring-application.yaml

```
# Define a 'Service' To Expose mysql to Other Services
apiVersion: v1
kind: Service
metadata:
  name: mysql # DNS name
  labels:
    app: mysql
    tier: database
spec:
  ports:
    - port: 3306
      targetPort: 3306
  selector: # mysql Pod Should contain same labels
    app: mysql
    tier: database
  clusterIP: None # We Use DNS, Thus ClusterIP is not relevant
---
# Define a 'Persistent Voulume Claim'(PVC) for Mysql Storage, dynamically provisioned by cluster
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pv-claim # name of PVC essential for identifying the storage data
  labels:
    app: mysql
    tier: database
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
---
```



# Kubernetes

```
# Configure 'Deployment' of mysql server
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mysql
  labels:
    app: mysql
    tier: database
spec:
  selector: # mysql Pod Should contain same labels
  matchLabels:
    app: mysql
    tier: database
  strategy:
    type: Recreate
  template:
    metadata:
      labels: # Must match 'Service' and 'Deployment' selectors
      app: mysql
      tier: database
    spec:
      containers:
        - image: mysql:5.7 # image from docker-hub
          args:
            - "--ignore-db-dir=lost+found" # Workaround for https://github.com/docker-library/mysql/issues/186
          name: mysql
          env:
            - name: MYSQL_ROOT_PASSWORD # Setting Root Password of mysql From a 'Secret'
              valueFrom:
                secretKeyRef:
                  name: db-root-credentials # Name of the 'Secret'
                  key: password # 'key' inside the Secret which contains required 'value'
            - name: MYSQL_USER # Setting USER username on mysql From a 'Secret'
              valueFrom:
                secretKeyRef:
                  name: db-credentials
                  key: username
```



# Kubernetes

```
- name: MYSQL_PASSWORD # Setting USER Password on mysql From a 'Secret'  
  valueFrom:  
    secretKeyRef:  
      name: db-credentials  
      key: password  
- name: MYSQL_DATABASE # Setting Database Name from a 'ConfigMap'  
  valueFrom:  
    configMapKeyRef:  
      name: db-conf  
      key: name  
ports:  
- containerPort: 3306  
  name: mysql  
volumeMounts:    # Mounting volume obtained from Persistent Volume Claim  
- name: mysql-persistent-storage  
  mountPath: /var/lib/mysql  
volumes:  
- name: mysql-persistent-storage # Obtaining 'volume' from PVC  
  persistentVolumeClaim:  
    claimName: mysql-pv-claim
```

## backend-deployment.yaml

```
# Define 'Service' to expose backend application deployment
apiVersion: v1
kind: Service
metadata:
  name: to-do-app-backend
spec:
  selector: # backend application pod lables should match these
    app: to-do-app
    tier: backend
  ports:
    - protocol: "TCP"
      port: 80
      targetPort: 8080
  type: LoadBalancer # use NodePort, if you are not running Kubernetes on cloud
---
# Configure 'Deployment' of backend application
apiVersion: apps/v1
kind: Deployment
metadata:
  name: to-do-app-backend
  labels:
    app: to-do-app
    tier: backend
spec:
  replicas: 2 # Number of replicas of back-end application to be deployed
  selector:
    matchLabels: # backend application pod labels should match these
      app: to-do-app
      tier: backend
```

```
template:
  metadata:
    labels: # Must macth 'Service' and 'Deployment' labels
      app: to-do-app
      tier: backend
  spec:
    containers:
      - name: to-do-app-backend
        image: kubernetesdemo/to-do-app-backend # docker image of backend application
        env: # Setting Enviornmental Variables
          - name: DB_HOST # Setting Database host address from configMap
            valueFrom:
              configMapKeyRef:
                name: db-conf # name of configMap
                key: host
          - name: DB_NAME # Setting Database name from configMap
            valueFrom:
              configMapKeyRef:
                name: db-conf
                key: name
          - name: DB_USERNAME # Setting Database username from Secret
            valueFrom:
              secretKeyRef:
                name: db-credentials # Secret Name
                key: username
          - name: DB_PASSWORD # Setting Database password from Secret
            valueFrom:
              secretKeyRef:
                name: db-credentials
                key: password
    ports:
      - containerPort: 8080
```



```
# ConfigMap to expose configuration related to backend application
apiVersion: v1
kind: ConfigMap
metadata:
  name: backend-conf # name of configMap
data:
  server-uri: 34.66.207.42 # enternal ip of backend application 'Service'
```

### frontend-deployment.yaml

```
# Define 'Service' to expose FrontEnd Application
apiVersion: v1
kind: Service
metadata:
  name: to-do-app-frontend
spec:
  selector: # pod labels should match these
    app: to-do-app
    tier: frontend
  ports:
    - protocol: "TCP"
      port: 80
      targetPort: 8080
  type: LoadBalancer # use NodePort if you are not running Kubernetes on Cloud
---
```

```
# 'Deployment' to manage of configuration of frontEnd Deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: to-do-app-frontend
  labels: # pod labels should match these
    app: to-do-app
    tier: frontend
spec:
  replicas: 2 # number of replicas of frontEnd application
  selector:
    matchLabels:
      app: to-do-app
      tier: frontend
  template:
    metadata:
      labels: # Must match 'Service' and 'Deployment' labels
        app: to-do-app
        tier: frontend
    spec:
      containers:
        - name: to-do-app-frontend
          image: kubernetesdemo/to-do-app-frontend # docker image of frontend application
          env: # Setting Environmental Variables
            - name: SERVER_URI # Setting Backend URI from configMap
              valueFrom:
                configMapKeyRef:
                  name: backend-conf # Name of configMap
                  key: server-uri
          ports:
            - containerPort: 8080
```



UPPSALA  
UNIVERSITET

# Complete K8s online course



YouTube Link:

<https://www.youtube.com/watch?v=X48VuDVv0do&t=1s>



UPPSALA  
UNIVERSITET

# Application-as-a-Package



# Helm charts

- A package manager for Kubernetes
- Helm chart can contain any number of Kubernetes objects.
- Helm chart usually contains:
  - deployment
  - service
  - Ingress
  - Persistent Volume Claims
  - or any other Kubernetes object



UPPSALA  
UNIVERSITET

# Serverless architecture



Serverless architecture is a concept to build and run software and services without having to manage servers. In serverless architecture, the server will manage by provider and not by consumer.

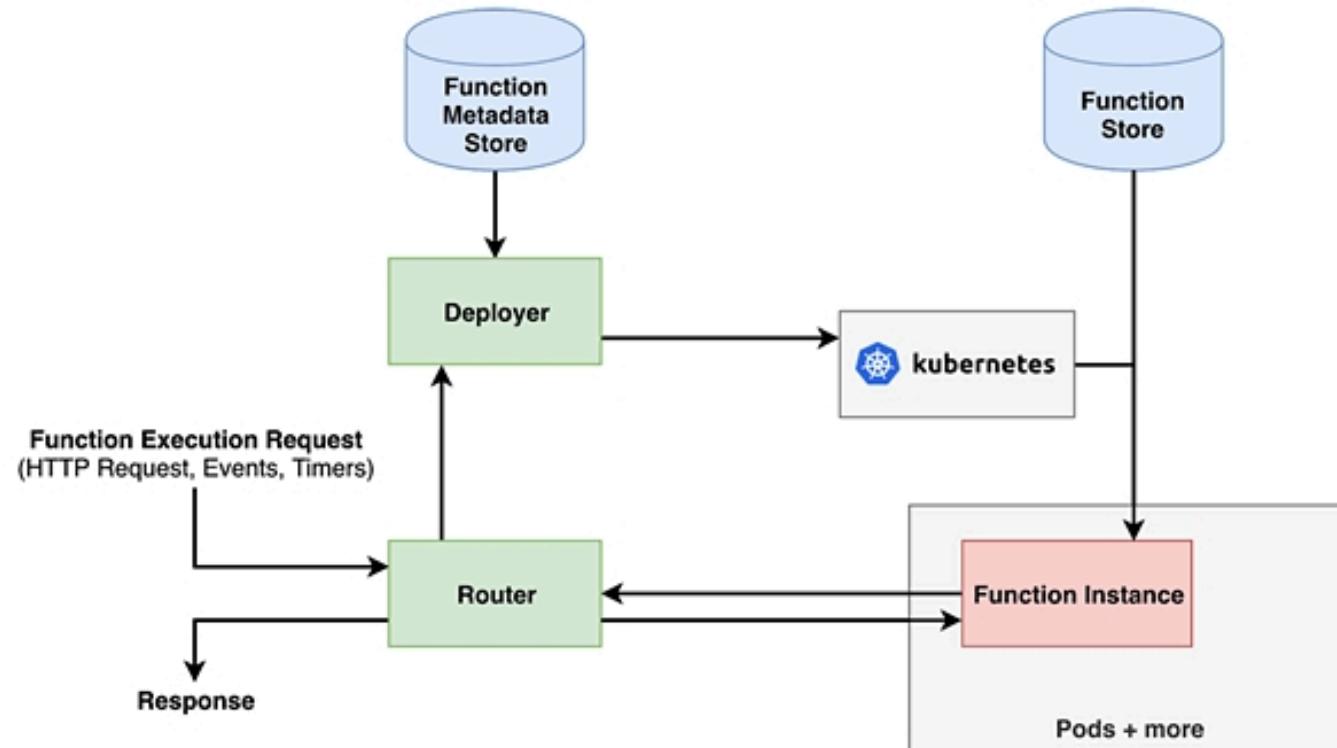
# FaaS: Function as a Service

- FaaS is a realization of the serverless architecture
- FaaS allows to execute code in response to events without the complex infrastructure typically associated with building and launching microservices applications.
- FaaS consists of lightweight functions that subdivide the application's logic and only executed based on generated events.

# FaaS: Function as a Service

- Characteristics:
  - Rapid development
  - Event driven execution
  - Cost effective computing environment (Pay-as-you-run)
  - No resource selection, configuration and management
  - Better elastic behaviour

# FaaS: Function as a Service



The anatomy of the runtime of a FaaS platform.

# FaaS: Function as a Service

- FaaS Platforms
  - Firecracker, AWS Lambda and AWS Fargate (Amazon)
  - Fission (Platform 9)
  - Fn (Oracle)
  - Knative/Kubernetes (Google)
  - OpenFaaS