



UPPSALA  
UNIVERSITET

# Lecture-5: Distributed Infrastructures (Part-II)

Salman Toor

[salman.toor@it.uu.se](mailto:salman.toor@it.uu.se)



UPPSALA  
UNIVERSITET

# Distributed Computing Infrastructures: Part-I

- Infrastructure as Code (IaC)
  - Contextualization
  - Orchestration
- Kubernetes
- Application-as-a-Package
  - Helm charts



UPPSALA  
UNIVERSITET

# Distributed Computing Infrastructures: Part-II

- Serverless architecture
  - FaaS
- Continuous Integration and Continuous Delivery (CI/CD)
  - Git Hooks
- Model Serving
  - Model as Code (MaC)
  - Model as Data (MaD)
- Lab-2/Assignment discussion
- Case studies
  - KubeNow
  - Stackn



UPPSALA  
UNIVERSITET

# Serverless architecture



Serverless architecture is a concept to build and run software and services without having to manage servers. In serverless architecture, the server will manage by provider and not by consumer.

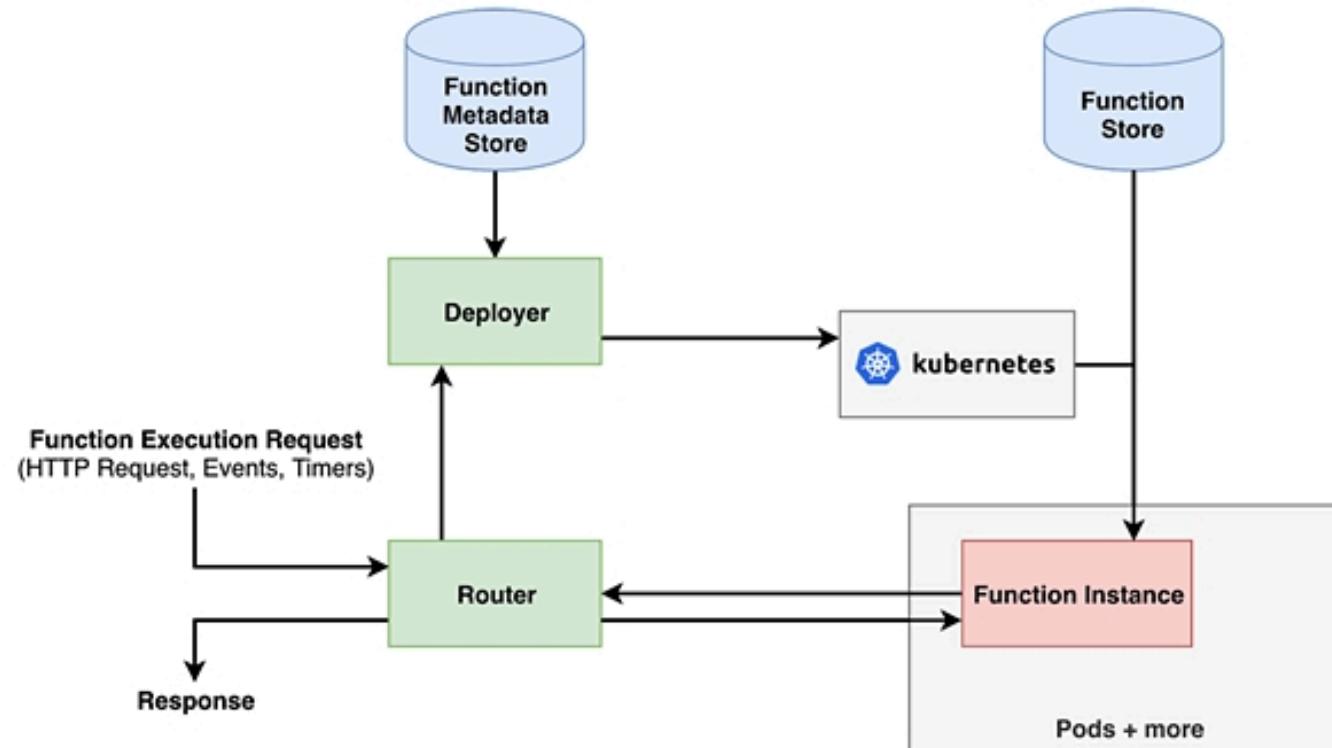
# FaaS: Function as a Service

- FaaS is a realization of the serverless architecture
- FaaS allows to execute code in response to events without the complex infrastructure typically associated with building and launching microservices applications.
- FaaS consists of lightweight functions that subdivide the application's logic and only executed based on generated events.

# FaaS: Function as a Service

- Characteristics:
  - Rapid development
  - Event driven execution
  - Cost effective computing environment (Pay-as-you-run)
  - No resource selection, configuration and management
  - Better elastic behaviour

# FaaS: Function as a Service



The anatomy of the runtime of a FaaS platform.

# FaaS: Function as a Service

- FaaS Platforms
  - Firecracker, AWS Lambda and AWS Fargate (Amazon)
  - Fission (Platform 9)
  - Fn (Oracle)
  - Knative/Kubernetes (Google)
  - OpenFaaS



UPPSALA  
UNIVERSITET

# Continuous Integration and Continuous Delivery (CI/CD)

- Continuous Integration (CI)
  - A software engineering process of avoiding enormous code conflicts at the end of a development cycle
- Continuous Delivery (CD)
  - A process of ensuring that software is always ready to be deployed
  - Continuous Delivery requires Continuous Integration

Continuous Delivery: <https://martinfowler.com/books/continuousDelivery.html>

[https://link.springer.com/chapter/10.1007/978-3-319-97925-0\\_33](https://link.springer.com/chapter/10.1007/978-3-319-97925-0_33)

# Continuous Integration(CI)

- Benefits
  - Systematic development with less number of bugs
  - Building the release is easy
  - Reduce the development time
  - Testing costs are reduced, a CI server can run hundreds of tests in parallel
  - Less time required by the quality assurance team, can focus on overall system improvement than minor application related details.

# Continuous Delivery(CD)

- Benefits
  - Faster and systematic deployment with checkpointing
  - Reduce the transition time between development and production environments
  - Frequent updates are possible in production pipelines
  - Increase confidence in software releases



UPPSALA  
UNIVERSITET

# Git Hooks

- Custom software modules that run automatically every time a particular event occurs in a Git repository
- Use to customize the Git's internal behavior to build reliable and systematic execution pipeline
- *hooks* directory is by default available in every git repo ".git/hooks"
- Some of the samples:
  - *pre-commit.sample*: invoked just before making the commit.
  - *post-receive.sample*: invoked after the remote repository has updated.



UPPSALA  
UNIVERSITET

# Continuous Integration and Continuous Delivery (CI/CD)

Local model training and development environment

Deployment Server

Git Client

0

Production Server

Jump directory  
.git/hooks

Project files

2

Test Server

Jump directory  
.git/hooks

Project files

1

Step – 0 & 1, Continuous Integration and Continuous Delivery

Step – 2 Continuous Deployment

# Tools for CI/CD development

- Jenkins: <https://jenkins.io/>
- CircleCI: <https://circleci.com/>
- TeamCity: <https://www.jetbrains.com/teamcity/>
- Bamboo: <https://www.atlassian.com/software/bamboo>
- GitLab: <https://about.gitlab.com/>

<https://www.katalon.com/resources-center/blog/ci-cd-tools/>

# What is model serving?

Model serving is a recently introduced concept to offer machine learning as a service (MLaaS). Mainly designed in following two ways:

1. Dedicate inference serving (service for a particular application)
2. Open inference serving (inference service for different applications within defined settings)

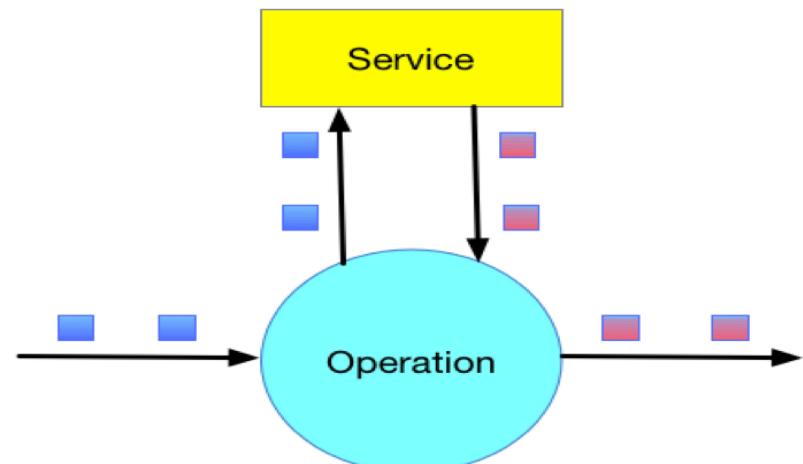
The aim is to simplify the process by offering reliable, easy to use and efficient predictions at scale

# Model serving and stream processing

*Model servers simplify the task of deploying machine learning at scale, the same way app servers simplify the task of delivering a web app or API to end users. The rise of model servers, will likely accelerate the adoption of user-facing machine learning in the wild.*

Two popular approaches for model server implementation:

- Model as Code (MaC)
- Model as Data (MaD)



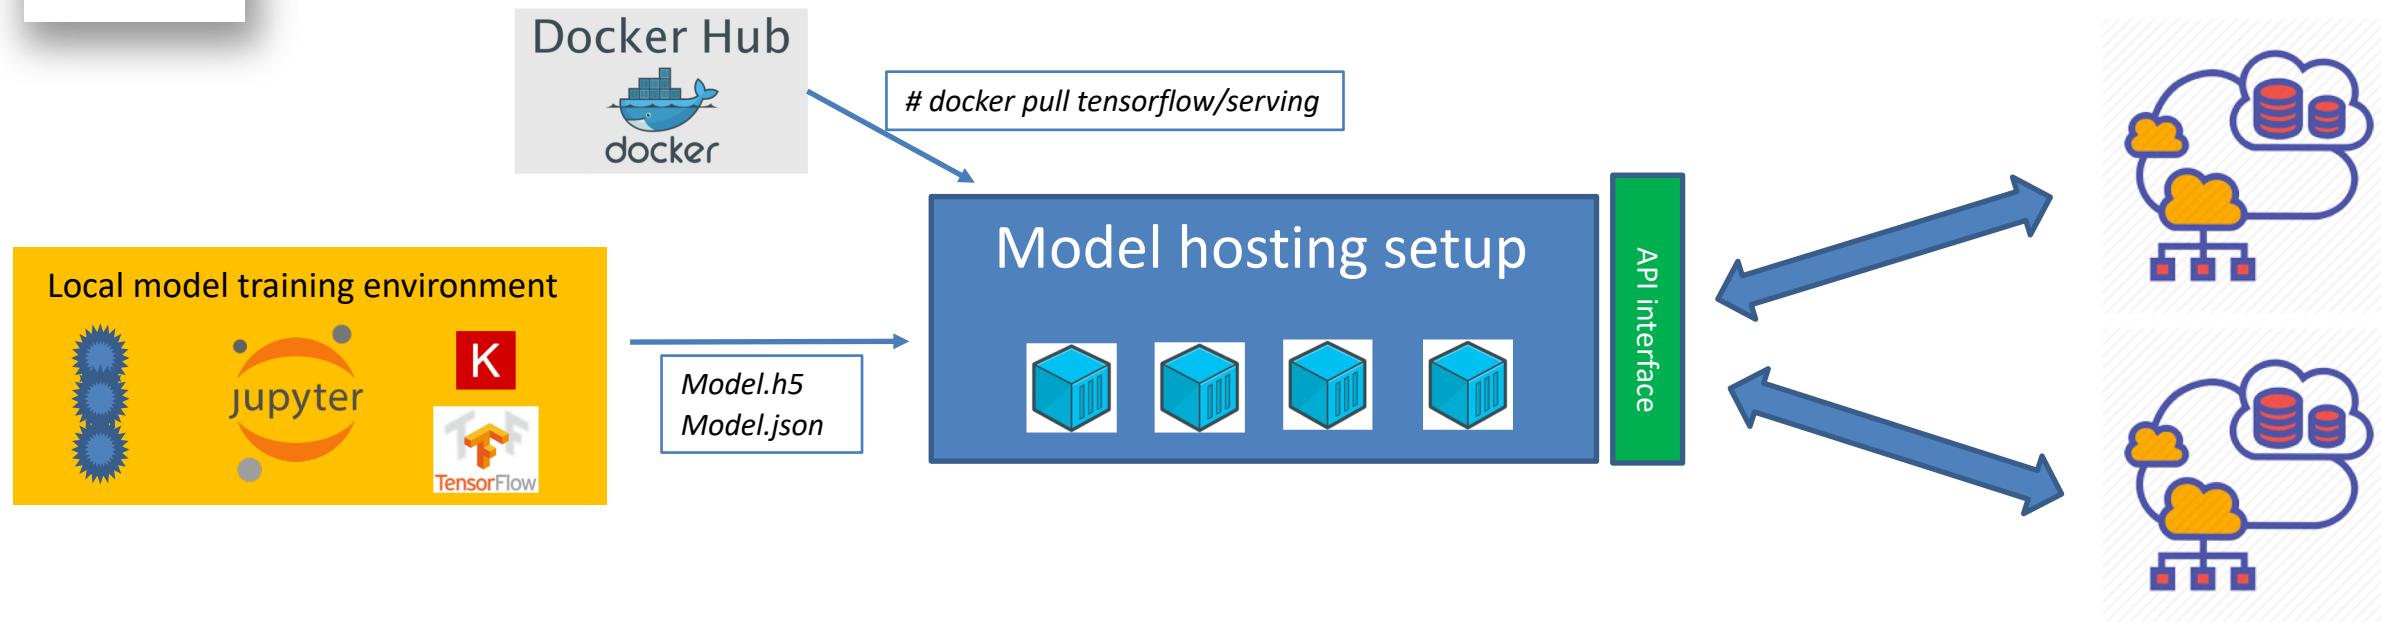
# Model Servers

Two popular approaches for model server implementation:

- Model as Code (MaC): Enables model as a function that take input and produce inference
- Model as Data (MaD): Allows use of different runtimes for machine learning and model serving



# TensorFlow model serving



1. Save model – both weights and architecture
  - For different versions, “saved\_model/<version>”, versions = 1,2 or 3
2. Load TensorFlow-serving container from the Docker Hub
3. Setup the model directory, config file, container parameters and start the container
4. Use both web interface or APIs to get the predictions

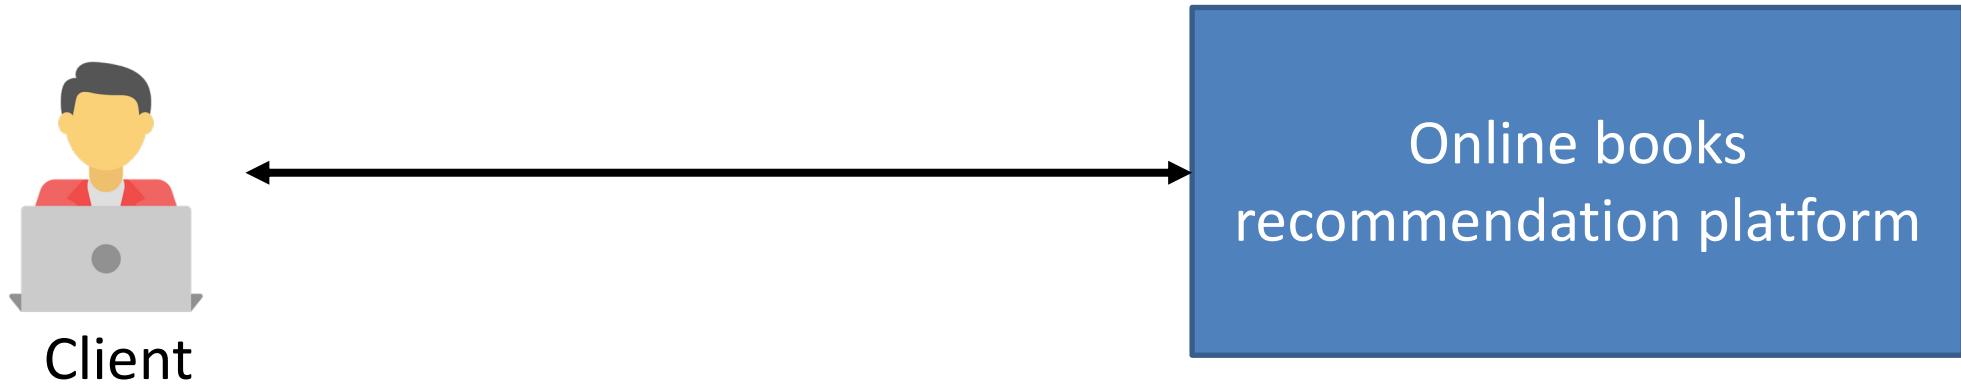
[https://www.tensorflow.org/tfx/serving/serving\\_basic](https://www.tensorflow.org/tfx/serving/serving_basic)

<https://hub.docker.com/r/tensorflow/serving>

# Model serving platforms

- TensorFlow Serving: [http://learningsys.org/nips17/assets/papers/paper\\_1.pdf](http://learningsys.org/nips17/assets/papers/paper_1.pdf)
- Clipper: <https://www.usenix.org/system/files/conference/nsdi17/nsdi17-crankshaw.pdf>
- Model Server for Apache MXNet: <https://aws.amazon.com/blogs/ai/introducing-model-server-for-apache-mxnet/>
- DeepDetect: <https://deepdetect.com>
- TensorRT: <https://developer.nvidia.com/tensorrt>

# Example



- Users interaction:
  - Search new books
  - Buy books
  - Get recommendations based on previous searches or purchases
  - Listen to online books
  - Advertisements for non-paying customers

# Deployment stage - 1

Can this architecture manage hundreds of requests per minute or second?



Client



Online books  
recommendation platform

Web Server

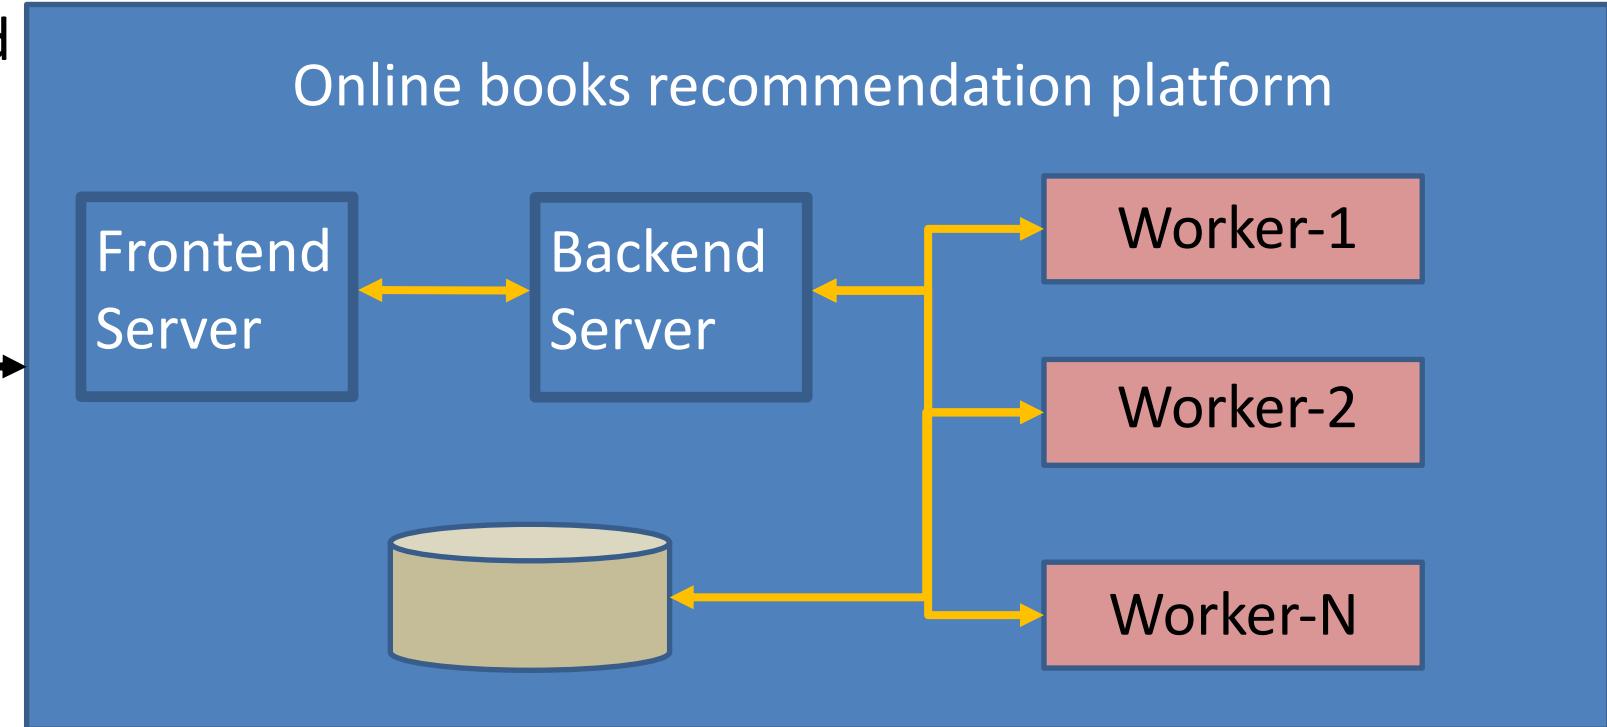
- Internal architecture
  - Frontend Server, manages incoming requests
  - Backend Server, runs the logic, maintains the communication with the internal components of the application

# Deployment stage - 1

What if the frontend or backend servers have some issues?

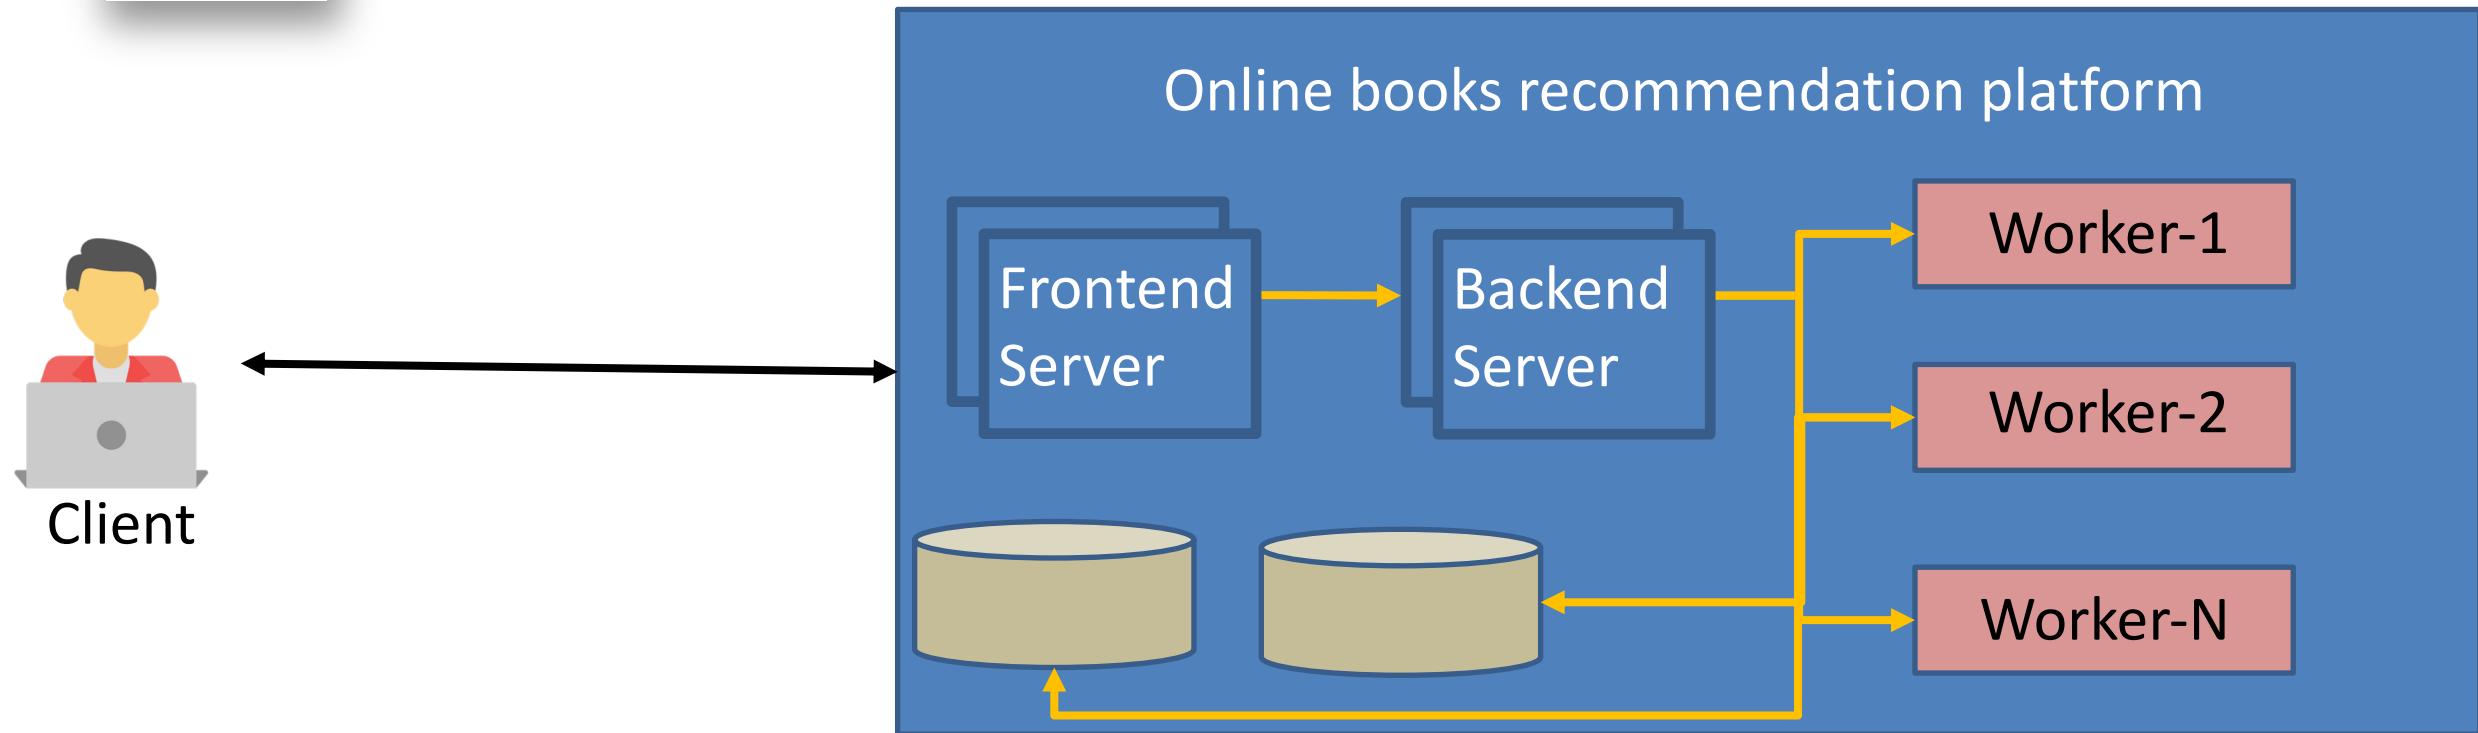


Client



- Internal architecture
  - Frontend Server, manages incoming requests
  - Backend Server, distribute the tasks to the available workers
  - Workers, execute the assigned tasks

# Deployment stage - 1

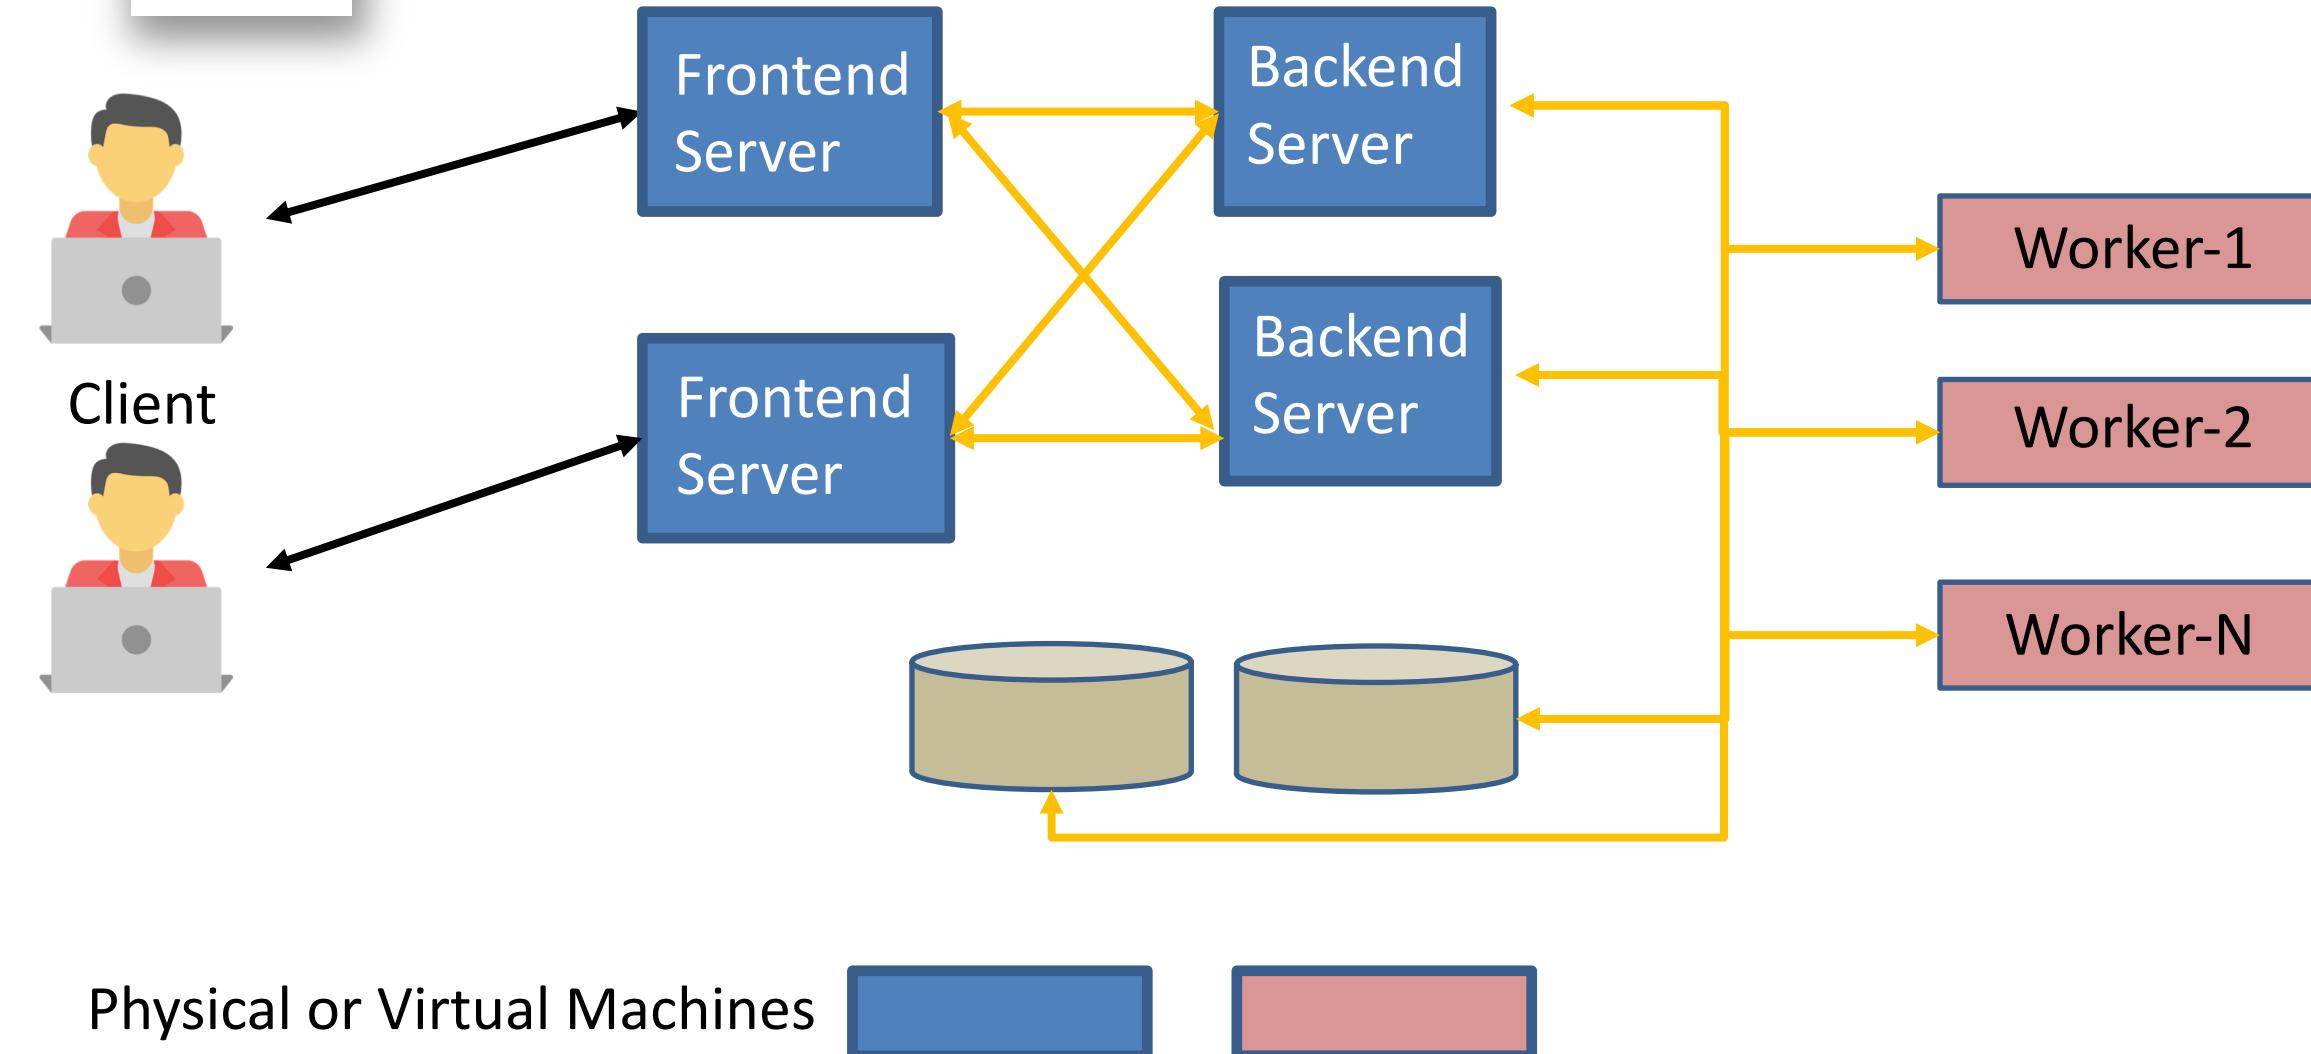


- Internal architecture
  - Frontend Server, manages incoming requests
  - Backend Server, distribute the tasks to the available workers
  - Workers, execute the assigned tasks



UPPSALA  
UNIVERSITET

# Deployment stage - 2





UPPSALA  
UNIVERSITET

# We need fast, systematic and efficient deployment

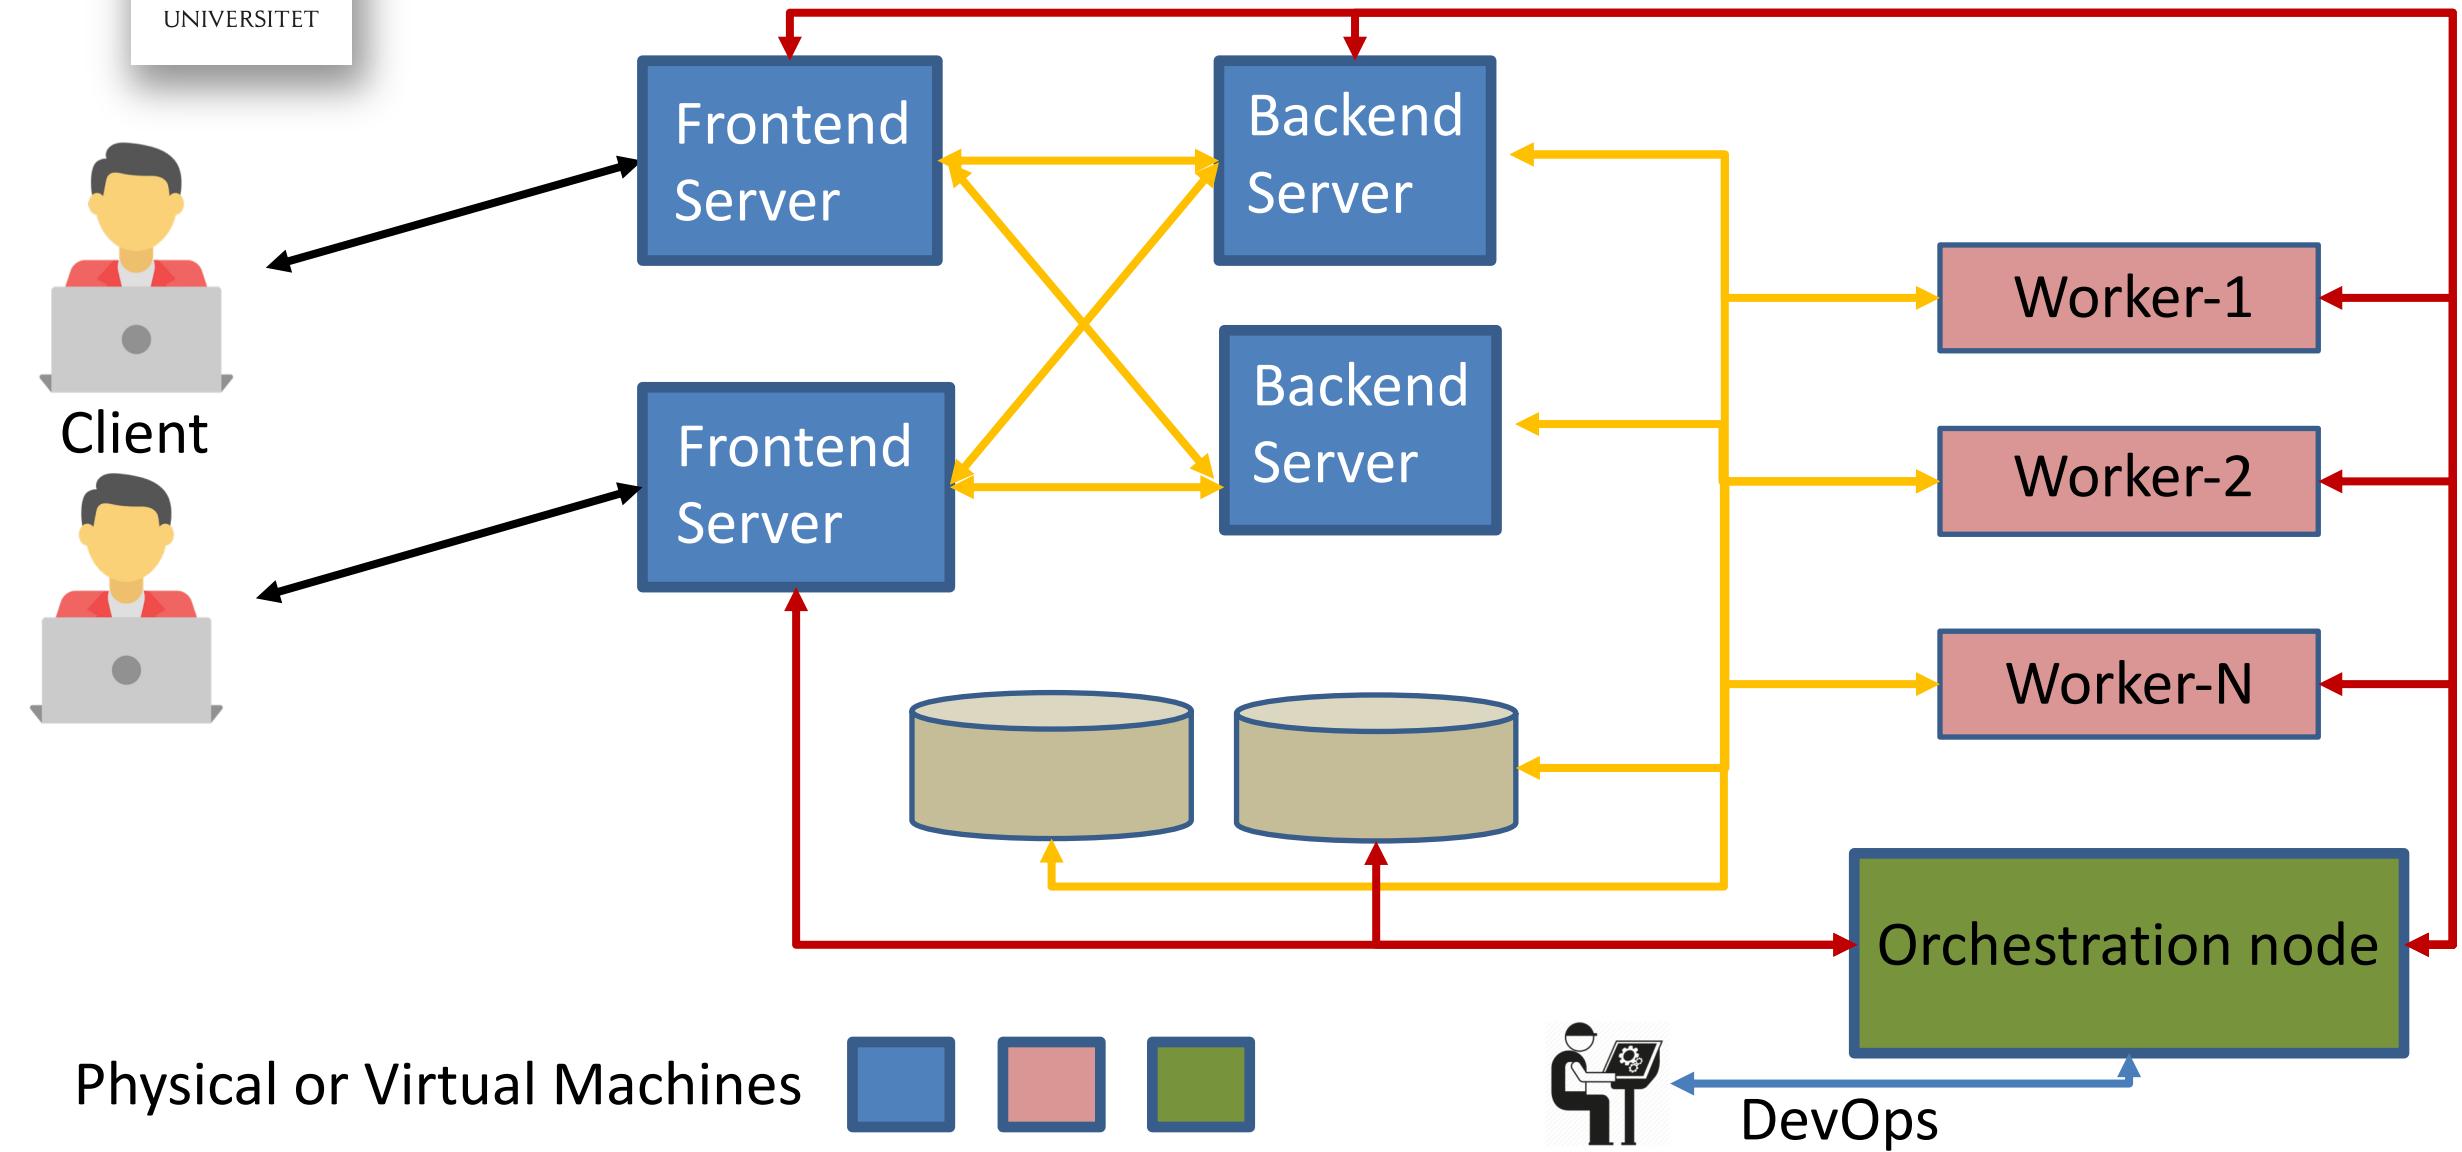
Cloud platforms are good but we need tools for

- Automation
- Contextualization
- Orchestration



UPPSALA  
UNIVERSITET

# Deployment stage - 2





UPPSALA  
UNIVERSITET

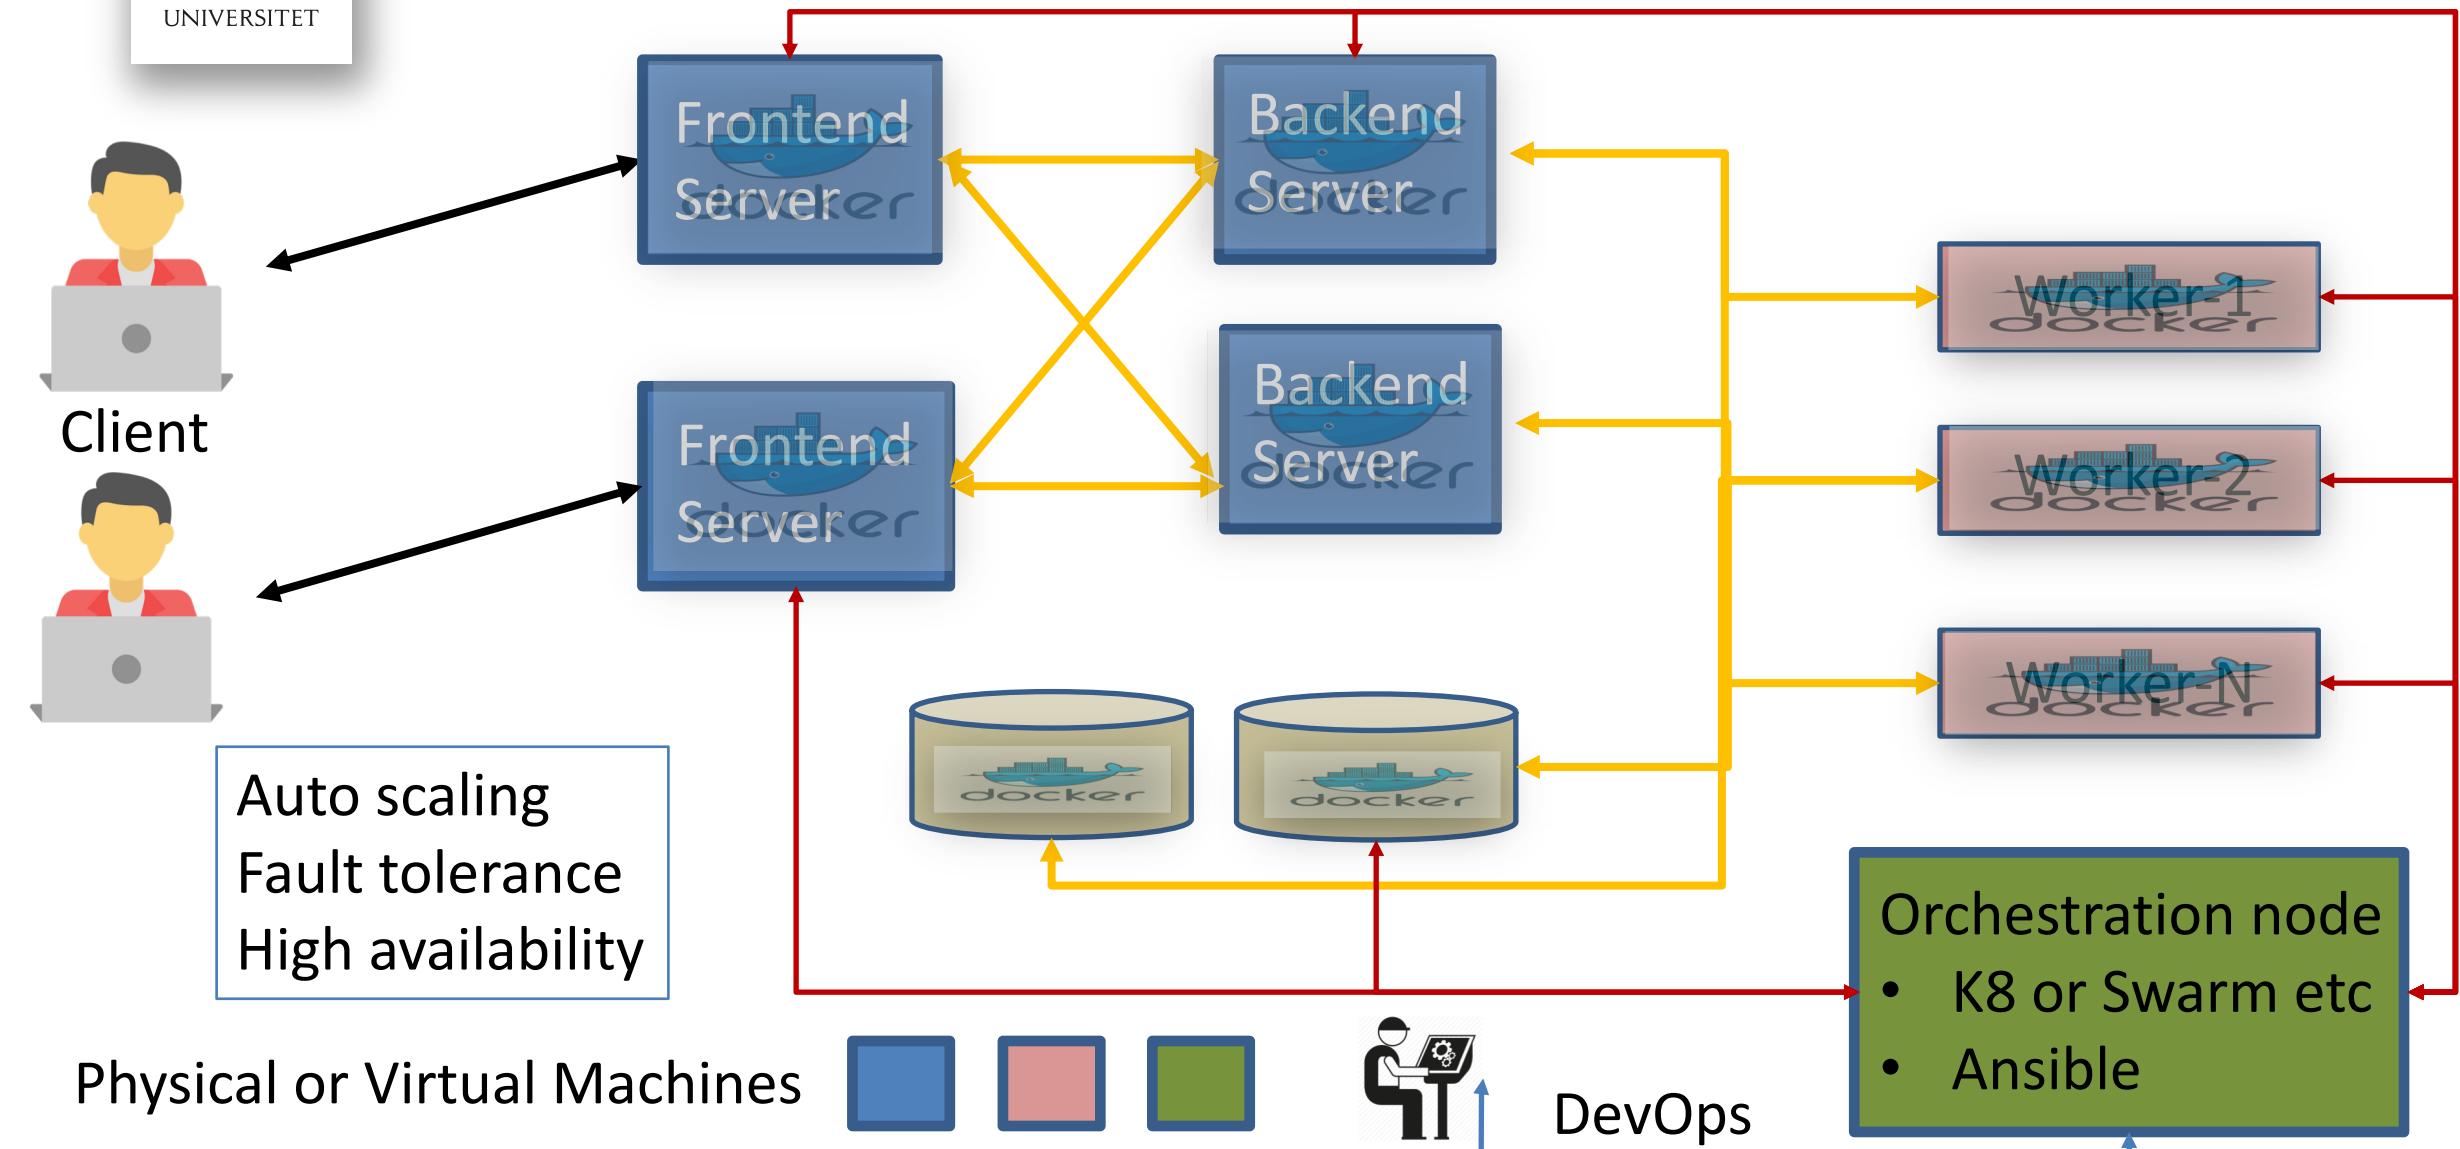
## More improvements to make the platform highly scalable and fault-tolerant

- Avoid tight coupling with the host operating system
- Checkpoint and serialize each step required to make modifications
- Geographically distributed setup for high availability



UPPSALA  
UNIVERSITET

# Deployment stage - 2



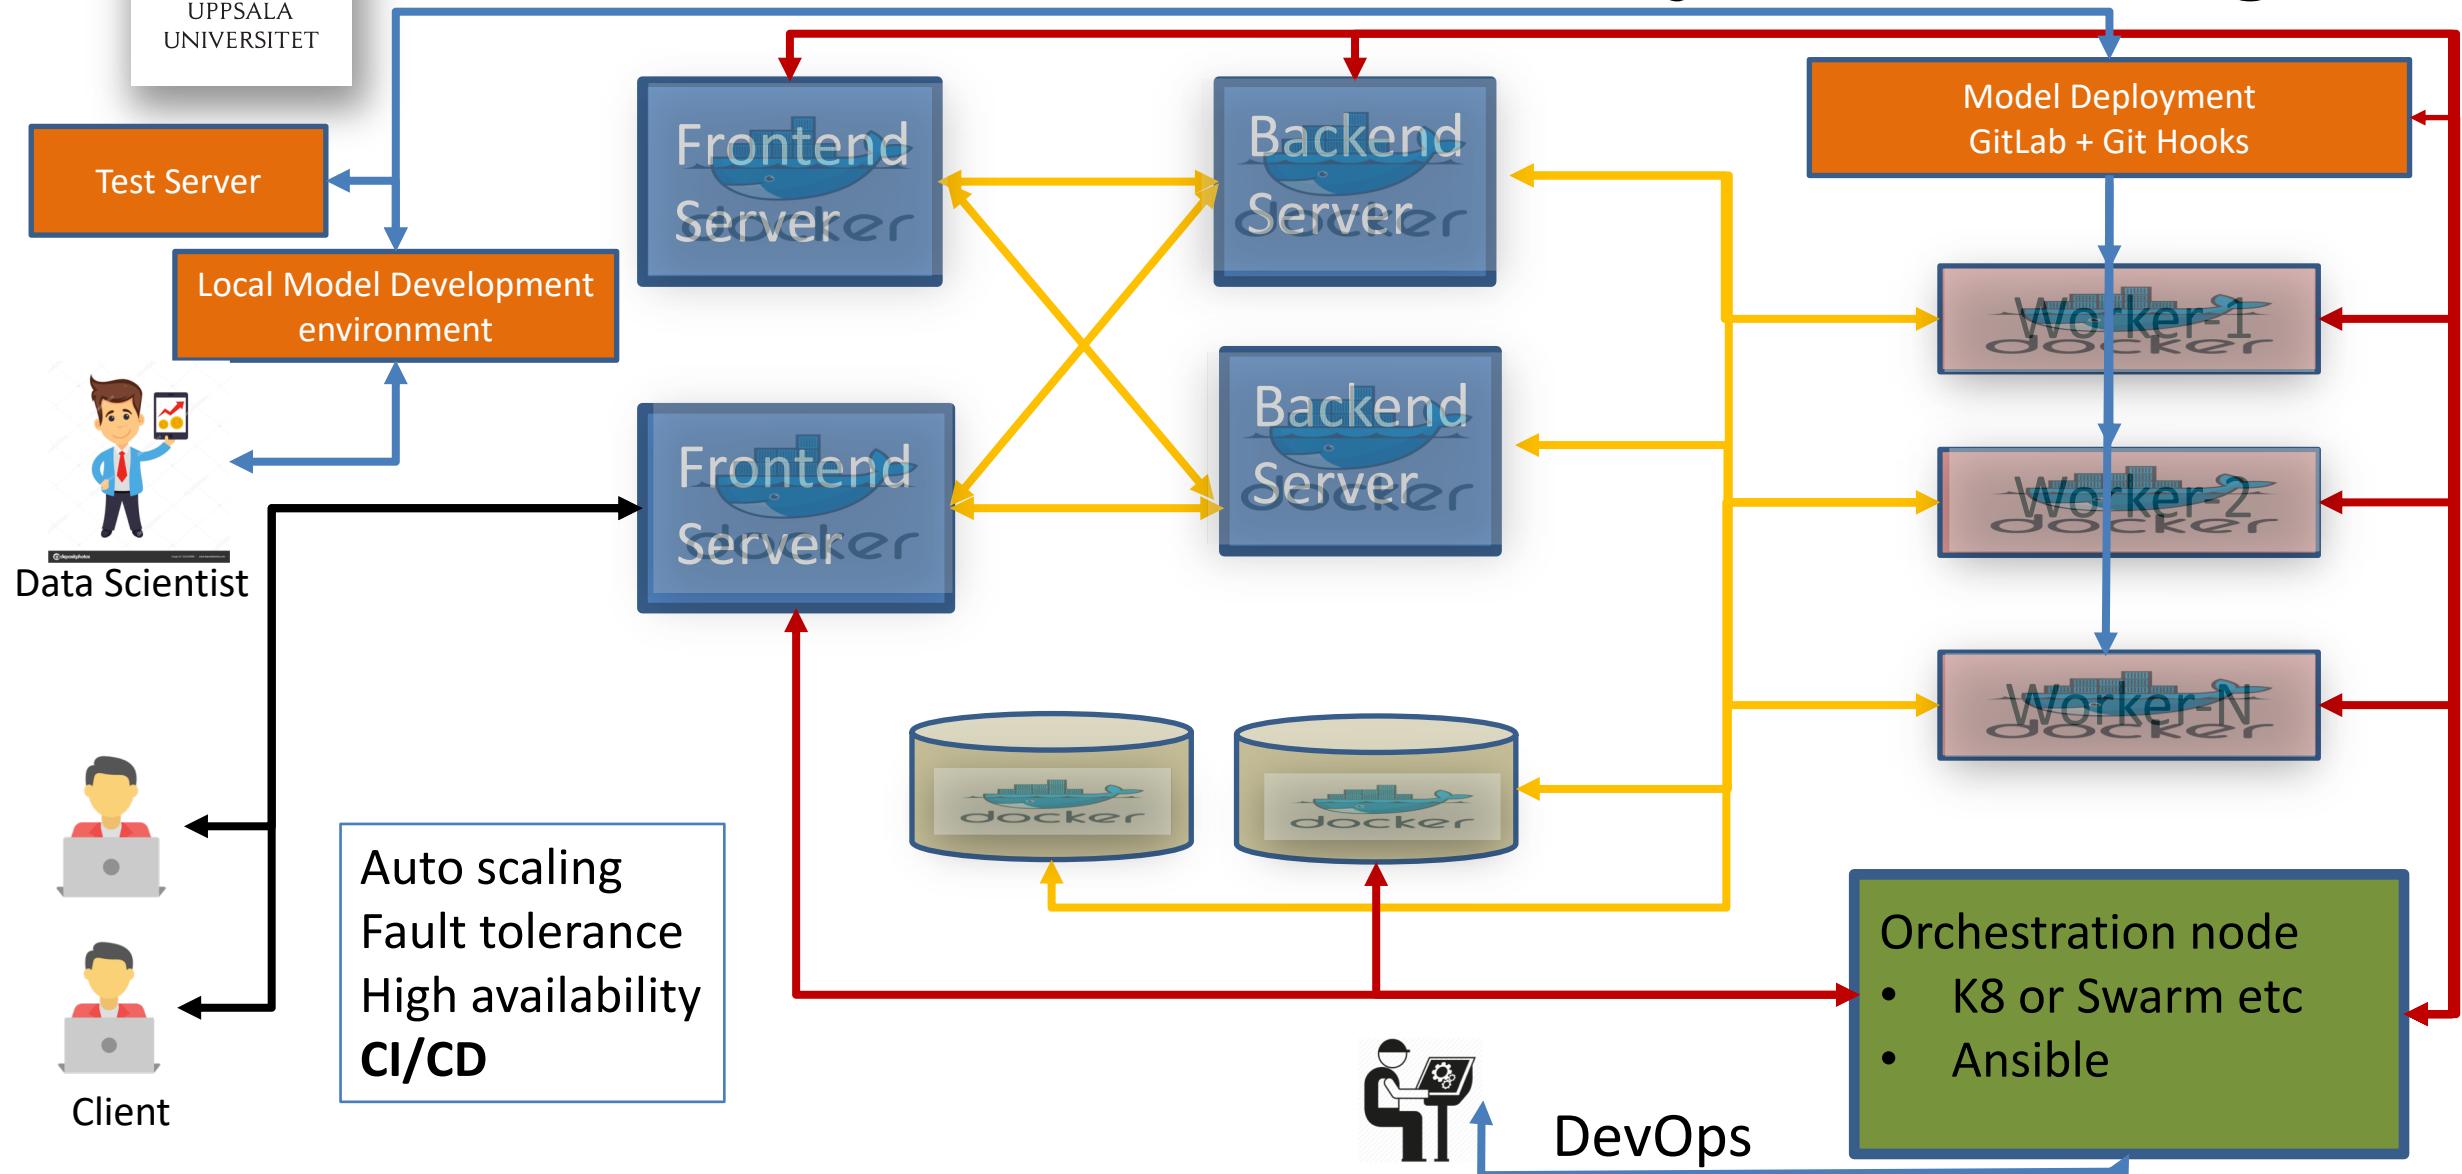


UPPSALA  
UNIVERSITET

Physical or Virtual Machines



# Deployment stage - 3





UPPSALA  
UNIVERSITET

# Deployment stage - 4



Geographically distributed execution environment

# Assignment

- For a simple neural network example, we have used PIMA Indian Diabetes Prediction dataset
  - <https://data.world/data-society/pima-indians-diabetes-database>
  - <https://towardsdatascience.com/pima-indian-diabetes-prediction-7573698bd5fe>
- The objective is to predict based on diagnostic measurements whether a patient has diabetes or not
- Neural network model is used to make the predictions

<https://machinelearningmastery.com/tutorial-first-neural-network-python-keras/>



# Tasks

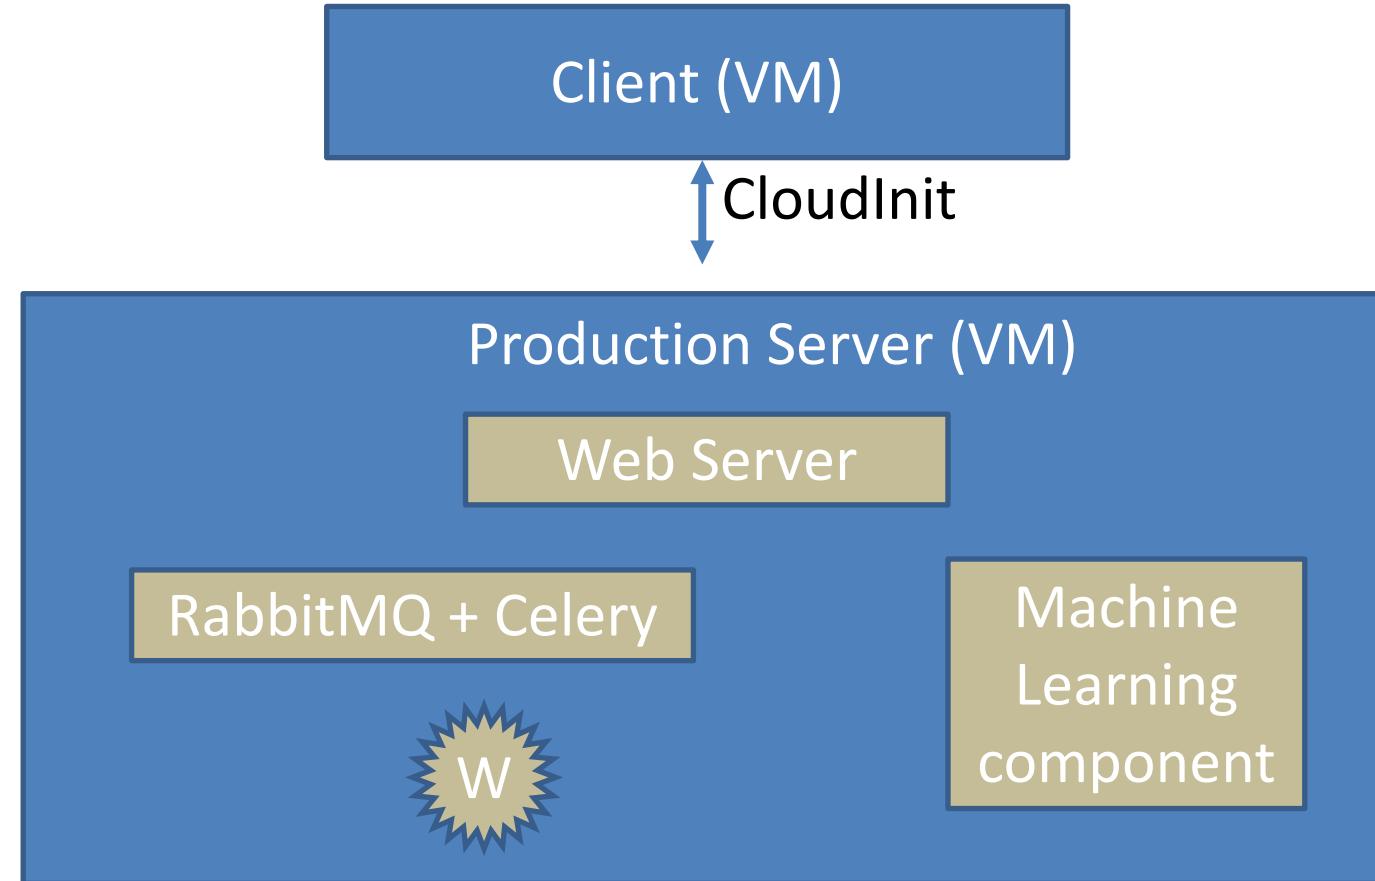
1. Dynamic contextualization using CloudInit package
  - Single node deployment
  - Containerized application deployment using docker-compose
  
2. Multi-node model serving setup
  - Model serving using Ansible playbooks
  - Continuous integration and delivery using Git Hooks



UPPSALA  
UNIVERSITET

# Task-1

Client contacts the cloud via OpenStack APIs and start a server by passing contextualization script.



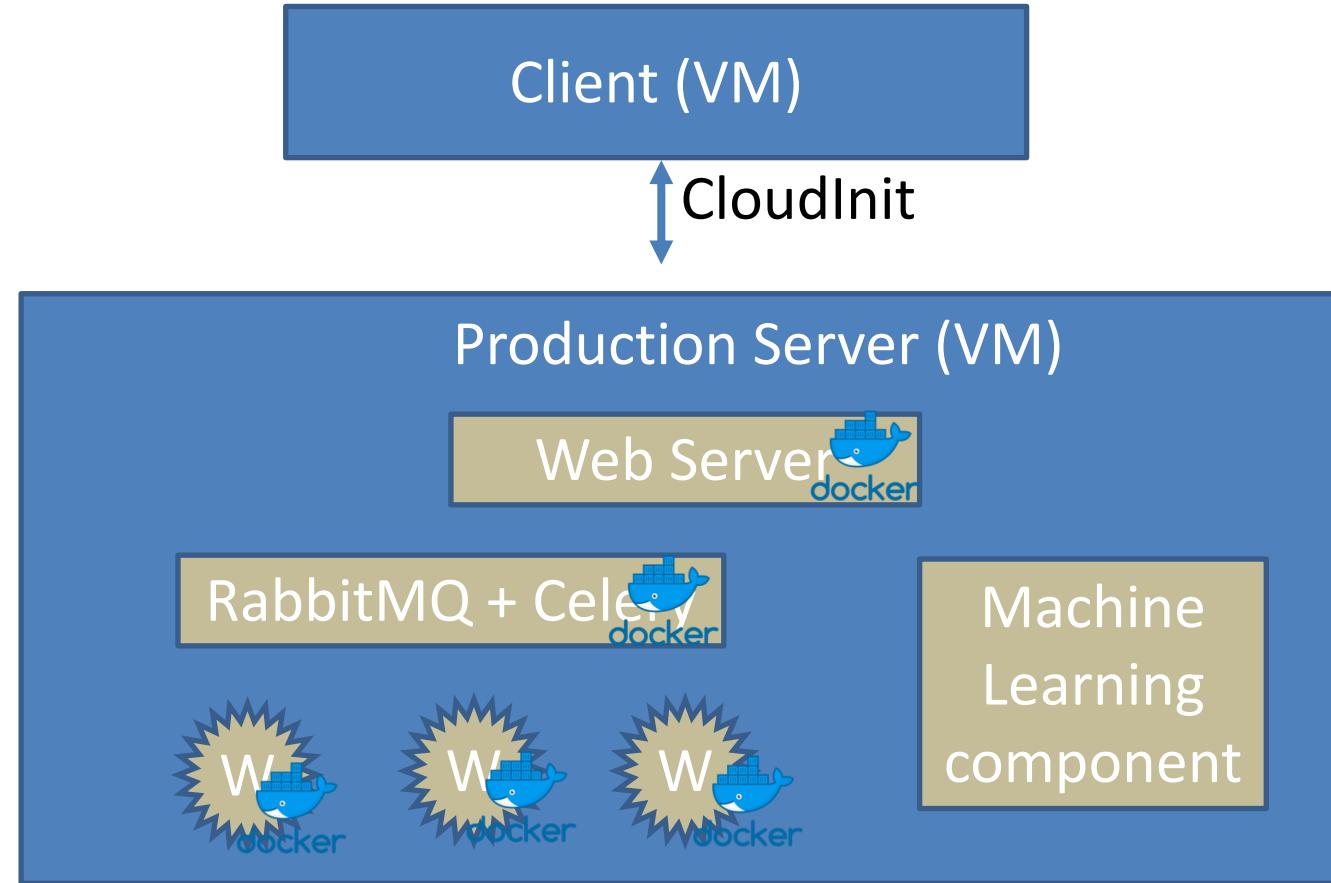
Production server hosts the complete application and provide a web interface to access the results.

Your task is to

1. start and configure a Client VM
2. send an API request from Client to initiate the Production Server
3. access the predictions via the web interface
4. answer the questions related to the Task-1 (available in Git repo)

# Task-2

Client contacts the cloud via OpenStack APIs and start a server by passing contextualization script.

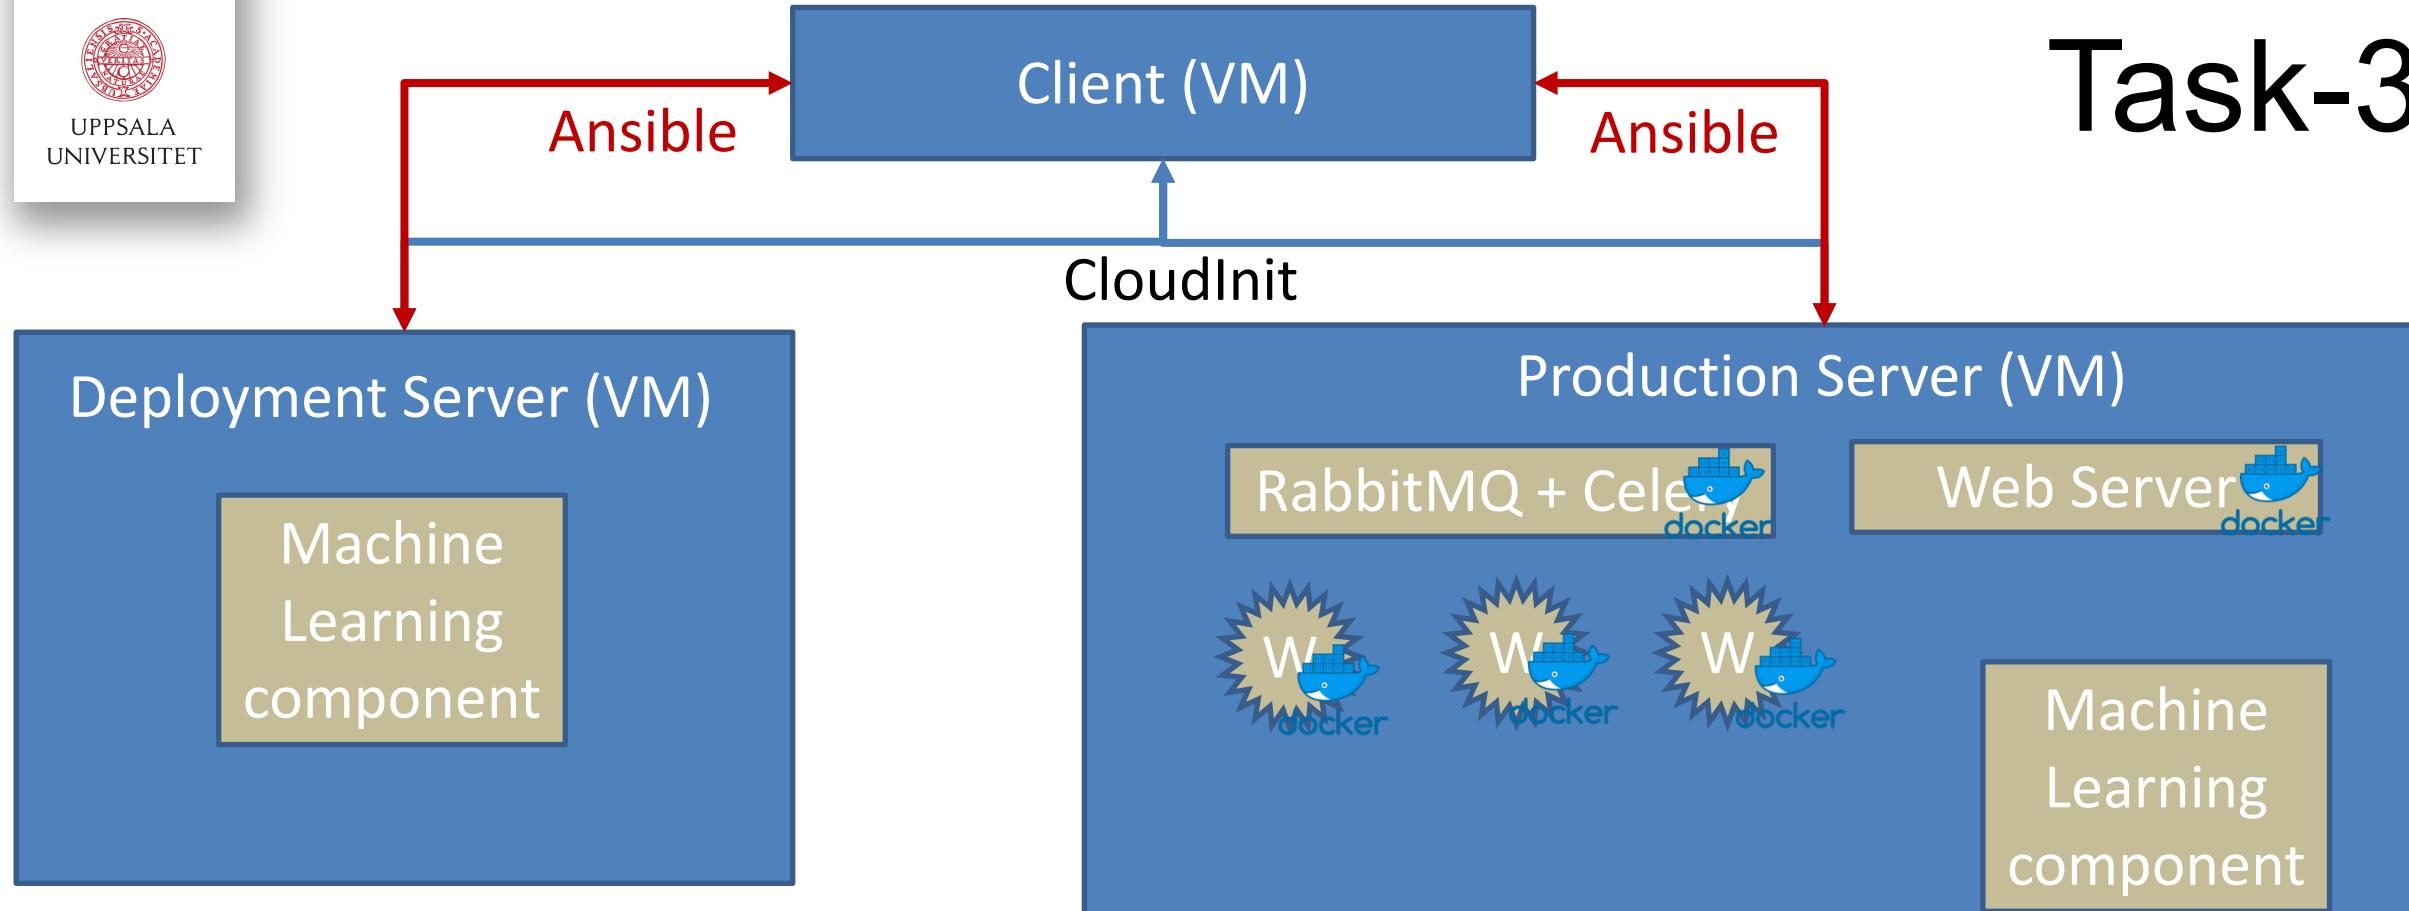


Production server hosts the complete application and provide a web interface to access the results.

Your task is to

1. send an API request from Client to initiate a **new** Production Server
2. access the predictions via the web interface
3. add more workers using docker commands
4. answer the questions related to the Task-2 (available in Git repo)

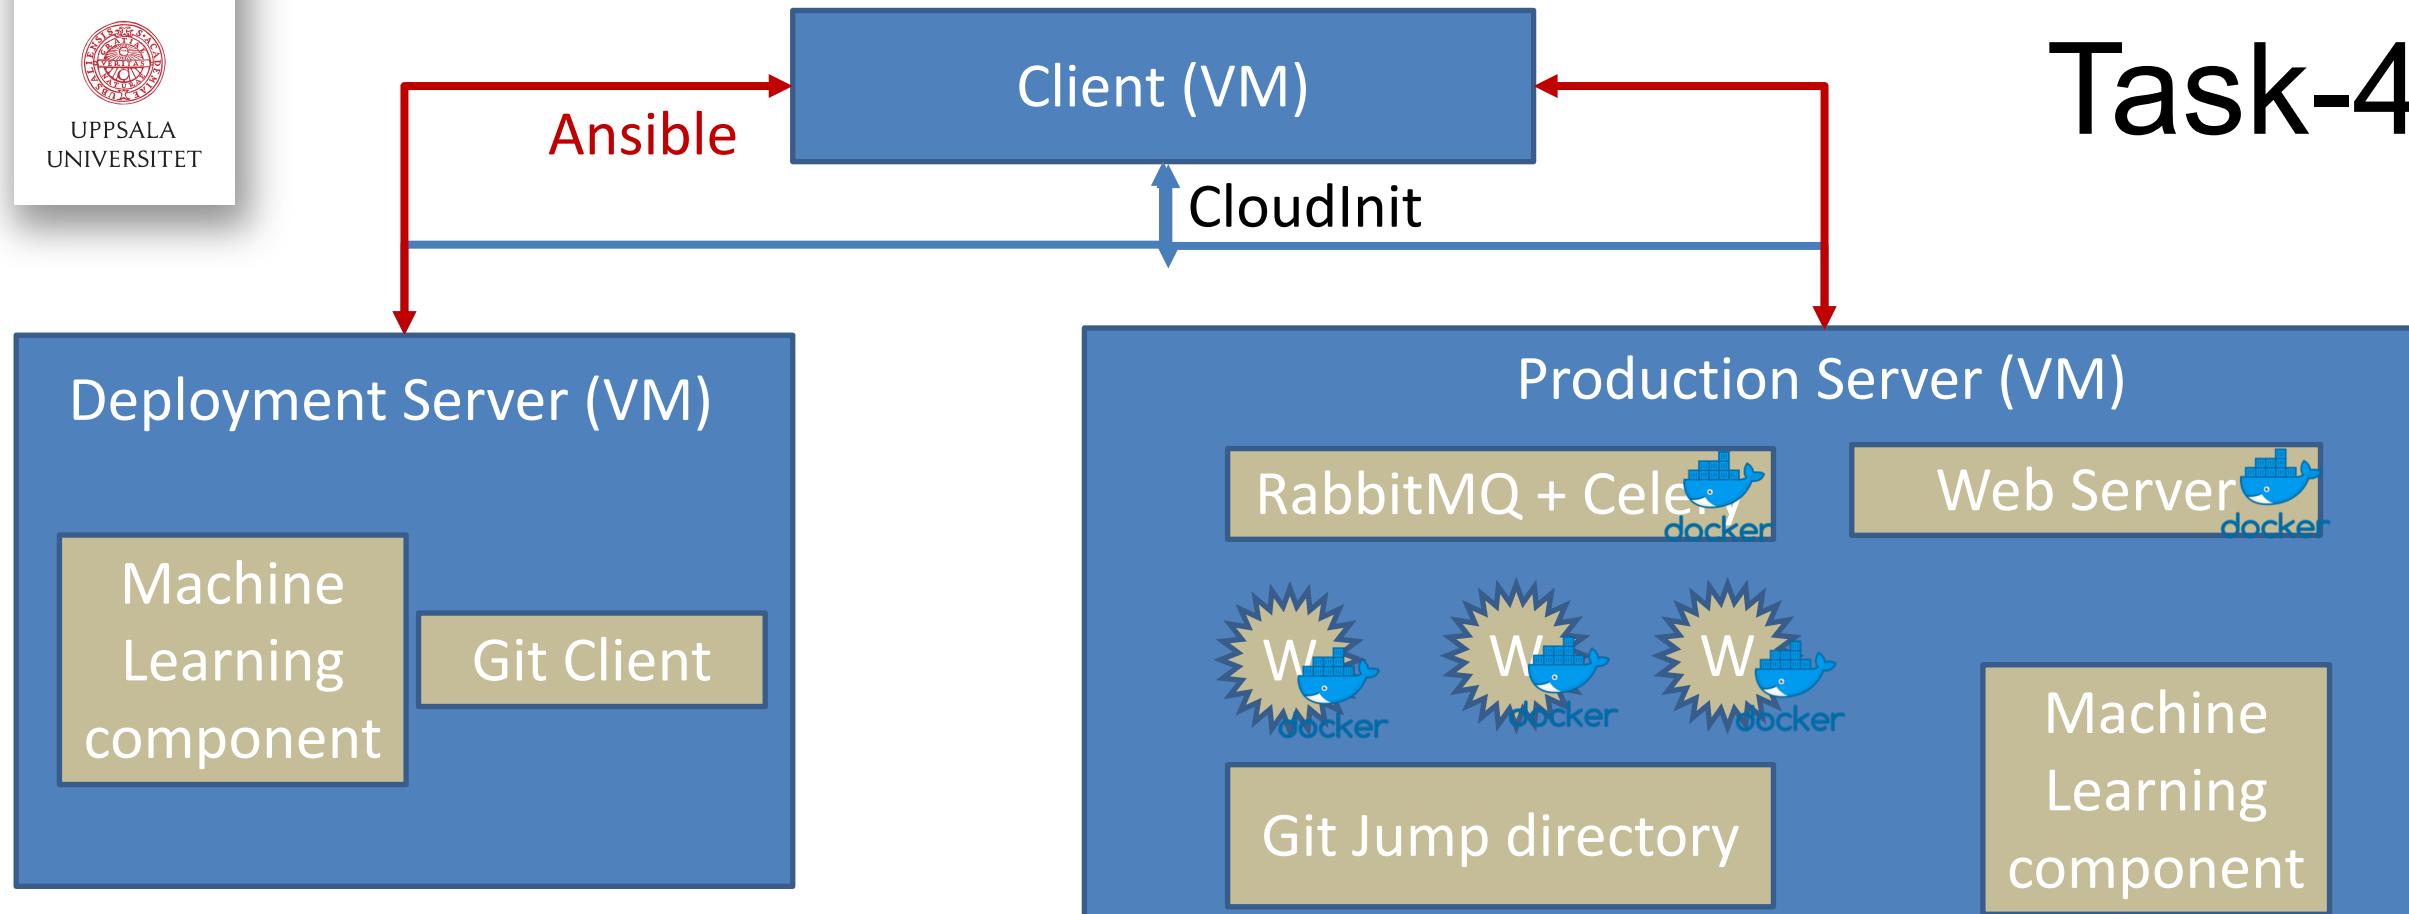
# Task-3



Your task is to

1. send an API request from Client to initiate both Production and Deployment Servers
2. contextualize the servers using Ansible playbook
3. access the predictions via the web interface
4. answer the questions related to the Task-3 (available in Git repo)

# Task-4



Your task is to

1. use the environment created in Task-3
2. Setup the Git Hook communication between Deployment and Production Servers
3. Make the changes in the machine learning code and access changes via web interface.
4. Answer the questions related to the Task-4 (available in Git repo)



UPPSALA  
UNIVERSITET

# Case Study: Stackn

<https://github.com/scaleoutsystems/stackn>

<https://scaleoutsystems.com/stackn>





UPPSALA  
UNIVERSITET

# Stackn

- STACKn is an open source collaborative ML platform
- It provides a highly flexible, cloud-native open toolkit for full-stack data science projects and the entire machine learning lifecycle.
- Key benefits:
  - Continuous Analytics
  - End-to-end workflow orchestration
  - Simple deployment
  - Open source (Apache2)



UPPSALA  
UNIVERSITET

Stackn

Demo



UPPSALA  
UNIVERSITET

# Case Study: KubeNow

<https://github.com/kubenz/KubeNow>

On-demand virtual research environments using  
microservices



UPPSALA  
UNIVERSITET

# KubeNow

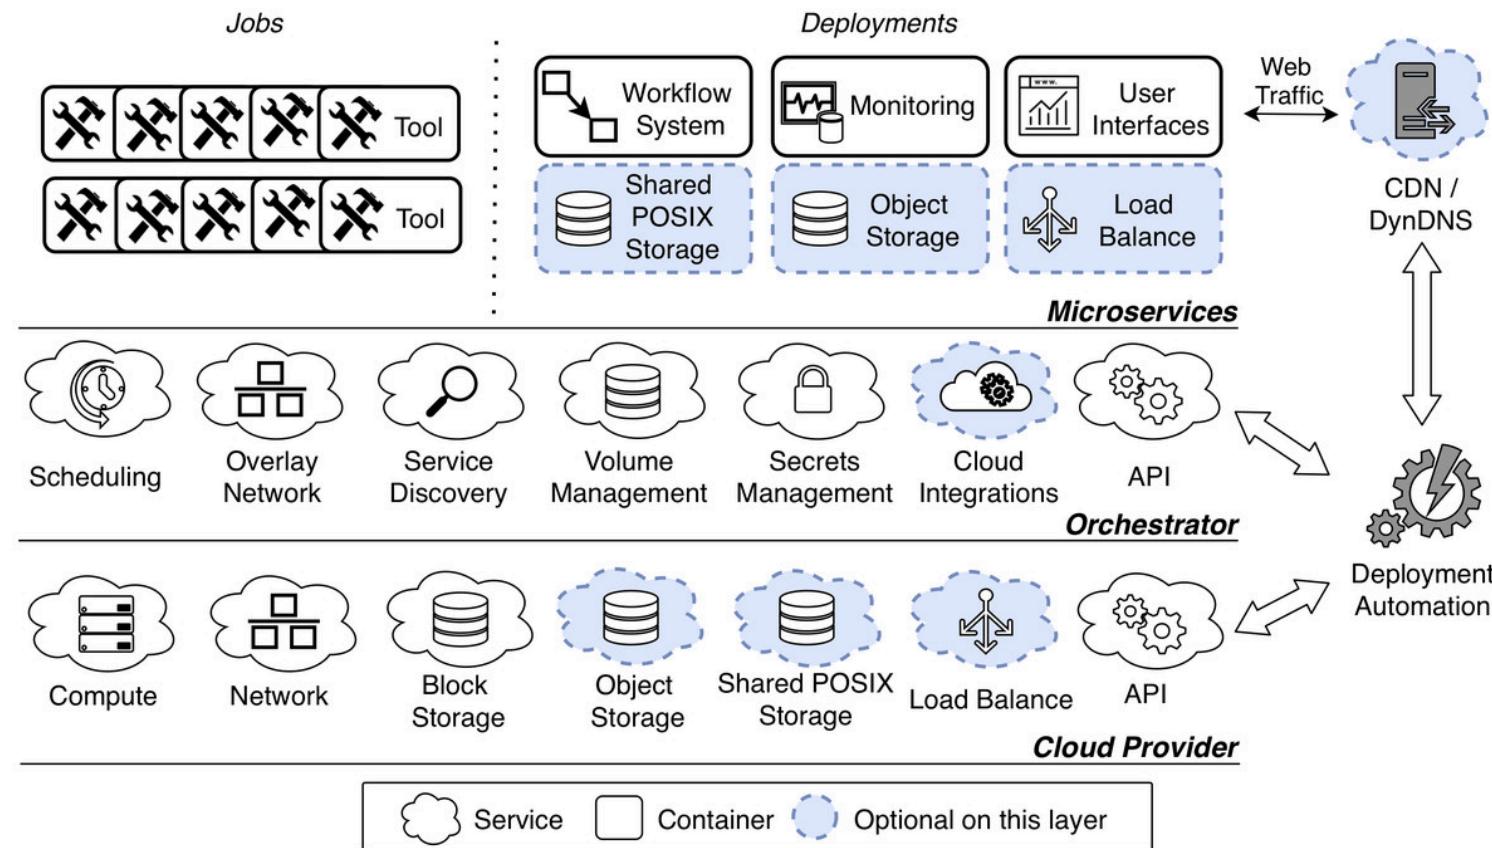
- Infrastructure-as-a-Service (IaaS) does not achieve the degree of flexibility required by the scientific community
- The framework enables scientific applications to run in a distributed orchestration platform as software containers, referred to as on-demand, virtual research environments (VRE)
- KubeNow is vendor-agnostic, enables on-demand VREs on the major cloud providers



UPPSALA  
UNIVERSITET

# KubeNow architecture

- **Deployment automation**
- **Content delivery network and dynamic domain name system**
- **Microservices**
- **Orchestrator**
- **Cloud Provider**

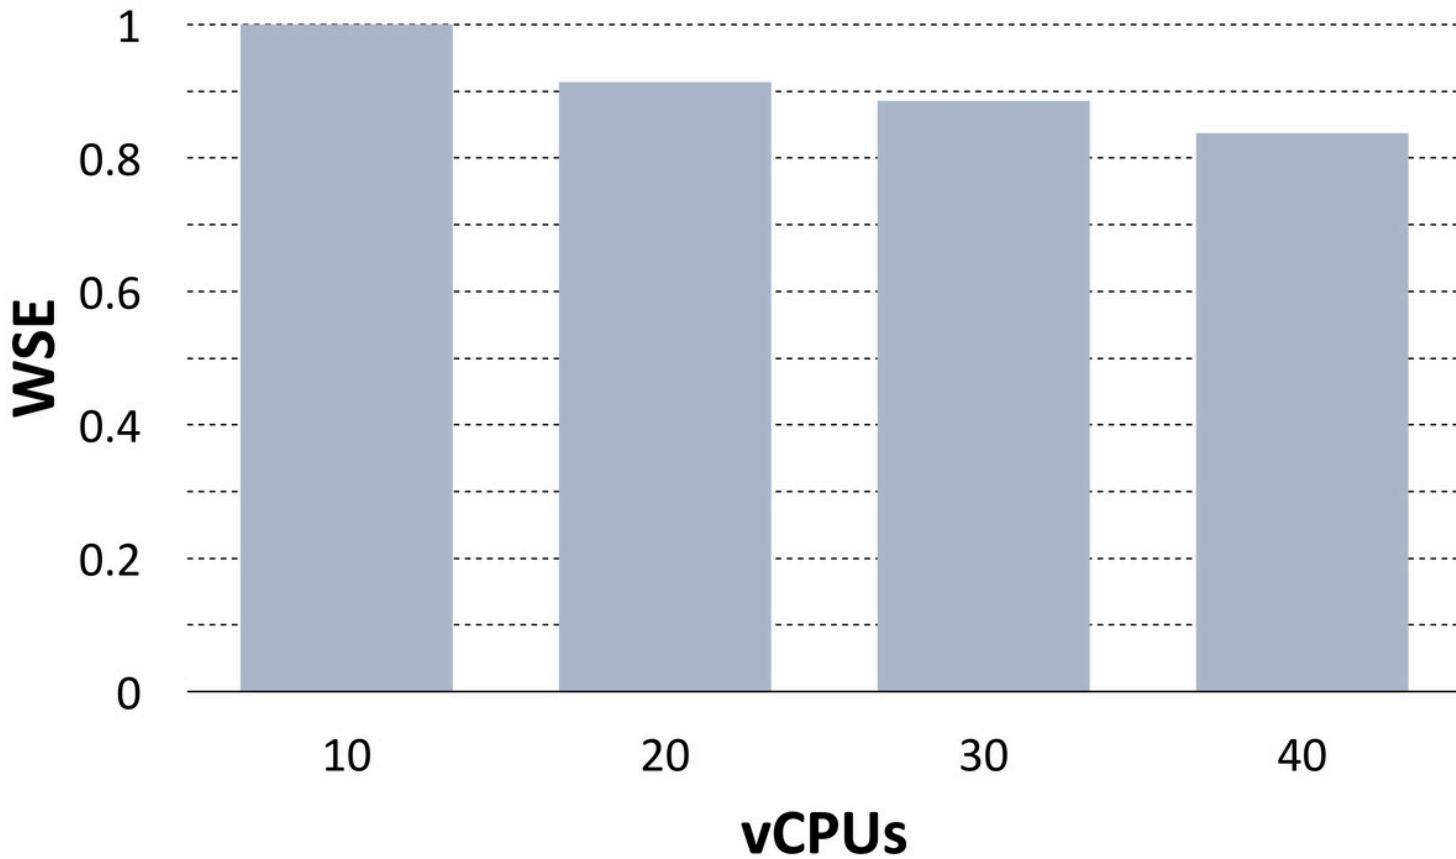


# KubeNow, framework evaluation

- Weak Scale Efficiency (WSE)

$$\text{WSE} = T_{10} / T_N$$

where  $T_{10}$  was the measured running time on 10 vCPUs and  $T_N$  was the measured running time on  $N$  vCPUs



# KubeNow tools

- **Docker:** the open source de facto standard container engine.
- **Kubernetes:** the orchestration platform that has collected the largest open source community.
- **GlusterFS:** an open-source distributed file system that provides both shared POSIX file spaces and object storage.
- **Traefik:** an open-source HTTP reverse proxy and load balancer.
- **Cloudflare:** a service that provides CDN and DynDNS.
- **Terraform:** an open-source IaC tool that enables provisioning at infrastructure level.
- **Ansible:** an open-source automation tool that enables provisioning of VMs and Kubernetes.

# KubeNow, framework evaluation

- Comparison
  - KubeNow
  - Kubespray

